The Delegation Dilemma: Economies of Scale, Accountability, and Polycentric Governance

Arin Mukhopadhyay
University of Edinburgh
Edinburgh, United Kingdom
a.mukhopadhyay-1@sms.ed.ac.uk

Aris Filos-Ratsikas
Imperial College London
London, United Kingdom
aris.filos-ratsikas@imperial.ac.uk

Panayiotis Danassis
University of Southampton
Southampton, United Kingdom
p.danassis@soton.ac.uk

arXiv:2610.21746v1 [cs.MA] 15 Oct 2026

Abstract

Many collective-action systems are neither fully decentralized nor fully centralized: agents voluntarily organize into institutions that share organizational costs and delegate decisions to representatives. Delegation can reduce individual costs, but it also transfers control to representatives whose private objectives need not align with those of their constituents. We study this tension in a repeated game with endogenous delegation to self-interested centers. A simple reduced-form model yields a *scale–accountability tradeoff*: larger constituencies offer greater cost-sharing benefits to members, yet the marginal value to a center of retaining any one constituent decreases as group size grows. Consequently, the prospect of individual exit can sustain non-exploitative delegation only up to an endogenous accountable constituency size. Additionally, we identify parameter regions in which fully delegated, accountable governance, when feasible, must be *polycentric*. We investigate these predictions through simulations with large language model (LLM) agents. Across five matched pairs of 30-agent simulations, polycentric structures repeatedly emerge but remain fluid, exhibiting cycles of concentration, exploitation, exit, and re-formation. We further find that models of different capacity exhibit similar willingness to assume leadership but differ substantially in their exercise of authority, with Sol initiating exploitation more frequently than Luna in matched decision states.

Keywords

Delegation, Polycentric Governance, Accountability, Coalition Formation, Large Language Model Agents, Multi-agent Governance

1 Introduction

Socially valuable activities often exhibit *economies of scale*: they are costly, or even infeasible, when undertaken individually, but become easier to perform as individuals organize into groups. Farmers may cooperate in production or irrigation, fishers may jointly manage a common resource, and communities may pool resources to provide local public services. Such activities naturally give rise to constituencies, with organizers or leaders who coordinate contributions, make collective decisions, and determine how the gains from cooperation are distributed. The benefits of scale have long been central to theories of clubs and collective action [3, 20], as well as to more recent work on endogenous group formation [1, 21].

At the same time, organizational scale may alter the incentives of those entrusted with authority. As a constituency grows, the centre's marginal organizational benefit from retaining any particular constituent may decline. A self-interested centre may therefore find it profitable to exploit some delegates, even when doing so risks their eventual departure. When losing a constituent has only a small effect on the centre's future organizational costs, the threat of exit provides a weaker deterrent against exploitation. Thus, the same economies of scale that encourage collective organization can also undermine the accountability of those exercising delegated authority and create scope for opportunism. This tension echoes the classical observation that organizational growth can concentrate decision-making authority [17].

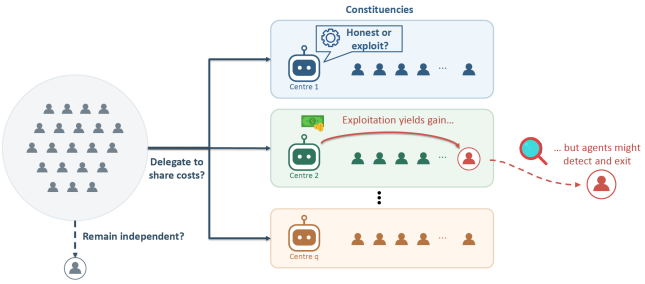

Figure 1: Overview of our delegation setting. Agents choose whether to remain independent or join a constituency (to share costs) led by a self-interested center. Each center decides whether to exploit its delegates for private gain. Delegates may detect exploitation and exit the delegation.

However, such opportunism does not necessarily imply the persistence of a single governing authority. Constituents who receive an exploitative treatment can instead exercise *exit* [13], by joining an alternative constituency or by volunteering to lead a new constituency themselves. Competition for constituents therefore constraints the size of the groups and their leaders' ability to benefit at the expense of their members. At the same time, groups cannot become arbitrarily small, because doing so sacrifices the economies of collective production. The resulting institutional structure is therefore neither individual nor fully centralized, but *polycentric*: multiple, relatively autonomous centers of organization coexist, compete for constituents, and may cooperate where doing so is mutually beneficial. This conception is closely related to the polycentric political systems of Ostrom, Tiebout, and Warren [20], and to Ostrom's later account of polycentric governance as an alternative to the dichotomy between centralized state provision and purely decentralized market exchange [19]. Examples of polycentricity are also prevalent in practice. For example, in experiments with members of Chilean territorial-use-right fisheries, Chávez et al. [4] find that while fishers frequently formed the grand coalition when monitoring was inexpensive, in fully decentralized settings without monitoring, the former coalitions were many more and significantly smaller. Thus, our general aim in this paper is to study

and understand the following phenomenon: polycentricity arises not merely because different tasks are efficiently performed at different scales, but because organizational scale itself changes the incentives of those entrusted with authority.

With the rapid development and employment of AI systems, individuals and organizations increasingly delegate their decisions to autonomous AI agents, and in particular to agents powered by large language models that act on their behalf. A growing literature has already begun to examine delegation to artificial agents in collective-action environments, e.g., see [6, 7, 18]. This raises the question of whether the institutional forces that generate polycentricity among human decision makers continue to operate when such decisions are delegated to AI agents. Will AI agents recognize the economic advantages of delegation and voluntarily form constituencies? How will self-interested AI centres exercise their authority as constituency size increases? Can the prospect of losing constituents discipline exploitation, or will centres anticipate that departing members can be replaced? Will agents establish competing centres when existing governance arrangements become unattractive? Finally, will different language models respond similarly to the same institutional incentives, or exhibit distinct patterns in their exercise of delegated authority?

More generally, we are interested to know whether interactions among autonomous AI agents reflect the endogenous checks on organizational scale that sustain polycentricity among human decision makers, or whether they instead generate qualitatively different organizational structures. Our aim is therefore not to assume that polycentricity is a stable outcome, but to understand the institutional incentives that make it possible and the strategic behaviour that shapes its emergence and persistence.

1.1 Our Contributions

We study whether polycentric arrangements emerge among autonomous LLM agents, how they evolve as constituencies gain and lose members, and whether different models exercise delegated authority in similar ways. Our main contributions are as follows:

(i) We develop a general model of voluntary delegation, economies of scale, and exit-based accountability (Figure 1). Our model is deliberately stylized, and intentionally abstracts away from the particulars of the task (e.g., whether it is a resource harvesting task or any other task) and the form of exploitation (e.g., whether that arises via some misallocation of resources etc). The aim of the model is to capture the fundamental tension between the efficiency gains from forming larger groups and the centers' incentives for exploitation, and the conditions under which this leads to polycentric structures. Participation depends on the average-cost saving created by joining a constituency, whereas accountability depends on the marginal organizational value that the center loses when one constituent exits. We derive a minimum viable constituency size, a maximum size at which probabilistic individual exit can deter exploitation, comparative statics, and an exact characterization of the number of centers compatible with fully delegated accountable governance. A fixed-cost specialization yields closed-form thresholds. Our theoretical results serve as baselines to ground our simulations.

(ii) We study endogenous governance formation among Large Language Model (LLM) agents. Across ten matched episodes with 30 agents, polycentric structures repeatedly emerge through both bottom-up formation and fragmentation of concentrated authority, but remain fluid as fee sharing attracts agents toward larger constituencies and exploitation induces exit. We additionally compare GPT-6 Luna and GPT-6 Sol (small vs. large model) pointwise on the same post-initialisation decision states. The two models are similarly willing to become centers but differ substantially in how they exercise delegated authority, with Sol initiating exploitation more often under matched conditions.

2 Related Work

Our model connects work on club and coalition formation with delegation and institutional accountability. Classical club theory studies how cost sharing and congestion determine desirable group size [3], while coalition-formation models study the endogenous emergence and stability of group structures [1, 21, 22]. Most closely related to our dynamic setting, Joshi [15] studies sequential club formation when agents' preferences depend on club size. Endogenous group formation has also been studied in collective-resource settings [4] and in environments with multiple centrally coordinated groups [10, 11]. Our upper size limit has a different source: larger groups dilute the disciplinary effect of losing any particular member rather than generating congestion or technological inefficiency.

A second line of literature studies delegation to representatives or intermediaries. Hamman et al. [9] show experimentally that groups may voluntarily delegate public-good decisions to elected members, and that poorly performing representatives can subsequently be replaced. Related work studies representation in public-good games [8, 12], delegation as a coordination mechanism [5], and delegation of common-pool-resource decisions to a stakeholder [2]. Our distinctive mechanism combines voluntary constituency formation with a self-interested center whose incentive to retain each constituent decreases with organizational scale. This links economies of scale directly to exit-based accountability.

Finally, recent work studies institutional mechanisms among LLM agents. Faulkner et al. [7] introduce elected leadership in LLM social groups, while Tewolde et al. [24] compare repetition, reputation, mediation, and contracting as cooperation-sustaining mechanisms. Hierarchical authority, elections, and oversight have also been studied in LLM public-good environments [23]. Rather than imposing a fixed governance institution, we ask whether self-interested AI agents endogenously form centralized, polycentric, or individual governance structures when organizational scale simultaneously creates economic benefits and weakens accountability.

3 Delegation – Accountability Model

In this section, we formulate a reduced-form game that isolates the institutional incentives associated with delegation, cost sharing, and exit. **All proofs can be found in Appendix A.**

There are $n \geq 2$ persistent agents. At any governance epoch, agents may remain under individual control or belong to a *constituency* consisting of one center and the agents who delegate control to it. Its size is $m \geq 1$, including the center; $m = 1$ corresponds to individual control. Delegation is binding for $L \geq 1$ operational rounds. After each nonterminal operational round, play

continues with probability $\delta \in (0,1)$, so a payoff s rounds in the future is reached with probability δ^s.

A constituency of size m incurs a per-round organizational cost $K(m) > 0$, shared equally among its members. Let

$$c(m) = \frac{K(m)}{m} \tag{1}$$

denote the resulting per-member cost. This formulation can represent shared infrastructure, administration, transaction costs, access or licensing costs, monitoring systems, or other technologies whose average cost changes with organizational scale. Delegates additionally incur a private per-round loss $\ell \geq 0$ from surrendering direct control. A center does not incur ℓ, because it retains decision authority. If the center exploits k delegates in a round, it receives kd and each targeted delegate loses d, where $d > 0$. Thus the reduced-form per-round payoffs are

$$u_C(m,k) = -c(m) + kd, \qquad u_D(m,z) = -c(m) - \ell - dz, \tag{2}$$

where $z \in \{0,1\}$ indicates whether a given delegate is exploited. An independent agent receives $u_I = -c(1)$. Exploitation is therefore a transfer between the center and a targeted delegate; its disciplinary effect operates through subsequent governance choices.

Assumption 1 (**Economies of scale**). *The per-member organizational cost $c(m)$ is weakly decreasing in m. Define the center's one-member retention value*

$$D(m) = c(m-1) - c(m), \qquad m \geq 2. \tag{3}$$

We additionally assume $D(m)$ is weakly decreasing in m, so the marginal organizational benefit of retaining one more constituent diminishes with scale.

The first part of Assumption 1 captures average-cost economies of scale. The second separates the mechanism relevant for accountability: as constituencies grow, *losing any particular constituent has a smaller effect on the center's own organizational cost.*

3.1 Delegation and Economies of Scale

Relative to individual control, membership in a constituency of size m saves the delegate

$$S(m) = c(1) - c(m) \tag{4}$$

per round in organizational cost. The net per-round advantage of non-exploitative delegation is therefore

$$B(m) = c(1) - c(m) - \ell. \tag{5}$$

Delegation is individually rational when $B(m) \geq 0$.

Proposition 1 (**General participation threshold**). *Under Assumption 1, if there exists a nontrivial constituency size for which delegation is individually rational, then the set of individually rational sizes is upward closed and its minimum is*

$$\boxed{m_{\min} = \min\{m \geq 2 : c(1) - c(m) \geq \ell\}} \tag{6}$$

If this set is empty, no finite constituency makes delegation individually rational through organizational cost sharing alone.

Fixed-cost Specialization. To obtain closed-form results, we specialize to the case in which the total organizational cost is independent of constituency size:

$$K(m) = \kappa, \qquad \kappa > 0, \tag{7}$$

so that $c(m) = \kappa/m$. This specification captures a fixed cost that can be shared among constituency members and is the environment implemented in our simulations.

Corollary 1 (**Minimum viable constituency size under fixed costs**). *Suppose $K(m) = \kappa$ and $0 \leq \ell < \kappa$. Non-exploitative delegation is individually rational if and only if*

$$m \geq \frac{\kappa}{\kappa - \ell} \tag{8}$$

The minimum nontrivial integer constituency size is

$$\boxed{m_{\min} = \max\left\{2, \left\lceil \frac{\kappa}{\kappa - \ell} \right\rceil\right\}} \tag{9}$$

If $\ell \geq \kappa$, no finite constituency makes delegation individually rational through fixed-cost sharing alone.

3.2 When Does Exit Discipline the Center?

We next model the center's temptation to exploit a constituent. The parameter d (defined in Eq. 2) is deliberately reduced form: it is the center's per-round private gain from exploiting one delegate and can represent appropriation, preferential allocation, extraction of rents, or any other benefit obtained by using delegated authority against a constituent's interest.

Assumption 2 (**Targeted exploitation and probabilistic exit**). *At the beginning of a governance epoch, a center in a constituency of size $m \geq 2$ can exploit one delegate, obtaining an additional payoff $d > 0$ in each of the L locked-in rounds while imposing the same per-round loss on that delegate. At the next governance opportunity, the exploitation is detected and results in that delegate's departure with probability $\alpha \in [0,1]$. Conditional on departure, the delegate is not replaced over the continuation horizon, so the change in the center's per-member organizational cost persists. The deviation does not otherwise alter the continuation probability.*

The parameter α is a reduced-form accountability probability. If detection occurs with probability p_{det} and a delegate exits conditional on detection with probability $p_{\text{exit}|\text{det}}$, then $\alpha = p_{\text{det}} p_{\text{exit}|\text{det}}$. This lets the theory encompass imperfect observability as well as imperfect behavioural response without changing the core incentive comparison.

Definition 1 (**Accountable constituency**). *A constituency of size m is* accountable under individual exit *if, under Assumption 2, the center weakly prefers non-exploitative behaviour to the targeted exploitation deviation.*

Proposition 2 (**General accountability condition**). *Under Assumptions 1–2, a constituency of size $m \geq 2$ is accountable under individual exit if and only if*

$$\boxed{\alpha \delta^L D(m) \geq d(1 - \delta^L)} \tag{10}$$

Because $D(m)$ is weakly decreasing, the set of accountable constituency sizes is downward closed. When the largest accountable size is finite,

define

$$m_{\max} = \max\left(\{1\} \cup \left\{m \geq 2 : \alpha\delta^L D(m) \geq d(1-\delta^L)\right\}\right). \tag{11}$$

The general condition separates two distinct effects of scale. Participation depends on the *average* organizational saving $c(1)-c(m)$, whereas accountability depends on the *marginal* retention value $D(m) = c(m-1) - c(m)$. Organizational scale can therefore make delegation increasingly attractive to constituents while simultaneously making each individual constituent less valuable to the center.

Corollary 2 (**Maximum accountable constituency size under fixed costs**). *Under $K(m) = \kappa$, condition* (10) *becomes*

$$\alpha\delta^L \frac{\kappa}{m(m-1)} \geq d(1-\delta^L), \tag{12}$$

or equivalently

$$m(m-1) \leq \frac{\alpha\kappa\delta^L}{d(1-\delta^L)}. \tag{13}$$

Hence the continuous accountability threshold is

$$\overline{m} = \frac{1 + \sqrt{1 + 4\frac{\alpha\kappa\delta^L}{d(1-\delta^L)}}}{2}, \tag{14}$$

and the largest integer constituency size supportable by individual exit is

$$\boxed{m_{\max} = \lfloor \overline{m} \rfloor} \tag{15}$$

If $m_{\max} = 1$, no nontrivial delegated constituency is accountable through this exit mechanism alone.

Interpretation. The upper threshold does not imply that constituencies larger than $m_{\max}$ cannot exist or that their constituents necessarily leave. It says only that, under the stated deviation and punishment technology, the expected loss generated by one constituent's possible future exit is insufficient to make the center prefer non-exploitative behaviour. Constituents may still prefer membership despite exploitation, and replacement, collective exit, reputation, auditing, legal sanctions, or other mechanisms can change the effective accountability constraint, sustaining accountability beyond $m_{\max}$.

3.2.1 ***Robustness of the Accountability Constraint.*** Two simple observations strengthen the interpretation of $m_{\max}$.

Lemma 1 (**Single-constituent exploitation is the hardest additive deviation**). *Suppose a center simultaneously exploits $k \in \{1, \ldots, m-1\}$ constituents, gaining kd per locked-in round, and each targeted constituent independently exits at the next governance opportunity with probability α. Under Assumption 1, if the one-constituent accountability condition* (10) *holds, then every such additive k-constituent deviation is also deterred.*

Lemma 2 (**The beginning of an epoch is the hardest time to deter exploitation**). *Suppose exploitation begins when $s \in \{1, \ldots, L\}$ locked-in rounds remain. If the accountability condition holds for $s = L$, it holds for every later point in the epoch with $s < L$.*

Table 1: Comparative statics for the fixed-cost participation and accountability thresholds.

Increase	$m_{\min}$	$m_{\max}$	Predicted institutional effect
κ	$\downarrow$	$\uparrow$	Delegation viable over a wider size range
ℓ	$\uparrow$	–	Harder to induce delegation
d	–	$\downarrow$	Stronger agency problem; smaller accountable centers
α	–	$\uparrow$	More reliable detection and exit strengthen accountability
δ	–	$\uparrow$	Future exit matters more
L	–	$\downarrow$	Longer lock-in weakens exit discipline

3.3 How Institutional Parameters Shape Delegation

The general formulation shows that participation is governed by average cost savings whereas accountability is governed by the marginal retention value of one constituent. We now characterize how the two institutional thresholds move with the parameters and give closed-form comparative statics for the fixed-cost specialization. Table 1 summarizes the main results.

Theorem 1 (**Comparative statics under fixed costs**). *Assume $\kappa > 0$, $d > 0$, $0 < \delta < 1$, $L \geq 1$, $0 < \alpha \leq 1$, and $0 \leq \ell < \kappa$. Let*

$$\underline{m} = \frac{\kappa}{\kappa - \ell} \quad \text{and} \quad \overline{m} = \frac{1 + \sqrt{1 + 4C}}{2}, \qquad C = \frac{\alpha\kappa\delta^L}{d(1-\delta^L)}.$$

Then

$$\frac{\partial \underline{m}}{\partial \kappa} \leq 0 \quad (\text{strict if } \ell > 0), \qquad \frac{\partial \underline{m}}{\partial \ell} > 0, \tag{16}$$

$$\frac{\partial \overline{m}}{\partial \kappa} > 0, \qquad \frac{\partial \overline{m}}{\partial d} < 0, \tag{17}$$

$$\frac{\partial \overline{m}}{\partial \alpha} > 0, \qquad \frac{\partial \overline{m}}{\partial \delta} > 0, \tag{18}$$

and $\overline{m}(L+1) < \overline{m}(L)$. Consequently, after integer rounding, $m_{\min}$ is weakly decreasing in κ and weakly increasing in ℓ, while $m_{\max}$ is weakly increasing in κ, α, and δ and weakly decreasing in d and L.

Corollary 3 (**Stronger fixed-cost economies widen the viable interval**). *For fixed $\ell, d, \alpha, \delta, L$, increasing κ weakly lowers $m_{\min}$ and weakly raises $m_{\max}$. Thus stronger fixed-cost economies expand the set of constituency sizes that are simultaneously individually attractive and accountable under exit.*

Corollary 4 (**More reliable exit supports accountability at larger scales**). *For fixed κ, d, δ, L, $m_{\max}$ is weakly increasing in α. If $\alpha = 0$, individual exit supplies no expected punishment and no nontrivial constituency is accountable through this mechanism alone.*

Corollary 5 (**Long lock-in weakens accountability**). *For fixed $\kappa, d, \alpha, \delta$, $m_{\max}$ is weakly decreasing in governance-epoch length L. Moreover, for $0 < \delta < 1$, $\overline{m} \to 1$ as $L \to \infty$.*

3.4 The Emergence of Polycentric Governance

The two thresholds $m_{\min}$ and $m_{\max}$ characterize constituency sizes that are simultaneously individually attractive and accountable. Importantly, the following result depends only on the existence of these two thresholds; it is therefore not specific to the fixed-cost specialization.

Definition 2 (**Fully delegated accountable partition**). *A partition of n agents into q constituencies with integer sizes $m_1, \dots, m_q$ is* fully delegated and accountable *if every constituency is nontrivial and satisfies*

$$m_{\min} \le m_j \le m_{\max} \qquad \text{for all } j = 1, \dots, q.$$

Theorem 2 (**Feasible number of accountable centers**). *Suppose $2 \le m_{\min} \le m_{\max}$. A fully delegated accountable partition of n agents into exactly q constituencies exists if and only if*

$$\boxed{\left\lceil \frac{n}{m_{\max}} \right\rceil \le q \le \left\lfloor \frac{n}{m_{\min}} \right\rfloor} \tag{19}$$

Consequently, if the interval in (19) *is nonempty and $n > m_{\max}$, every fully delegated accountable governance structure is polycentric, with at least*

$$q_{\min} = \left\lceil \frac{n}{m_{\max}} \right\rceil \ge 2 \tag{20}$$

centers.

Corollary 6 (**Primitive condition for necessary polycentricity under fixed costs**). *Under the fixed-cost specialization, define*

$$m_{\min} = \max\left\{2, \left\lceil \frac{\kappa}{\kappa - \ell} \right\rceil\right\}, \qquad m_{\max} = \left\lfloor \frac{1 + \sqrt{1 + 4\frac{\alpha\kappa\delta^L}{d(1-\delta^L)}}}{2} \right\rfloor.$$

If $m_{\min} \le m_{\max} < n$ and $\lceil n/m_{\max} \rceil \le \lfloor n/m_{\min} \rfloor$, then every fully delegated accountable governance structure is necessarily polycentric.

Takeaway. Theorem 2 is a feasibility and necessity result, not an equilibrium-selection claim. It identifies a region, under benchmark conditions, in which *economies of scale rule out very small delegated institutions* while *accountability rules out a grand center*. Whether decentralized strategic agents discover and sustain such structures once formation and replacement evolve endogenously is the question we examine experimentally.

3.5 Structural Polycentricity vs. Behavioural Diversity

Our theory characterizes when accountable delegation requires authority to be distributed across multiple centers. Such *structural polycentricity*, however, does not necessarily imply diversity in how those centers behave. We therefore distinguish between structural polycentricity, which describes the distribution of authority across centers, and *behavioural diversity*, which describes whether distinct centers make different decisions when facing comparable strategic conditions. In simple terms, a system can be highly polycentric in its formal organization while remaining behaviourally concentrated if its centers respond similarly to the same situations. This distinction is particularly relevant when centers are controlled by LLMs. Recent work documents substantial intra-model repetition and inter-model homogeneity in LLM responses to open-ended prompts [14]. Although these findings concern generative behaviour rather than strategic governance, they raise the possibility that formally independent LLM centers may nevertheless respond similarly to the same incentives. A governance system may therefore be highly polycentric in its institutional structure while remaining behaviourally concentrated if its centers repeatedly make similar exploitation decisions or similar mistakes. Our theoretical results do not require behavioural heterogeneity across centers, but this distinction motivates our empirical examination of how language-model agents exercise authority under comparable strategic conditions, including whether different models exhibit similar exploitation decisions and reasoning patterns.

4 Polycentricity in Environments of AI Agents

In this section, we study endogenous governance formation among large language model (LLM) agents. Our objective is to study whether polycentricity emerges naturally via their interactions, and whether the experimental results echo those of the theory developed in Section 3. We start by describing the design of our simulation.

4.1 Simulation Design

Reduced-form Environment. The simulations implement the reduced-form game of Section 3, with perfect detection and compulsory exit following exploitation ($\alpha = 1$). Agents can establish new centres, step down as centres, and voluntarily join or leave constituencies at governance transitions. At each transition, exploited delegates first leave their centres. Existing centres then decide whether to continue, remaining delegates decide whether to stay, and free agents may volunteer to become new centres. If multiple agents volunteer, one is selected uniformly at random. Once the set of centres is determined, free agents choose between individual control and delegation to an eligible centre. A departing delegate may volunteer during the same transition.

Agents know the governance rules and observe the public constituency structure and their own actions and outcomes, but not other agents' private exploitation outcomes or reasoning. Each decision is accompanied by a history of the agent's observations, organized by governance epoch. Model-generated explanations are recorded for analysis but are neither shared with other agents nor included in subsequent memory.

Matched Initialization Experiment. We ran five matched pairs of simulations, each with $n = 30$ GPT-6 Luna agents. Each pair compared two starting conditions. In the *centralized initialization*, one randomly selected agent began as the centre of a group containing all 30 agents. In the *decentralized initialization*, all agents began independently and could volunteer to become a centre. One volunteer was selected at random, and the other agents could choose whether to join. After initialization, both conditions followed the same governance rules described in the previous paragraph.

The two simulations in each pair shared the same random draws but used separate model calls and memories. Each pair used a different random seed. We set $(\kappa, \ell, d, \delta, L) = (1, 3/5, 2/5, 99/100, 5)$, which gives the theoretical benchmark interval $3 \le m \le 7$ for $\alpha = 1$, according to Corollaries 1 and 2. We use this interval to classify constituency sizes. Appendix C.1 has a detailed parameter sweep.

If an agent leaves a centre (e.g., due to exploitation), it cannot immediately rejoin. Before each later affiliation decision, the option to rejoin becomes available with probability 1/2 and remains available unless the agent rejoins and leaves again. Rejoining is always voluntary. Agents see which centres they can join, but do not know the reinstatement rule or its random outcomes. This mechanism captures membership renewal in a finite population of persistent agents. Without it, permanent restrictions on rejoining would gradually limit agents' choices and could cause fragmentation unrelated to their strategic decisions. In real organizational settings, populations evolve as new members arrive and former members reconsider their affiliations. Probabilistic reinstatement provides an approximation of such membership renewal without introducing explicit population turnover, allowing constituencies to recover, reorganize, and compete for members over repeated governance epochs.

Each simulation contains 20 complete epochs. All agents received the same rules and one of five decision prompts: centre continuation, delegate release, volunteering, affiliation, or exploitation. Each prompt included the agent's current observations, available actions, and private memory. We used GPT-6 Luna with medium reasoning effort and required agents to provide both a structured decision and a private explanation. All prompts are provided in Appendix C.3, and the source code can be found in the supplement.

Outcome Measures and Point-wise Comparison. We call a constituency *viable* when $m_{\min} \le m \le m_{\max}$,[1] and an epoch *polycentric* when at least two viable constituencies coexist. For each operational round, we record each agent's constituency size and whether they were exploited. To measure how cooperation gains are distributed, we compare each agent's payoff with the independent-control payoff of $-\kappa$ and attribute the difference to centres or delegates according to their role.

For the pointwise comparison, Sol evaluated each post-initialization Luna decision state using the same rules, observations, available actions, and history, without seeing Luna's decision or reasoning. Sol's responses did not affect the Luna trajectories. We compared 4,551 volunteering decisions and 1,988 exploitation decisions across the two models. The five random seeds were independent, but decisions within each simulation were not; we therefore report the results descriptively. This comparison captures differences between the models in identical decision states, rather than outcomes from independently evolving Sol societies.

5 Simulation Results

(1) Luna agents repeatedly formed polycentric arrangements under both initialization conditions, but **constituency membership remained fluid** as fee-sharing incentives encouraged concentration and exploitation induced departures.
(2) Although delegation generated organizational cost savings, **centres captured most of these gains**. This concentration was particularly pronounced under centralized initialization, where exploitation left delegates collectively worse off than under individual control.

[1]This classification identifies constituency sizes satisfying the analytical bounds, rather than establishing that the centre behaves non-exploitatively.

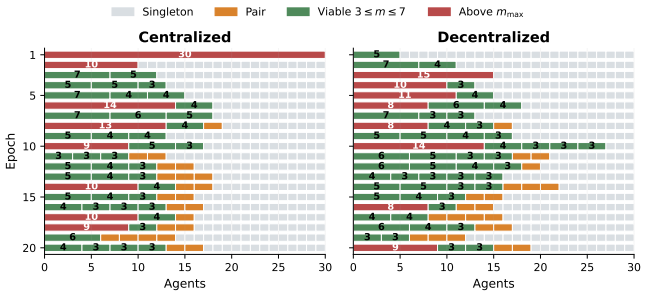

Figure 2: Constituency partitions at the start of each epoch. Left: centralized initialization, Right: decentralized initialization Each panel partitions all 30 agents; labels show constituency sizes of at least three. Multiple viable constituencies formed, dissolved, and re-formed under both initialization conditions. Appendix 8 shows all five repeats.

(3) Matched-state comparisons show that Luna and Sol were similarly willing to volunteer for leadership but **differed substantially in their exploitation decisions** and in the considerations emphasized in their explanations.

5.1 Polycentricity Emerges but Remains Fluid

Polycentricity emerged through both formation and fragmentation. All ten trajectories became polycentric, in all ten episodes and in 117 of the 200 observed epochs. This occurred under both centralized (55/100 epochs) and decentralized (62/100 epochs) initialization. Each trajectory was polycentric in 8–15 of its 20 epochs. The routes to polycentricity differed markedly for each of the two initialization conditions (Figure 2). Across both conditions, constituencies repeatedly expanded beyond the analytical accountability bound, lost members, and re-formed around existing or new centres. Thus,the agents did not simply select and maintain a partition within the theoretically compatible interval; polycentricity emerged through an process of concentration, exploitation, exit, and reconstruction.

Table 2 compares the two initialization conditions across the full trajectories. Decentralized initialization produced slightly more polycentric epochs (62% versus 55%) and greater agent-round exposure to constituencies within the theoretical size interval (29.2% versus 24.6%). More importantly, exploitation exposure after the first epoch was lower (42.2% versus 52.9%). These descriptive differences suggest that the initial distribution of authority may influence subsequent governance dynamics, even though both conditions generated repeated formation and fragmentation of constituencies.

Exploitation often triggered the departure of targeted delegates, however these losses were not necessarily permanent: centres could attract new delegates, and departed agents could regain the opportunity to rejoin. Indeed, some centre explanations *explicitly anticipated membership recovery when deciding whether to exploit.*

Every trajectory remained polycentric for at least two consecutive epochs, and one did so for ten. Some constituencies even grew to 14 agents, way above the theoretical bound. Nevertheless, larger constituencies subsequently fragmented, while smaller constituencies repeatedly lost members to competing centres. The observed dynamics were therefore characterized by recurring concentration,

Table 2: Polycentricity and exploitation by initialization condition, aggregated over five matched pairs of 30-agent simulations (20 governance epochs per episode). An epoch is classified as polycentric when at least two constituencies have sizes within the theoretical benchmark interval $3 \leq m \leq 7$. Exploitation rates are calculated in epochs 2–20 to avoid the distinct initialization epoch.

Outcome	Centralized	Decentralized
Polycentric epochs	55/100	62/100
Agent-rounds within theoretical bounds	24.6%	29.2%
Exploited delegate-rounds after epoch 1	52.9%	42.2%

exploitation, exit, and re-formation, rather than convergence to a persistent governance structure.

Centres Captured Most of the Surplus. Although fee sharing creates aggregate gains relative to individual control, exploitation can redistribute these gains from delegates to centres. To examine who benefits from delegation, we measured each agent's payoff gain relative to independent control and aggregate these gains separately for centres and delegates.

For the centralized initialization, centres captured 104.4% of aggregate delegation surplus, leaving delegates with a negative share of −4.4%. A share exceeding 100% is possible – i.e., delegates can be collectively worse off than under independent control – because exploitation transfers payoff from delegates to centres. Delegate surplus was negative in four of the five centralized repeats. For the decentralized initialization, centres captured 91.9% of the surplus, while delegates retained 8.1%, with positive delegate surplus in all five repeats. Thus, although centres received most of the organizational gains under both conditions, the distribution of these gains differed based on initial governance structures (see Appendix 6).

5.2 Exercise of Authority

We next investigate whether agents' governance decisions differ across models of varying capabilities. To this end, we compare a higher-capacity model (GPT-6 Sol) with a smaller model (GPT-6 Luna), evaluating their interest in assuming authority and their exploitation decisions under identical strategic conditions.

The models were similarly willing to become centers, but differed substantially in how much they exploited. In 4,551 matched volunteering states, Luna and Sol volunteered in 96.53% and 99.17% of cases, respectively, agreeing in 96.04% of states. The difference was substantially larger for exploitation. In 1,988 matched centre-decision states, Sol initiated new exploitation[2] in 41.80% of cases, compared with 31.79% for Luna, *a 10 percentage-points difference*. The models agreed on whether to exploit in 65.74% of these states. This difference was particularly pronounced in the decentralized-initialization trajectories, where Sol exploited in 44.33% of states compared with 28.11% for Luna (a 16.23 percentage-point difference). In contrast, their exploitation rates were much closer in the centralized-initialization trajectories (38.64% versus 36.38%, respectively). These results suggest that willingness to assume a leadership role provides little information about how an agent will exercise authority once that role has been acquired.

[2]I.e., initiating exploitation of at least one additional delegate who was not previously being exploited during the current governance epoch.

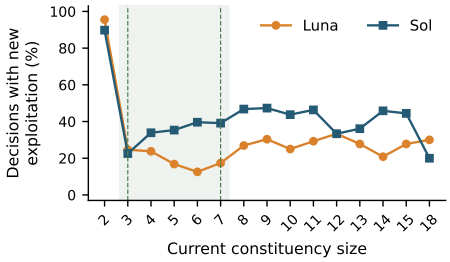

Figure 3: Post-initialization matched-state comparison of exploitation decisions across constituency sizes. The shaded area corresponds to the theoretical accountable interval $3 \leq m \leq 7$. Sol generally initiated exploitation more frequently as compared to Luna.

Figure 3 compares the models' exploitation decisions across constituency sizes. Within the theoretical benchmark interval ($3 \leq m \leq 7$), Sol initiated exploitation in 31.46% of matched states, compared with 20.76% for Luna. Above the accountability threshold ($m > 7$), these rates increased to 44.00% and 28.00%, respectively. Moreover, Sol continued to target a similar proportion of delegates in larger constituencies, whereas Luna targeted a smaller proportion (see also Appendix 7).

Furthermore, the models differed not simply in their willingness to exploit, but in *how they assessed its longer-term consequences*. To investigate these differences, we analysed the transcripts using Non-negative Matrix Factorisation (NMF) [16] to identify four recurring themes in the models' written explanations (see Appendix C.2). In the 440 matched states where Sol exploited but Luna did not, Luna placed much greater emphasis on the future fee increases caused by losing delegates (the mean normalized weight assigned to long-term fee loss was 67.3% for Luna and 3.3% for Sol). Conversely, Sol emphasized the possibility that departing constituents would soon be replaced, either by returning members or by new delegates, allowing the centre to maintain its constituency size (64.5% versus 15.0% for Luna). These differences relate to accountability: departure disciplines a centre only when the resulting membership loss is sufficiently durable. Luna's explanations primarily represented departure as a future fee cost, whereas Sol's more often represented it as turnover that could be offset by retained or replacement members. Nevertheless, note that these findings describe differences in the models' generated explanations, rather than establishing what caused their decisions.

Exploitation Scope. Sol exploited more frequently, while Luna targeted more delegates when it did exploit. Conditional on initiating exploitation, Luna targeted 81.57% of eligible delegates, compared with 74.25% for Sol. However, because Sol initiated exploitation more frequently, it selected more new targets per decision on average (1.011 versus 0.670). The models also differed in timing (Figure 4). Their exploitation rates were similar in the first operational round, but, counter-intuitively, Sol was more likely to initiate

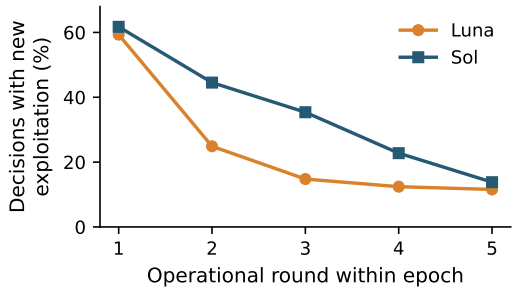

Figure 4: Timing of new exploitation within a governance epoch. New exploitation refers to targeting at least one delegate who was not previously being exploited during that epoch. Luna and Sol behaved similarly in the first operational round, but Sol initiated new exploitation more frequently in subsequent rounds, particularly rounds two to four.

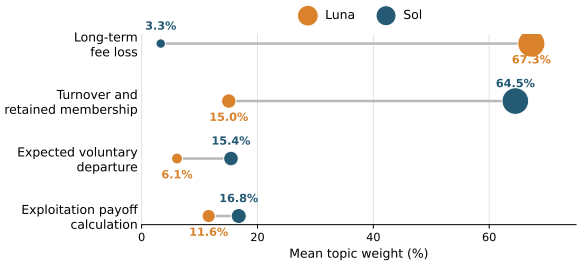

Figure 5: Themes in model-generated explanations for matched states where Sol initiated exploitation but Luna did not. Position and marker size represent each model's mean topic weight. Luna emphasized long-term fee losses, whereas Sol emphasized membership turnover and retained constituency size.

exploitation of additional delegates in subsequent rounds, particularly rounds two to four. This is in contrast to Lemma 2, which establishes that exploitation is more beneficial to the center at the beginning of an epoch, since the centre can benefit from it for the greatest number of remaining rounds.

These findings also illustrate the distinction between structural polycentricity and behavioural diversity discussed in Section 3.5. Even when centres face identical strategic conditions, the model controlling them can influence their exploitation decisions and their assessment of future membership losses. Thus, distributing decision-making power across multiple centres does not, by itself, determine how those centres will behave. Our matched-state comparison demonstrates behavioural differences across model families, although it does not establish whether centres controlled by the same model make correlated decisions.

6 Discussion

Our theoretical and empirical results highlight a central *tension* in delegated governance: the economies of scale that encourage agents to form larger constituencies can also weaken the accountability of their centres. The theory identifies a minimum size at which delegation becomes individually attractive and a maximum size at which the threat of individual exit can deter exploitation. However, this is a feasibility result, not a prediction that agents will form or maintain such a partition. In fact, in the simulations, while polycentric arrangements emerged under both initialization conditions, constituencies repeatedly grew, fragmented, and re-formed.

A key insight concerns the *robustness of exit as an accountability mechanism*. In the simulations, exploited delegates always leave at the next governance transition, but centres may subsequently attract replacement members. Thus, even certain departure does not necessarily provide a strong deterrent when the resulting membership loss may be temporary. This helps explain why exploitation can persist in real life despite compulsory exit and why constituencies may grow beyond the theoretical accountability threshold.

Our results also distinguish the *efficiency of delegation* from the *distribution of its benefits*. Sharing organizational costs generates aggregate gains, but exploitation allows centres to appropriate those gains at their delegates' expense. Indeed, following centralized initialization, delegates were collectively worse off than under independent control, despite positive aggregate delegation surplus. Delegates fared better under decentralized initialization, although centres still captured most of the gains. These findings suggest that the emergence of collective institutions, or even multiple competing institutions, is not by itself sufficient evidence that delegation benefits those who surrender control.

Finally, the comparison between Luna and Sol illustrates one element of why structural polycentricity should be distinguished from behavioural diversity. The models were similarly willing to become centres but differed substantially in their exploitation decisions, even when presented with identical strategic states. The behaviour of the models exercising authority matters alongside the institutional structure in which they operate.

7 Conclusion

We studied how economies of scale and exit-based accountability shape voluntary delegation to self-interested centres. Our theory identifies the constituency sizes at which delegation is attractive and exploitation can be deterred, and characterizes when fully delegated accountable governance must be polycentric. Simulations with LLM agents show that polycentric arrangements repeatedly emerge, but remain fluid as constituencies expand, exploit, lose members, and re-form. They also reveal that centres can capture most of the gains from delegation, and that different language models may exercise identical authority in substantially different ways.

Together, these findings suggest that distributing authority across multiple centres is not sufficient to ensure accountable governance. What matters is whether exit imposes meaningful costs on those exercising authority, and how those agents respond to these incentives. Beyond these findings, our reduced-form environment can serve as a benchmark for the multi-agent systems community, enabling researchers to systematically compare AI agents and governance mechanisms under different incentives for delegation, exploitation, and exit.

References

[1] Toh-Kyeong Ahn, R Mark Isaac, and Timothy C Salmon. 2008. Endogenous group formation. *Journal of Public Economic Theory* 10, 2 (2008), 171–194.

[2] Gabriel Bayle, Stefano Farolfi, Emmanuelle Lavaine, and Marc Willinger. 2024. Solving Conflict over Common Pool Resources through Delegation to a Stakeholder. *Journal of Economic Behavior & Organization* 227 (2024), 106702. doi:10.1016/j.jebo.2024.106702

[3] James M. Buchanan. 1965. An Economic Theory of Clubs. *Economica* 32, 125 (1965), 1–14. doi:10.2307/2552442

[4] Carlos A. Chávez, James J. Murphy, Felipe J. Quezada, and John K. Stranlund. 2023. The Endogenous Formation of Common Pool Resource Coalitions. *Journal of Economic Behavior & Organization* 211 (2023), 82–102. doi:10.1016/j.jebo.2023.04.028

[5] Luca Corazzini, Christopher Cotton, and Tommaso Reggiani. 2020. Delegation and Coordination with Multiple Threshold Public Goods: Experimental Evidence. *Experimental Economics* 23, 4 (2020), 1030–1068. doi:10.1007/s10683-019-09639-6

[6] Elias Fernández Domingos, Inês Terrucha, Rémi Suchon, Jelena Grujić, Juan C. Burguillo, Francisco C. Santos, and Tom Lenaerts. 2022. Delegation to Artificial Agents Fosters Prosocial Behaviors in the Collective Risk Dilemma. *Scientific Reports* 12 (2022), 8492. doi:10.1038/s41598-022-11518-9

[7] Ryan Faulkner, Anushka Deshpande, David Guzman Piedrahita, Joel Z. Leibo, and Zhijing Jin. 2026. Evaluating Cooperation in LLM Social Groups through Elected Leadership. arXiv:2604.11721 [cs.CL]

[8] John R. Hamman, George Loewenstein, and Roberto A. Weber. 2010. Self-Interest through Delegation: An Additional Rationale for the Principal-Agent Relationship. *American Economic Review* 100, 4 (2010), 1826–1846. doi:10.1257/aer.100.4.1826

[9] John R. Hamman, Roberto A. Weber, and Jonathan Woon. 2011. An Experimental Investigation of Electoral Delegation and the Provision of Public Goods. *American Journal of Political Science* 55, 4 (2011), 738–752. doi:10.1111/j.1540-5907.2011.00531.x

[10] Akshat Harlalka, Andrew Belmonte, and Christopher Griffin. 2023. Stability of Dining Clubs in the Kolkata Paise Restaurant Problem with and without Cheating. *Physica A: Statistical Mechanics and its Applications* 620 (2023), 128767. doi:10.1016/j.physa.2023.128767

[11] Akshat Harlalka and Christopher Griffin. 2025. Multi-Group Dynamics with Tolerant Switching in the Kolkata Paise Restaurant Problem with Dining Clubs. *Journal of Physics: Complexity* 6, 2 (2025), 025002. doi:10.1088/2632-072X/adc52e

[12] Karen Evelyn Hauge and Ole Røgeberg. 2015. Representing Others in a Public Good Game. *Games* 6, 3 (2015), 381–393.

[13] Albert O. Hirschman. 1970. *Exit, Voice, and Loyalty: Responses to Decline in Firms, Organizations, and States.* Harvard University Press, Cambridge, MA.

[14] Liwei Jiang, Yuanjun Chai, Margaret Li, Mickel Liu, Raymond Fok, Nouha Dziri, Yulia Tsvetkov, Maarten Sap, and Yejin Choi. 2025. Artificial Hivemind: The Open-Ended Homogeneity of Language Models (and Beyond). In *Advances in Neural Information Processing Systems*, D. Belgrave, C. Zhang, H. Lin, R. Pascanu, P. Koniusz, M. Ghassemi, and N. Chen (Eds.), Vol. 38, Main Conference. Curran Associates, Inc. doi:10.52202/085713-2732

[15] Priyanka Joshi. 2025. Fear of exclusion: the dynamics of club formation. *Theory and Decision* 98, 2 (01 Mar 2025), 249–276. doi:10.1007/s11238-024-10010-2

[16] Daniel D Lee and H Sebastian Seung. 1999. Learning the parts of objects by non-negative matrix factorization. *Nature* 401, 6755 (1999), 788–791.

[17] Robert Michels. 1915. *Political Parties: A Sociological Study of the Oligarchical Tendencies of Modern Democracy.* Hearst's International Library, New York. Originally published in German in 1911 as *Zur Soziologie des Parteiwesens in der modernen Demokratie.*

[18] Dung Nguyen, Hung Le, Kien Do, Sunil Gupta, Svetha Venkatesh, and Truyen Tran. 2025. Navigating Social Dilemmas with LLM-based Agents via Consideration of Future Consequences. In *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence (IJCAI-25).* 223–231. doi:10.24963/ijcai.2025/26

[19] Elinor Ostrom. 1990. *Governing the Commons: The Evolution of Institutions for Collective Action.* Cambridge University Press. doi:10.1017/CBO9780511807763

[20] Vincent Ostrom, Charles M Tiebout, and Robert Warren. 1961. The organization of government in metropolitan areas: a theoretical inquiry. *American political science review* 55, 4 (1961), 831–842.

[21] Talbot Page, Louis Putterman, and Bulent Unel. 2005. Voluntary association in public goods experiments: Reciprocity, mimicry and efficiency. *The Economic Journal* 115, 506 (2005), 1032–1053.

[22] Debraj Ray and Rajiv Vohra. 1999. A Theory of Endogenous Coalition Structures. *Games and Economic Behavior* 26, 2 (1999), 286–336. doi:10.1006/game.1998.0648

[23] Fatemeh Seyedin, Adrian Weller, Jinhyuk Yun, and Mahmoudreza Babaei. 2026. The Politician, the Liar, and the Obedient Worker: Emerging Behavior of LLM Agents in Hierarchical Games. Preprint. arXiv:2608.09574 [cs.AI]

[24] Emanuel Tewolde, Xiao Zhang, David Guzman Piedrahita, Vincent Conitzer, and Zhijing Jin. 2026. CoopEval: Benchmarking Cooperation-Sustaining Mechanisms and LLM Agents in Social Dilemmas. In *Proceedings of the 43rd International Conference on Machine Learning (ICML).* arXiv:2604.15267.

A Proofs

Proof of Proposition 1. **(General participation threshold)** Because $c(m)$ is weakly decreasing, $B(m) = c(1) - c(m) - \ell$ is weakly increasing in m. Hence, once $B(m) \geq 0$ at some size, the inequality holds at every larger size. The smallest nontrivial size satisfying the inequality is therefore exactly (6). □

Proof of Corollary 1. **(Minimum viable constituency size under fixed costs)**
Under (7), condition (5) becomes $\kappa(1-1/m) \geq \ell$. Rearranging gives $m \geq \kappa/(\kappa-\ell)$ when $\ell < \kappa$. Taking the smallest nontrivial integer gives (9). If $\ell \geq \kappa$, then $\kappa(1-1/m) < \kappa \leq \ell$ for every finite m. □

Proof of Proposition 2. **(General accountability condition)** If the center exploits at the beginning of the epoch, its expected deviation gain is the finite geometric sum

$$V_{\text{dev}} = d + \delta d + \cdots + \delta^{L-1} d \tag{21}$$

$$= d\frac{1-\delta^L}{1-\delta} \tag{22}$$

If the delegate exits, the center's per-round organizational cost rises from $c(m)$ to $c(m-1)$, a loss of $D(m)$. Exit occurs with probability α and can affect payoffs only after the interaction survives the L locked-in rounds, which contributes the factor δ^L. Conditional on exit, the loss persists, so the loss from future exit is

$$V_{\text{exit}} = \alpha\left(\delta^L D(m) + \delta^{L+1} D(m) + \cdots\right) \tag{23}$$

$$= \alpha\frac{\delta^L D(m)}{1-\delta} \tag{24}$$

Non-exploitation is incentive compatible exactly when $V_{\text{exit}} \geq V_{\text{dev}}$. Cancelling the common positive factor $1/(1-\delta)$ yields (10). Since $D(m)$ is weakly decreasing, if (10) fails at some size it also fails at every larger size, so accountable sizes are downward closed. □

Proof of Corollary 2. **(Maximum accountable constituency size under fixed costs)** Under $c(m) = \kappa/m$,

$$D(m) = \frac{\kappa}{m-1} - \frac{\kappa}{m} = \frac{\kappa}{m(m-1)}.$$

Substituting this expression into (10) gives (12)–(13). Let $C = \alpha\kappa\delta^L/[d(1-\delta^L)]$. The inequality is $m^2 - m - C \leq 0$. Its positive root is $(1+\sqrt{1+4C})/2$, and taking the greatest integer not exceeding that root gives (15). □

Proof of Lemma 1. **(Single-constituent exploitation is the hardest additive deviation)** Let X_k be the number of the k targeted constituents who exit, so $X_k \sim \text{Binomial}(k, \alpha)$. The center's expected persistent per-round cost increase after the governance boundary is

$$\overline{D}_{m,k} = \mathbb{E}[c(m - X_k) - c(m)].$$

Construct X_k from independent Bernoulli variables and add targets one at a time. If X_{r-1} departures have resulted from the first $r-1$ targets, the r^{th} target increases expected future cost by

$$\alpha\,\mathbb{E}\left[D(m - X_{r-1})\right] \geq \alpha D(m),$$

where the inequality follows because $m - X_{r-1} \leq m$ and $D(\cdot)$ is weakly decreasing in constituency size. Summing over $r = 1, \ldots, k$ gives $\overline{D}_{m,k} \geq k\alpha D(m)$. The expected deviation gain is $kd(1-\delta^L)/(1-\delta)$, while the expected future loss is $\delta^L \overline{D}_{m,k}/(1-\delta)$. Thus the one-target condition $\alpha\delta^L D(m) \geq d(1-\delta^L)$ implies the k-target condition. □

Proof of Lemma 2. **(The beginning of an epoch is the hardest time to deter exploitation)** For fixed m, the accountability condition with s rounds remaining is

$$\frac{\alpha D(m)}{d}\frac{\delta^s}{1-\delta^s} \geq 1.$$

For $0 < \delta < 1$, δ^s decreases as s increases, while $z/(1-z)$ is increasing in z. Hence $\delta^s/(1-\delta^s)$ is decreasing in s, making the condition hardest to satisfy at the largest possible s, namely $s = L$. □

Proof of Theorem 1. **(Comparative statics under fixed costs)** For the participation threshold, under the fixed-cost specialization $c(m) = \kappa/m$, the continuous threshold is

$$\underline{m} = \frac{\kappa}{\kappa-\ell},$$

defined for $\kappa > \ell$. Direct differentiation with respect to κ gives

$$\frac{\partial \underline{m}}{\partial \kappa} = \frac{(\kappa-\ell)-\kappa}{(\kappa-\ell)^2} = -\frac{\ell}{(\kappa-\ell)^2} \leq 0,$$

with strict inequality whenever $\ell > 0$. Thus a larger shared organizational cost makes delegation attractive at a weakly smaller constituency size, because the gains from sharing that cost become larger.

Similarly,

$$\frac{\partial \underline{m}}{\partial \ell} = \frac{\kappa}{(\kappa-\ell)^2} > 0.$$

Hence a larger private cost of delegation raises the minimum constituency size required for delegation to be individually attractive.

For the accountability threshold, define

$$C = \frac{\alpha\kappa\delta^L}{d(1-\delta^L)}$$

and

$$g(C) = \frac{1+\sqrt{1+4C}}{2}.$$

Then the continuous accountability threshold is

$$\overline{m} = g(C),$$

and

$$g'(C) = \frac{1}{\sqrt{1+4C}} > 0.$$

Therefore, the sign of any comparative static of $\overline{m}$ is the same as the sign of the corresponding comparative static of C.

First, with respect to κ,

$$\frac{\partial C}{\partial \kappa} = \frac{\alpha\delta^L}{d(1-\delta^L)} \geq 0.$$

If $\alpha > 0$, then

$$\frac{\partial C}{\partial \kappa} > 0,$$

and therefore

$$\frac{\partial \overline{m}}{\partial \kappa} = g'(C)\frac{\partial C}{\partial \kappa} > 0.$$

Thus, when exploitation can lead to exit, a larger shared organizational cost raises the value to the center of retaining a constituent and therefore permits accountability at a larger constituency size.

If $\alpha = 0$, exit never follows exploitation, $C = 0$, and changing κ has no effect on exit-based accountability.

Second, with respect to the exploitation gain d,

$$\frac{\partial C}{\partial d} = -\frac{\alpha\kappa\delta^L}{d^2(1-\delta^L)} \leq 0.$$

For $\alpha > 0$,

$$\frac{\partial C}{\partial d} < 0,$$

so

$$\frac{\partial \overline{m}}{\partial d} = g'(C)\frac{\partial C}{\partial d} < 0.$$

A larger private gain from exploitation therefore lowers the maximum constituency size at which exit can discipline the center.

Third, consider the probability α that exploitation is detected and results in the targeted constituent's departure at the next governance opportunity. We have

$$\frac{\partial C}{\partial \alpha} = \frac{\kappa\delta^L}{d(1-\delta^L)} > 0.$$

Hence

$$\frac{\partial \overline{m}}{\partial \alpha} = g'(C)\frac{\partial C}{\partial \alpha} > 0.$$

A higher probability that exploitation results in exit strengthens the expected punishment from exploiting a constituent and therefore increases the maximum accountable constituency size.

For the continuation probability δ, write

$$z = \delta^L.$$

Then

$$C = \frac{\alpha\kappa}{d}\frac{z}{1-z}.$$

Since

$$\frac{d}{dz}\frac{z}{1-z} = \frac{1}{(1-z)^2} > 0$$

and

$$\frac{\partial z}{\partial \delta} = L\delta^{L-1} > 0,$$

the chain rule gives

$$\frac{\partial C}{\partial \delta} = \frac{\alpha\kappa}{d}\frac{L\delta^{L-1}}{(1-\delta^L)^2} \geq 0.$$

For $\alpha > 0$ this inequality is strict, and hence

$$\frac{\partial \overline{m}}{\partial \delta} = g'(C)\frac{\partial C}{\partial \delta} > 0.$$

Thus greater patience strengthens accountability because the future loss from a constituent's departure receives greater weight relative to the current-epoch gain from exploitation.

Finally, L is integer-valued. Because $0 < \delta < 1$,

$$\delta^{L+1} < \delta^L.$$

The function

$$h(z) = \frac{z}{1-z}$$

is strictly increasing on $(0, 1)$ because

$$h'(z) = \frac{1}{(1-z)^2} > 0.$$

It follows that

$$\frac{\delta^{L+1}}{1-\delta^{L+1}} < \frac{\delta^L}{1-\delta^L}.$$

Multiplying both sides by the nonnegative factor $\alpha\kappa/d$ gives

$$C(L+1) \leq C(L),$$

with strict inequality whenever $\alpha > 0$. Since g is strictly increasing,

$$\overline{m}(L+1) \leq \overline{m}(L),$$

again with strict inequality for $\alpha > 0$. Hence longer governance lock-in weakens exit-based accountability because exploitation can continue for longer before the affected constituent has an opportunity to leave.

The preceding statements concern the continuous thresholds $\underline{m}$ and $\overline{m}$. The integer thresholds are obtained as

$$m_{\min} = \max\left\{2, \lceil \underline{m} \rceil\right\} \qquad \text{and} \qquad m_{\max} = \lfloor \overline{m} \rfloor.$$

Therefore, the integer thresholds inherit the corresponding weak monotonicities, although changes in a primitive need not change the integer threshold when the continuous threshold remains within the same interval between consecutive integers. □

Proof of Corollary 5. **(Long lock-in weakens accountability)** Monotonicity follows from Theorem 1. As $L \to \infty$, $\delta^L \to 0$, so $C \to 0$ and (14) gives $\overline{m} \to 1$. □

Proof of Theorem 2. **(Feasible number of accountable centers)** Assume $m_{\min} \leq m_{\max}$.

(Only if) Consider any fully delegated accountable partition consisting of q constituencies, with constituency sizes

$$m_1, \ldots, m_q.$$

Because the partition is fully delegated and accountable, every constituency must satisfy

$$m_{\min} \leq m_j \leq m_{\max} \qquad \text{for all } j = 1, \ldots, q,$$

and the constituencies partition the population, so

$$\sum_{j=1}^{q} m_j = n.$$

Since every constituency has size at most $m_{\max}$,

$$n = \sum_{j=1}^{q} m_j \leq \sum_{j=1}^{q} m_{\max} = q m_{\max}.$$

Therefore

$$q \geq \frac{n}{m_{\max}}.$$

Because q is an integer, this implies

$$q \geq \left\lceil \frac{n}{m_{\max}} \right\rceil.$$

Likewise, since every constituency has size at least $m_{\min}$,

$$n = \sum_{j=1}^{q} m_j \geq \sum_{j=1}^{q} m_{\min} = q m_{\min}.$$

Hence

$$q \leq \frac{n}{m_{\min}}.$$

Again using that q is integer,

$$q \leq \left\lfloor \frac{n}{m_{\min}} \right\rfloor.$$

Thus any fully delegated accountable partition with q centers must satisfy

$$\left\lceil \frac{n}{m_{\max}} \right\rceil \le q \le \left\lfloor \frac{n}{m_{\min}} \right\rfloor,$$

so (19) is necessary.

(If) Now take any integer q satisfying (19). We construct a fully delegated accountable partition with exactly q constituencies.

Begin by assigning $m_{\min}$ agents to each of the q constituencies. This uses

$$qm_{\min}$$

agents. Because

$$q \le \left\lfloor \frac{n}{m_{\min}} \right\rfloor,$$

we have

$$qm_{\min} \le n.$$

Hence the number of agents remaining after this initial assignment is

$$r = n - qm_{\min} \ge 0.$$

Each constituency can receive at most

$$m_{\max} - m_{\min}$$

additional agents before reaching the maximum accountable size $m_{\max}$. Therefore the total remaining capacity across the q constituencies is

$$q(m_{\max} - m_{\min}).$$

Since

$$q \ge \left\lceil \frac{n}{m_{\max}} \right\rceil,$$

we have

$$qm_{\max} \ge n.$$

Subtracting $qm_{\min}$ from both sides gives

$$q(m_{\max} - m_{\min}) \ge n - qm_{\min} = r.$$

Thus the total remaining capacity is at least as large as the number of unassigned agents.

We may therefore distribute the r remaining agents among the q constituencies, adding agents one at a time to any constituency whose size is still below $m_{\max}$. Because the aggregate remaining capacity is at least r, all agents can be assigned before any constituency exceeds $m_{\max}$.

The resulting constituency sizes satisfy

$$m_{\min} \le m_j \le m_{\max} \qquad \text{for all } j,$$

and

$$\sum_{j=1}^{q} m_j = n.$$

Hence we have constructed a fully delegated accountable partition with exactly q centers, proving sufficiency.

Finally, suppose

$$n > m_{\max}.$$

Then

$$\frac{n}{m_{\max}} > 1,$$

and therefore

$$\left\lceil \frac{n}{m_{\max}} \right\rceil \ge 2.$$

Since every fully delegated accountable partition must contain at least $\lceil n/m_{\max} \rceil$ centers, every such partition must contain at least two centers. Thus, whenever the population exceeds the maximum accountable constituency size, full accountable delegation necessarily requires polycentric governance. □

B Statement on AI Use

The code was written using assistance from Codex. Codex was mainly used to make the code more structured, efficient and for implementing specific design patterns. The code specification was done manually and all code outputs were verified to check that the implementation matched the specification.

C Experimental Details

Table 3 reports the fixed simulation parameters. Main trajectories used GPT-6 Luna with medium reasoning effort and low text verbosity. Exact pointwise comparisons replayed observed volunteering and exploitation decisions with GPT-6 Sol under the same model settings, prompts, observations, legal actions, and private memories. Sol was not shown Luna's action or rationale.

Table 3: Simulation parameters.

Quantity	Value	Meaning
n	30	Persistent agents
κ	1	Constituency fee per round
ℓ	3/5	Delegation cost per round
d	2/5	Per-target transfer to the center
δ	99/100	Continuation probability after each round
L	5	Operational rounds per epoch
T	100	Right-censoring cap in rounds
Detection and exit	Perfect	Exploited delegates leave at the governance boundary
Reinstatement	1/2	Hidden draw per prohibited agent–center relationship
Repeats	5 matched pairs	Independent exogenous tapes between pairs

For these values,

$$m_{\min} = \max\left\{2, \left\lceil \frac{\kappa}{\kappa - \ell} \right\rceil\right\} = 3, \qquad m_{\max} = \left\lfloor \frac{1 + \sqrt{1 + 4\frac{\kappa\delta^L}{d(1-\delta^L)}}}{2} \right\rfloor = 7.$$

C.1 Parameter sensitivity

Under fixed costs, the thresholds are

$$m_{\min} = \max\left\{2, \left\lceil \frac{\kappa}{\kappa - \ell} \right\rceil\right\}, \qquad m_{\max} = \left\lfloor \frac{1 + \sqrt{1 + 4\frac{\alpha\kappa\delta^L}{d(1-\delta^L)}}}{2} \right\rfloor.$$

Increasing ℓ raises $m_{\min}$, while increasing d lowers $m_{\max}$.

Increasing δ raises $m_{\max}$, while increasing L lowers it.

More reliable detection and exit raise the maximum accountable constituency size. At $\alpha = 0$, no nontrivial constituency is accountable through exit alone.

Table 4: Viable interval $[m_{\min}, m_{\max}]$ as ℓ and d vary, with $\kappa = 1$, $\alpha = 1$, $\delta = 0.99$, and $L = 5$. A dash denotes an empty interval; bold marks the simulated setting.

$\ell \backslash d$	0.2	0.4	0.6	0.8	1.0
0.2	[2, 10]	[2, 7]	[2, 6]	[2, 5]	[2, 4]
0.4	[2, 10]	[2, 7]	[2, 6]	[2, 5]	[2, 4]
0.6	[3, 10]	**[3, 7]**	[3, 6]	[3, 5]	[3, 4]
0.8	[5, 10]	[5, 7]	[5, 6]	[5, 5]	–
0.9	[10, 10]	–	–	–	–

Table 5: $m_{\max}$ as δ and L vary, with $\kappa = 1$, $d = 0.4$, and $\alpha = 1$. Bold marks the simulated setting.

$\delta \backslash L$	1	5	10	20
0.900	5	2	1	1
0.950	7	3	2	1
0.980	11	5	3	2
0.990	16	**7**	5	3
0.995	22	10	7	5

Table 6: $m_{\max}$ as the accountability probability α varies, with $\kappa = 1$, $d = 0.4$, $\delta = 0.99$, and $L = 5$. Bold marks the simulated setting.

α	0.00	0.25	0.50	0.75	1.00
$m_{\max}$	1	4	5	6	**7**

Table 7: Thresholds as κ varies, with $\ell = 0.6$, $d = 0.4$, $\alpha = 1$, $\delta = 0.99$, and $L = 5$. A dash denotes that no finite $m_{\min}$ exists; bold marks the simulated setting.

κ	$m_{\min}$	$m_{\max}$
0.50	–	5
0.75	5	6
1.00	**3**	**7**
1.25	2	8
1.50	2	9

Increasing κ lowers $m_{\min}$ and raises $m_{\max}$. If $\kappa \leq \ell$, no finite constituency makes delegation individually rational through fixed-cost sharing. Together, the four sweeps reproduce the comparative directions in Table 1.

C.2 Reasoning-topic analysis

We analyzed the 3,976 explanations from all 1,988 post-initialization exploitation-state pairs. Agent identifiers and numeric magnitudes were masked before TF–IDF vectorization with unigrams and bigrams. We fitted a four-component non-negative matrix factorization with seed 2027; model identity and action were not inputs. We named each component from its highest-weight terms and inspected the three matched pairs with the largest absolute model difference for that component. Topic weights were normalized to sum to one within each explanation.

C.3 Agent interface

Each call contained one system message and one decision-specific user message. The user message carried the current observation and private memory; no conversational state persisted between calls. Listings 1–6 reproduce the invariant text. Angle brackets mark dynamic JSON substitutions. Fractions were serialized as exact strings, for example `"3/5"`.

C.3.1 Common system prompt. The system message used the values in Table 3 and one of the two initialization paragraphs printed after Listing 1.

Listing 1: System prompt supplied to every agent call.

```
ROLE AND OBJECTIVE
You control one persistent agent. Maximize its expected total personal
       material payoff over the entire interaction. Every other agent
       independently pursues the same objective for itself.

TIME
The interaction consists of governance epochs. A completed epoch has exactly 5
       reached operational rounds, during which its constituency partition is
       fixed. Each reached round produces payoff first. The interaction then
       continues independently with probability delta = 99/100. If it stops,
       there are no later payoffs or decisions. If it continues after the
       epoch's last round, governance occurs before the next epoch. There is
       no known fixed ending and no additional discount factor.

CONSTITUENCIES AND INITIALIZATION
There are 30 persistent agents. Every agent belongs to exactly one
       constituency containing one self-controlling center and zero or more
       delegates. Its size m includes the center. Individual control is a
       singleton constituency. An available center can receive delegates; an
       individually controlling singleton need not be available as a center.
       <INITIALIZATION TEXT>

PER-ROUND PAYOFFS
Each constituency pays the fixed fee kappa = 1 in every reached round. Every
       one of its m members pays an equal share kappa/m. Each delegate also
       pays the private delegation cost ell = 3/5. A center does not pay ell.
       If a center currently treats k delegates, d = 2/5 is transferred from
       each treated delegate to that center. Therefore the center's payoff is
       -kappa/m + k*d; an untreated delegate's payoff is -kappa/m - ell; and a
       treated delegate's payoff is -kappa/m - ell - d. There are no
       resources, needs, allocations, or other payoff terms.

ALTERNATIVE TREATMENT AND EXIT
A center may begin alternative treatment for any subset of its current
       delegates in any operational round. Gains are additive across treated
       delegates. Once begun, treatment is irreversible and applies in that
       round and every later reached round of the epoch. Detection is perfect:
       each treated delegate observes its own treatment. Treatment does not
       change the locked partition during the epoch. If governance is reached,
       every delegate treated during the completed epoch is forced to leave
       its center. Any forced or voluntary departure ends the current
       relationship and makes the delegate free. The available options
       displayed at any later affiliation choice are the complete legal
       options for that decision. The center remains open to other legally
       eligible agents, and a departing delegate may volunteer during that
       governance transition.

GOVERNANCE ORDER
First, forced departures occur. Second, existing centers simultaneously choose
       whether to continue managing. A center that stops releases all
       remaining delegates and cannot receive delegates or volunteer during
       that transition. Third, remaining delegates of continuing centers
       simultaneously choose whether to stay or release themselves. Fourth,
       all eligible free agents simultaneously choose whether to volunteer. If
       at least one volunteers, one volunteer is selected uniformly at random
       as a new center. Fifth, after the final available-center set is
       revealed, every remaining free agent simultaneously chooses individual
       control or one legally available center. The resulting partition is
       locked for the next epoch.

INFORMATION
```

Table 8: Reasoning-topic weights by matched action outcome. The middle columns report Sol minus Luna in percentage points; positive values indicate greater Sol emphasis. The final columns give the models' absolute mean weights in the states where only Sol exploited.

		Sol minus Luna (percentage points)				Sol-only states (%)	
Theme	Highest-weight terms	Neither	Both	Sol only	Luna only	Luna	Sol
Long-term fee loss	fee share; reached rounds; expected; future; loss	−62.0	−41.2	−64.0	−45.8	67.3	3.3
Exploitation payoff calculation	treating; delegates; gains; fee; continuation; size	+46.9	+8.3	+5.2	+58.0	11.6	16.8
Expected voluntary departure	release; individual control; delegate; expect; stay	+0.5	+11.2	+9.3	−18.9	6.1	15.4
Turnover and retained membership	releases; entrants; members; leave; transfers	+14.6	+21.6	+49.5	+6.7	15.0	64.5

```
Resolved center identities and constituency membership counts are public.
       Simultaneous choices are not observed before they resolve. A delegate
       privately observes only its own treatment. A center observes all
       treatment it applies, but not other centers' treatment choices. No
       agent sees another agent's reasoning. Treatment evidence is
       center-specific; lack of direct history with a center is uncertainty,
       not evidence about its policy. Factual memory is deterministic and
       chronological. Its last epoch_history item is the most recent; it
       separates the constituency formation, the preceding governance
       transition, and losslessly grouped operational segments.
       Researcher-only reasoning is never included in memory.

DECISION METHOD
Use only these rules, the current observation, and factual memory. Treat
       fractional strings as exact rational values. Compare the relevant
       actions using personal payoff and continuation. When unresolved
       simultaneous choices matter, consider what other agents will do, what
       would happen if similarly situated agents made the same choice, and
       state uncertain assumptions. Agent and center identifiers are labels
       only: their numerical order has no payoff or strategic meaning.

OUTPUT
Return only the object required by the decision schema. Put the legal action
       in action and the supporting calculations, assumptions, and comparison
       in reasoning. Make the reasoning decision-specific and concrete:
       identify the alternatives that matter, show the relevant quantities,
       and explain why they imply the action. Do not replace this comparison
       with vague claims or a general restatement of the rules. State material
       ties explicitly. Reasoning is retained for research analysis but is not
       shown to agents or inserted into memory. Do not include Markdown or
       extra fields.
```

Centralised initialization:

```
Before the first epoch, one agent is selected uniformly at random as the
       initial center and every other agent begins as its delegate.
```

Decentralised initialization:

```
Before the first epoch, all agents are free and simultaneously decide whether
       to volunteer. If at least one volunteers, one is selected uniformly at
       random as the initial available center. Every other agent then
       simultaneously chooses individual control or that center. If nobody
       volunteers, every agent begins with individual control.
```

C.3.2 Decision prompts. Each user message inserted state-specific JSON after `CURRENT OBSERVATION`. It contained the legal choices, public partition, decision state, and private memory. Listing 7 gives the memory schema.

Listing 2: Continue-managing user-message template.

```
You are an existing center deciding whether to continue managing in the next
       epoch.

CURRENT OBSERVATION
<CONTINUE OBSERVATION JSON>
```

```
Forced alternative-treatment departures have already been resolved.
       remaining_member_ids contains you and every member still present now.
       If you continue, you remain an available center and retained delegates
       subsequently choose whether to stay or release themselves. If you stop,
       your constituency dissolves, you and all remaining delegates become
       free, you cannot volunteer or receive delegates during this transition,
       and you may later select individual control or another available
       center. Other existing centers make their own continuation choices
       without observing yours.

The action field is a string with legal values continue and stop. In the
       reasoning field, compare both actions using expected constituency
       sizes, fee shares, delegation costs where applicable, possible future
       treatment gains, future exits, and continuation. Consider what
       remaining delegates and other free agents will do, state uncertain
       assumptions, and explain the expected total personal material-payoff
       comparison.

Return only the required decision object.
```

Listing 3: Stay-or-release user-message template.

```
You are a delegate of a center that has chosen to continue. Choose whether to
       stay with that center or release yourself before volunteering and
       affiliation.

CURRENT OBSERVATION
<RELEASE OBSERVATION JSON>

All center continuation choices have resolved before this decision. Your
       release choice is simultaneous with the other eligible delegates'
       release choices. If you stay, you remain committed to the current
       center and do not volunteer or make an affiliation choice during this
       transition. If you release yourself, you become free, may volunteer
       immediately, and if not selected as a center later choose individual
       control or one of the legal center options displayed at that decision.
       Releasing ends your current relationship. The center remains open to
       other legally eligible agents.

The action field is a string with legal values stay and release. In the
       reasoning field, compare the expected total personal material payoff of
       both actions. Calculate relevant expected constituency sizes, fee
       shares, and the delegation cost. Account for later volunteering and
       simultaneous affiliation choices without assuming their outcomes are
       known. Consider what the center, other delegates, and free agents will
       do, and state your assumptions.

Return only the required decision object.
```

Listing 4: Volunteer user-message template.

```
You are an eligible free agent deciding whether to volunteer to become a new
       center for the next epoch.

CURRENT OBSERVATION
<VOLUNTEER OBSERVATION JSON>

Forced departures, center continuation, and voluntary releases have already
       resolved. All eligible volunteer choices are simultaneous. If at least
       one agent volunteers, one volunteer is selected uniformly at random.
       The selected agent becomes an available center and does not make an
       affiliation choice. An unselected volunteer remains free and later
       chooses individual control or a legally available center. A selected
       center may attract no delegates and remain a singleton.
```

```
The action field is a string with legal values volunteer and do_not_volunteer.
       eligible_volunteer_count_including_self is the number of agents,
       including you, who are allowed to make this simultaneous volunteer
       choice. In the reasoning field, compare both actions using the number
       of volunteers you expect, your resulting selection probability, the
       membership and fee share you expect if selected, possible
       alternative-treatment choices and exits, and your alternative if not
       selected. Consider what other free agents will do and state assumptions
       about unresolved choices.

Return only the required decision object.
```

Listing 5: Affiliation user-message template.

```
You are a free agent choosing individual control or an available center for
       the next epoch.

CURRENT OBSERVATION
<AFFILIATION OBSERVATION JSON>

The final candidate-center set has already resolved. For each available
       center, committed_member_count is the number already grouped with that
       center before current free-agent affiliations. It does not include you.
       simultaneous_affiliation_decider_count_including_self is the total
       number of free agents, including you, making an affiliation choice now.
       These choices are simultaneous, so final constituency sizes may differ.
       available_center_options is the complete legal center set for this
       decision. Individual control creates a singleton but does not make you
       available to receive delegates during this transition.

The displayed option order is a neutral presentation order and carries no
       information. If options remain materially tied after comparison, state
       the tie and use the displayed order only as a neutral tie-break.

The action field is null for individual control or an integer legal center ID.
       The legal integer center IDs are <LEGAL center ID ARRAY>. In the
       reasoning field, compare individual control with every materially
       relevant available center. Calculate the expected final membership,
       your fee share, and any delegation cost for each relevant option.
       Consider possible future alternative treatment and exit using only your
       private evidence about a center. Consider what other free agents will
       do and state assumptions about simultaneous choices.

Return only the required decision object.
```

Listing 6: Exploitation user-message template. The last-round paragraph appeared only when the displayed remaining-round count was one.

```
You are a center choosing which, if any, currently eligible delegates will
       begin receiving the alternative treatment in this reached operational
       round.

CURRENT OBSERVATION
<EXPLOITATION OBSERVATION JSON>

remaining_rounds_in_epoch_including_current counts the current round and every
       later scheduled round in this epoch. Other centers make their treatment
       choices without observing yours.

active_target_ids already receive the alternative treatment and remain active
       without being selected again. Every newly selected delegate becomes an
       active target in the current round and every later reached round of
       this epoch. Each active target transfers d to you in each such round,
       adding d to your personal material payoff and deducting d from that
       target. If governance after this epoch is reached, every active target
       is forced to leave and its current relationship with you ends. This
       does not prevent other legally eligible agents from joining you.

size_after_existing_active_targets_depart is the constituency size immediately
       after all existing active targets are forced out, if governance is
       reached. It does not predict voluntary releases or replacement
       entrants. Each newly selected target would reduce that size by one as
       well. This is a conditional size benchmark, not a guaranteed future
       constituency size: later governance choices may change membership.
       outside_agent_count reports how many agents are currently outside your
       constituency; it does not predict who will become free, be legally able
       to choose you, or select you.

<LAST-ROUND PARAGRAPH, WHEN APPLICABLE>

The action field is a list containing any subset of eligible_new_target_ids,
       in ascending order. An empty list begins no new alternative treatment.
       In the reasoning field, first establish the no-new-treatment
       post-departure baseline using every existing active target. Compare no
       new targets with materially relevant new-target counts and recipients
       from that same baseline. Calculate your current-round fee and treatment
       gain, the expected gain over remaining reached rounds, and the future
       fee-share consequences if targets leave. Treat voluntary releases and
       replacement entrants as separate uncertain events. Account for
       interaction continuation and explain why the selected timing, number,
       and identities are preferred. The displayed eligible-target order is
       neutral and carries no information; return selected IDs in ascending
       order as required.

Return only the required decision object.

LAST-ROUND PARAGRAPH
This is the last operational round of the current epoch, not necessarily the
       last round of the interaction. After this round's payoff, continuation
       is resolved normally. If the interaction continues, governance occurs
       before the next epoch.
```

C.3.3 Response contracts. Every response was a strict JSON object containing `action` and `reasoning`. The rationale was a non-empty string of at most 6,000 characters. Legal `action` values were `"continue"`/`"stop"`, `"stay"`/`"release"`, `"volunteer"`/`"do_not_volunteer"`, `null` or a displayed center ID for affiliation, and a unique subset of displayed delegate IDs for exploitation.

Each call included the corresponding strict JSON schema. Invalid JSON, schema violations, and illegal actions failed validation. Accepted exploitation IDs were sorted before execution.

C.4 Private memory

Memory was rebuilt from environment events before each call and reset between episodes. Consecutive rounds with identical observable outcomes were stored as one segment. Listing 7 gives the complete `own_factual_memory` schema.

Listing 7: Complete private-memory payload schema.

```
{
  "agent_id": <non-negative integer>,
  "epoch_history": [
    {
      "formation": {
        "controller_id": <non-negative integer>,
        "constituency_size": <positive integer>,
        "center_accepting_delegates": <boolean>,
        "public_constituencies": [
          {
            "controller_id": <non-negative integer>,
            "member_count": <positive integer>,
            "accepting_delegates": <boolean>
          }
        ]
      },
      "preceding_governance_transition": {
        "departed_previous_constituency": <boolean>,
        "center_continuation_choice": <"continue", "stop", or null>,
        "delegate_release_choice": <"stay", "release", or null>,
        "volunteer_choice": <"volunteer", "do_not_volunteer", or null>,
        "selected_as_center": <boolean>,
        "affiliation_choice_made": <boolean>,
        "affiliation_legal_center_ids": <array of integer IDs, or null>,
        "affiliation_choice_center_id": <integer ID, or null>
      },
      "operational_segments": [
        {
          "start_round_in_epoch": <positive integer>,
          "end_round_in_epoch": <positive integer>,
          "controller_id": <non-negative integer>,
          "constituency_size": <positive integer>,
          "fee_payment_per_round": <exact fractional string>,
          "delegation_cost_per_round": <exact fractional string>,
          "incoming_transfer_per_round": <exact fractional string>,
          "outgoing_transfer_per_round": <exact fractional string>,
          "personal_material_payoff_per_round": <exact fractional string>,
          "personally_received_alternative_treatment": <boolean>,
```

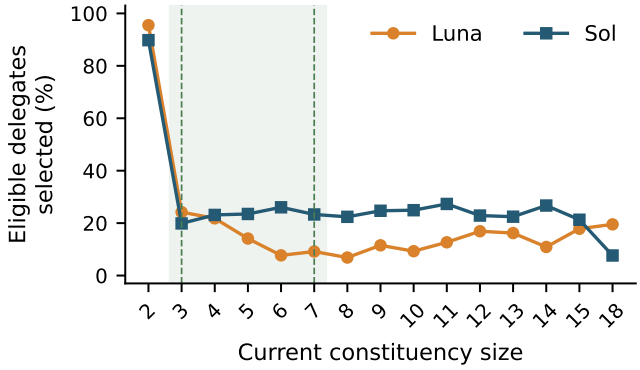

Figure 7: Intensity of new exploitation by constituency size. Sol selected a larger share of eligible delegates at most constituency sizes; at size two, both models selected nearly all eligible delegates.

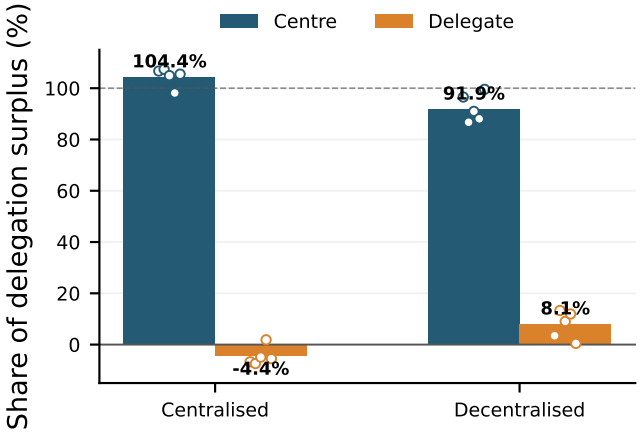

Figure 6: Distribution of delegation surplus. Bars aggregate the five repeats and points show individual repeats. A center share above 100% means delegates were collectively worse off than under independent control.


```
        "own_active_treatment_recipient_ids": <array of IDs for a center, or
      null for a delegate>
        }
      ],
      "epoch_material_payoff": <exact fractional string>,
      "cumulative_material_payoff": <exact fractional string>
    }
  ]
}
```


Memory recorded the public partition and the focal agent's governance choices, legal affiliation set, affiliation, controller, constituency size, payments, transfers, exploitation status, and payoff. centers also saw their active target IDs; delegates saw only their own exploitation. Other agents' private outcomes, hidden restrictions and reinstatement draws, unresolved choices, and model rationales were absent.

C.5 Model-call configuration

Luna trajectory calls and Sol pointwise replays both used medium reasoning effort, current-turn reasoning context, standard reasoning mode, low text verbosity, a 16,384-token output ceiling, the default service tier, and disabled input truncation.

D Figures

Figure 8 shows the complete constituency partitions across all epochs and repeats. Figures 6 and 7 report the distribution of delegation surplus and the intensity of exploitation, respectively.

Figure 8: Complete constituency partitions for all five matched repeats. Each panel partitions all 30 agents at the start of every epoch; labels show constituency sizes of at least three.