arXiv:2610.22293v1 [cs.LG] 15 Oct 2026

Towards Continual Learning: Mitigating Forgetting With Subspace Tuning

Arnav Chavan [1] **Kusha Sareen** [2,3] **Rishabh Tiwari** [1] **Abhranil Chandra** [5] **Johan Obando-Ceron** [2,4]
Kurt Keutzer [1] **Scott Niekum** [5] **Inderjit S Dhillon** [6] **Devvrit Khatri** [7,8]

[1]UC Berkeley [2]Mila [3]McGill University [4]Université de Montréal [5]UMass Amherst
[6]Google [7]UT Austin [8]Mirendil

Large language models (LLMs) acquire new skills through fine-tuning, but it can destroy capabilities they already have: catastrophic forgetting remains a central obstacle to continual learning in LLMs. Recent works connect forgetting to how far, and along which directions, fine-tuning moves the weights. On-policy methods such as reinforcement learning and on-policy distillation avoid updating the principal singular directions of weight matrices and thus forget less. In contrast, SFT concentrates updates along these directions and is susceptible to catastrophic forgetting. In this work, we introduce Subspace Tuning (SST), which restricts each weight update to the bottom singular subspace, eliminating changes along the principal directions. SST is optimizer-agnostic and composes well with fine-tuning recipes including SFT and self-distillation. SST combines tight subspace protection, dense gradient estimation, and a simple refresh mechanism, enabling stronger learning than prior orthogonal fine-tuning methods. With these mechanisms, protecting as much as 99% of the weight subspace substantially improves retention while preserving strong target-task performance. This behavior strengthens with model and data size and nearly eliminates forgetting in the large-model regime. SST also keeps the model close to its base distribution, inducing 2–4× less KL drift than the baseline recipes. Overall, SST advances the learning–forgetting Pareto frontier, making it a practical default for fine-tuning pre-trained language models.

Correspondence: arnavchavan@berkeley.edu, devvrit.03@gmail.com

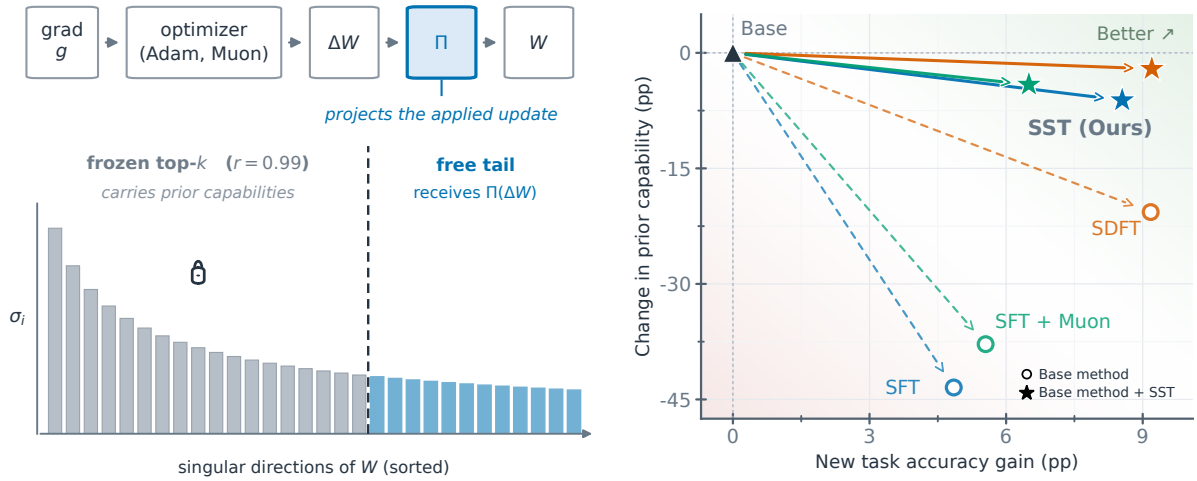

Figure 1: **SST freezes the top of each weight matrix's spectrum, and improves learning–forgetting trade-off across fine-tuning recipes. Left:** the method. SST projects the update on low energy singular directions of each weight matrix. **Right:** the result on Qwen3-8B, averaged over four domains (Table 4). Both axes show changes from the base model in percentage points (pp); higher is better. Circles denote training without SST, stars denote training with SST, and the black triangle marks the base model.

1 Introduction

Large language models (LLMs) are useful across a rapidly growing range of applications, including assistants, coding, search, science, and finance [1, 8, 20]. Many of these applications follow the same paradigm. A small number of strong base models are pre-trained once [7, 49, 56], and each use case then fine-tunes one of them on specialized data [6, 15, 16, 36, 46]. Fine-tuning has thus become a standard way to specialize an LLM for a target task or domain.

This paradigm comes at the cost of catastrophic forgetting. As a model learns a new task, capabilities acquired over trillions of training tokens, including instruction following, coding, reasoning, and general knowledge, can degrade severely [24, 28, 29]. This phenomenon is particularly difficult to address because the original pre-training data is typically unavailable [23], while post-hoc repair may recover only part of what was lost [28, 44] (Sec. 3). More broadly, forgetting stands in the way of continual learning. The next generation of foundation models should be able to improve over their deployed lifetime, adding new skills, domains, and corrections without sacrificing what they already do well.

Recent works provide insights into the origin of this forgetting. The leading singular directions of pretrained weight matrices have been associated with important model functionality [30, 43], while reinforcement learning has been found to fine-tune relatively small parameter subnetworks [33]. More directly, RLVR updates concentrate outside the principal directions and induce relatively little spectral drift compared with SFT [60]. On-policy RL also stays closer to the base policy and forgets less than SFT [45]. These observations suggest a simple hypothesis. Preserving the principal singular subspaces of pretrained weights may preserve prior capabilities while leaving sufficient capacity elsewhere in weight space to learn new ones. Existing methods induce this behavior indirectly. For example, SDFT [44] trains the model on its own rewrites of the training data, so that supervision stays close to the model's output distribution—approximately on-policy, and hence expected to forget less. This helps on non-reasoning instruct models, but we find substantial forgetting remains on longer reasoning traces (Sec. 3).

We introduce *Subspace Tuning* (SST), which restricts weight updates to a non-principal subspace of the weight spectrum. Concretely, let ΔW be the optimizer's final update to a weight matrix W, and let U_k and V_k contain the top-k left and right singular vectors of W. SST applies the projected update

$$\Pi(\Delta W) \;=\; (I - U_k U_k^\top)\, \Delta W\, (I - V_k V_k^\top) \tag{1}$$

in place of ΔW. The two projectors remove from the update any component along the principal directions, on the output and input sides respectively, so every applied step stays in the non-principal subspace (Sec. 2.1). Importantly, SST constrains the final optimizer update rather than the raw gradient. Stateful optimizers such as Adam rescale the gradient coordinate-wise before it reaches the weights, so a projected gradient does not produce a projected step and part of the update returns to the directions we mean to freeze (Sec. 2). Because the projection acts on the applied update, SST is optimizer-agnostic and composes with any fine-tuning recipe, with no other change to training.

We conduct extensive empirical analysis and establish the following properties of SST (Sec. 3):

- **SST learns new tasks and retains prior capabilities.** Despite protecting the top 99% of singular directions, SST matches the target-task accuracy of unconstrained fine-tuning, while retaining the prior capabilities that standard fine-tuning severely degrade (Figure 1). This behavior holds across Science, Medicine, Pathology, Finance, and Tool Use, across SFT and SDFT, and with both Adam and Muon; Tables 4 and 1 provide the full per-domain results. We hypothesize that such constrained learning is possible because LLMs are heavily over-parameterized [2, 22], leaving enough capacity in the non-principal subspace to fit a new task. When a fixed complement becomes restrictive, a refresh step recomputes the SVD on the current weights, realigning the protected subspace so adaptation can continue (Section 2.3).
- **SST scales well.** With more training data, SFT forgets progressively more while SST stays flat and still reaches higher target-task accuracy (Figure 3). With larger models, forgetting under SST shrinks even further, staying within $\sim 1\%$ of the base model from 8B onward (Figure 2).

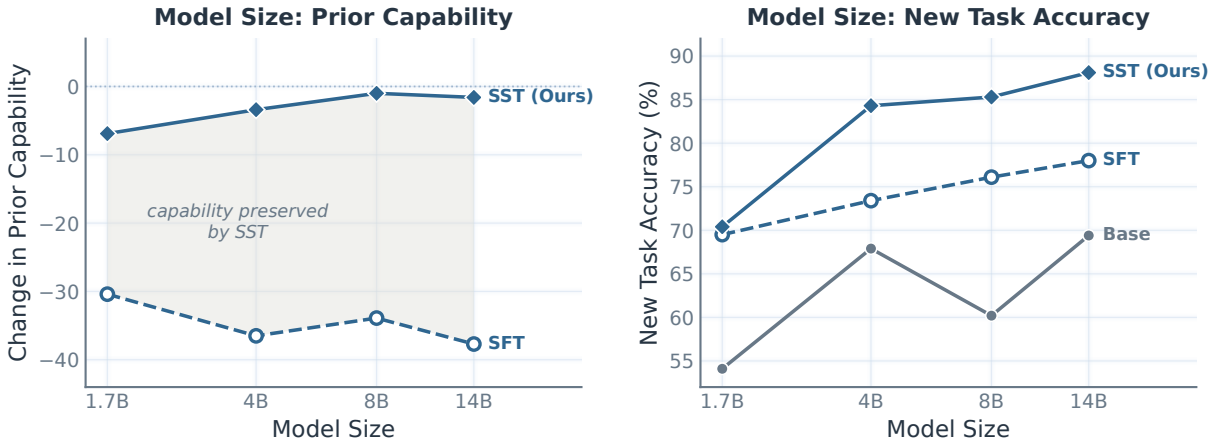

Figure 2: **Scaling model size.** Qwen3 models from 1.7B to 14B fine-tuned on Science for five epochs. SFT forgets heavily at every scale, whereas forgetting under SST shrinks as models grow and nearly vanishes from 8B onward. SST also reaches higher new-task accuracy at every size.

- **SST stays close to the base model.** It induces 2-4× less KL divergence from the pretrained model than unconstrained recipes and preserves behavior that accuracy alone does not capture. A thinking model fine-tuned with SST keeps its natural reasoning length where SFT collapses it, and adding SST to SDFT reduces fabricated reference-solution attributions (Figure 8).
- **SST makes sequential fine-tuning cumulative.** Over two- and three-stage fine-tuning chains of the above domains, each skill rises when it is trained and holds its level as later stages are learned, whereas under SFT each new stage degrades the dormant ones (Figure 9,10).

Overall, this work shows that catastrophic forgetting is not an inevitable consequence of learning, but depends on where the update lands in weight space, and provides a simple, practical method for improving capability retention while acquiring new capabilities..

2 Subspace Tuning

Across many complex systems, a small number of leading modes of a linear operator capture much of its dominant structure, while the remaining spectrum describes lower-energy variation [4, 42]. Deep neural network weight matrices exhibit similarly concentrated and heavy-tailed spectra [30], and interventions on their leading singular components can disproportionately affect model performance [43]. These observations motivate preserving the principal singular directions of pretrained weights during adaptation.

We first formalize this constraint (Sec. 2.1), show that it must be enforced on the applied optimizer update rather than the gradient (Sec. 2.2), and introduce a basis refresh that trades stricter protection for continued adaptability (Sec. 2.3).

2.1 Projection Formulation

Consider a single weight matrix $W \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$, which maps an input activation x to an output Wx. Its SVD, $W = \sum_i \sigma_i u_i v_i^\top$, reads off the component of x along the input direction v_i (via $v_i^\top x$) and writes it, scaled by σ_i, onto the output direction u_i. So the right singular vectors v_i are the *input* directions the layer responds to, and the left singular vectors u_i are the *output* directions it produces. The top-k of these, $V_k = [v_1, \ldots, v_k]$ and $U_k = [u_1, \ldots, u_k]$, are what we intend to protect. Let $P_U = U_k U_k^\top$ and $P_V = V_k V_k^\top$ be the orthogonal projectors onto the protected output and input subspaces, and $Q_U = I - P_U$, $Q_V = I - P_V$

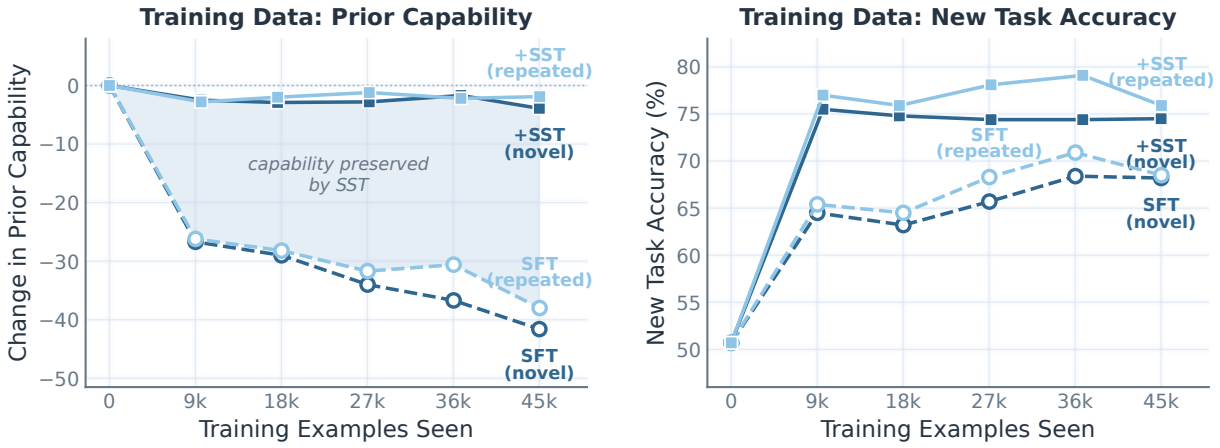

Figure 3: **Scaling training data.** Qwen3-8B fine-tuned on Science under a matched 45k-example budget: the *novel* arms see a fresh stream once, the *repeated* arms cycle 9k examples for five epochs. Under both exposures SFT drifts progressively farther from the base model, while SST stays flat and learns the task better.

be their complements (the free subspaces). Any output vector splits as protected-plus-free, $I = P_U + Q_U$, and likewise any input, $I = P_V + Q_V$. Inserting these identities on the left and right of ΔW decomposes *any* update into four blocks, one per (input-side, output-side) combination:

$$\Delta W = \underbrace{P_U \Delta W P_V}_{\text{(1) protected}\to\text{protected}} + \underbrace{P_U \Delta W Q_V}_{\text{(2) free}\to\text{protected}} + \underbrace{Q_U \Delta W P_V}_{\text{(3) protected}\to\text{free}} + \underbrace{Q_U \Delta W Q_V}_{\text{(4) free}\to\text{free}}. \tag{2}$$

Freezing the top directions therefore means keeping only block (4) and discarding (1)-(3) from the update ΔW. Keeping only block (4) is exactly the two-sided projection of Eq. 1, which we restate:

$$\Pi(\Delta W) \;=\; Q_U\, \Delta W\, Q_V \;=\; (I - U_k U_k^\top)\, \Delta W\, (I - V_k V_k^\top). \tag{3}$$

In particular, the updated matrix $W_{\text{ref}} + \Pi(\Delta W)$ acts identically to W_{ref} on every protected input direction and produces no new component along any protected output direction: the protected singular triplets (σ_i, u_i, v_i), $i \le k$, are preserved exactly. Any of blocks (1)-(3) would break this, which is why the projection must be two-sided. We fix $k = \lfloor r \cdot \min(d_{\text{out}}, d_{\text{in}}) \rfloor$ with a rank fraction r close to 1 so only the low-energy tail is free, and compute U_k, V_k from W_{ref}, the weights at the start of training (the anchor that refresh later updates, Sec. 2.3).

2.2 Projecting the Optimizer Update

SST applies Π to the optimizer's final update ΔW. The projection could instead act on the gradient before the optimizer, either alone or together with the applied update. These choices are not equivalent. For plain SGD, whose step is a scalar multiple of the gradient, projecting the gradient and projecting the update coincide. Under Adam and other stateful optimizers they differ. Adam rescales each coordinate by its running moments, and this diagonal preconditioning does not commute with Π, so a gradient confined to the free subspace can still produce an applied update with a component in the protected subspace.

We compare three possible enforcement points (full equations in Appendix B.8):

(1) **Project the gradient**: apply Π to the raw gradient before Adam, $g \leftarrow \Pi(g)$.

(2) **Project the applied update**: let Adam update its moments from the full gradient, then apply Π to the step that Adam produces (this is what SST does).

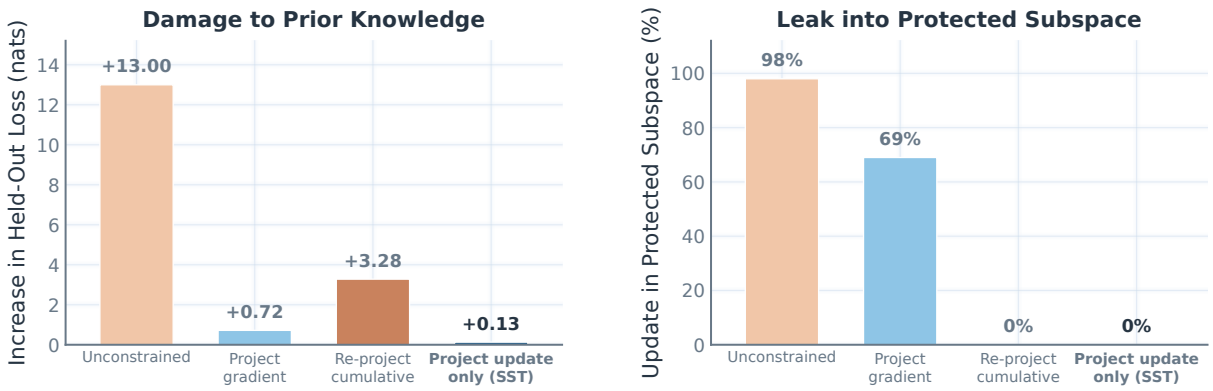

Figure 4: **Controlled study on Qwen2.5-0.5B.** Left: Increase in held-out NLL after fitting the toy task (lower is better). Right: Share of the weight change entering protected directions (lower is better). Only update projection achieves exact confinement while keeping the held-out NLL nearly unchanged.

(3) **Re-project the cumulative change**: project the gradient g and also project the applied update ΔW. This is equivalent to projecting the gradient and, after each step, re-projecting the total change from the initial anchor, $W - W_{\text{ref}}$.

To isolate this choice, we fine-tune Qwen2.5-0.5B for 60 Adam steps on a fixed eight-example next-token prediction task. We use $r{=}0.99$ and apply the constraint to the attention and MLP linear maps. Because the task is deliberately easy, every method reaches a final training loss of approximately 10^{-5}. We can therefore compare their collateral effects using two measurements: the increase in NLL on held-out generic sentences, and the fraction of the total weight change that falls outside the allowed free-to-free block.

Our results are in Figure 4. Unconstrained fine-tuning increases the held-out NLL by 13.0, with 98% of the weight change entering protected directions. Projecting the gradient reduces the NLL increase to 0.72, but Adam reintroduces substantial leakage: 69% of the final change still leaks into the protected directions. Re-projecting the cumulative change eliminates leakage, but raises the held-out NLL by 3.28. Projecting the applied update is the only scheme that both eliminates leakage and keeps the probe nearly unchanged, with an NLL increase of just 0.13. We therefore design SST to project the final applied update. The full procedure is given in Algorithm 1 in Appendix B.8.

2.3 Refreshing the Protected Basis

The default variant computes U_k, V_k once from the pretrained weights and keeps them fixed. During longer runs, the leading singular subspaces of the current weights can rotate away from this initial basis. SST can therefore re-anchor W_{ref} to the current weights and recompute U_k, V_k from it every K optimizer steps, using a Gram-based eigendecomposition for efficiency (Appendix B.13 discusses the minimal time overhead of SST compared to SFT).

A refresh defines the free complement relative to the current weights and opens new directions for adaptation. As the singular vectors rotate, a direction protected by the initial basis can acquire a component in a later free complement. Refresh can improve continued task learning, but it mildly weakens protection relative to the original model because SST now preserves a sequence of updated principal subspaces rather than one fixed pretrained subspace. We therefore treat K as an adaptation–retention control and report no refresh, R100 ($K{=}100$), and R10 ($K{=}10$) through many of our experiments.

3 Experiments

We ask the central question: *When a model is fine-tuned on a new task, does confining the update to the non-principal subspace let it learn the task while retaining what it already knew?*

Models and tasks. We fine-tune Qwen3-8B [56] on data containing explicit `<think>...</think>` traces in four domains: Science, Finance, Pathology, and Medicine. Science uses DeepSeek-R1-distilled traces from the `open-r1/Mixture-of-Thoughts` release [17]. Finance uses the reasoning-annotated FinQA/ConvFinQA release [12, 13]. Pathology and Medicine use MedMCQA questions [37] with traces generated by Claude Sonnet 4.6 [3].

Evaluation. *In-distribution accuracy* (ID) measures how well the new task was learned. *Prior capability* measures forgetting: the average over held-out general-purpose benchmarks not used for fine-tuning. We evaluate on a broad set of benchmarks before and after training, namely GPQA-Diamond [39], MMLU-Pro [53], GSM8K [14], HumanEval [11], and IFEval [59].

Baselines. We compare against standard supervised fine-tuning (SFT) and self-distillation fine-tuning [SDFT, 44], an on-policy method designed to reduce forgetting. We report SFT and SDFT alone, and each combined with SST at rank fraction r=0.99 unless specified otherwise (Appendix B.11 sweeps r and finds that stronger protection improves both retention and in-domain accuracy), optionally with basis refresh every 100 (R100) or 10 (R10) steps. We also include a Muon-optimizer variant [5] and the post-hoc KL-repair baseline ReInvoke [44] where relevant. Full training and evaluation details are in Appendix B.14.

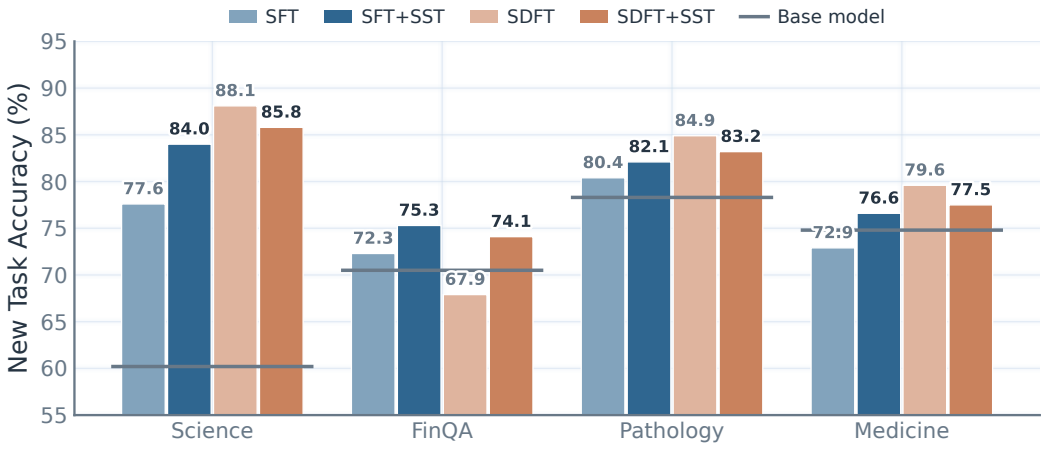

Figure 5: **Increase in in-domain accuracy** for each recipe and its +SST+R10 version (Table 4). Light bars are the recipe alone, full-color bars the +SST version, and the gray rule in each group is the base model. Refresh mechanism improves the in-domain accuracy matching SFT performance.

3.1 SST learns without forgetting

We begin by fine-tuning a thinking model on external chain-of-thought data from a new domain.

Unconstrained recipes collapse. Figure 6 shows broad losses in prior capability after both SFT and SDFT across the four domains. For example, as the model learns the new task, it often stops producing compilable code or respecting requested output formats. Post-hoc KL repair recovers only part of this capability and reduces in-domain gains, leaving a learning–retention trade-off rather than improving both simultaneously.

SST retains prior capability. Across domains, SST remains close to the base model while matching the in-domain performance of the corresponding unconstrained recipe. The benefit composes with SDFT and with the Muon optimizer, indicating it generalizes across training recipes. Figure 1 summarizes the improved learning–retention trade-off with R10 refresh. Figure 6 gives the per-benchmark breakdown for the no-refresh variant, whose protection is strictest; switching to R10 costs only $\sim 3-5$ additional points of prior capability

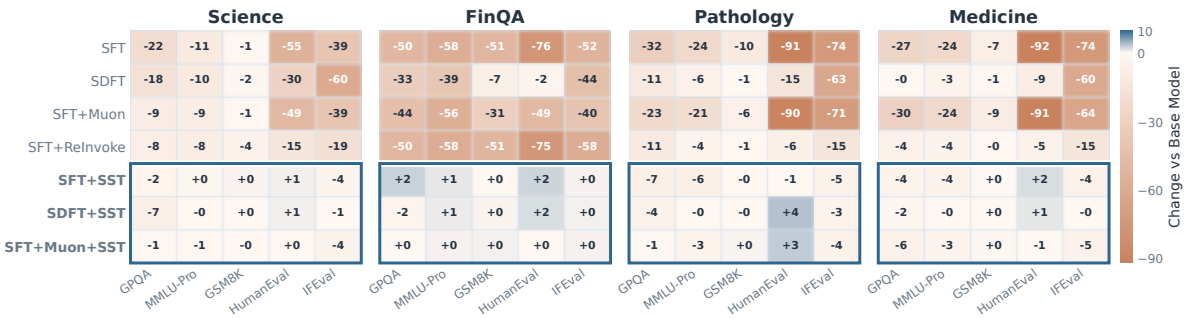

Figure 6: **Prior capability change compared to the base model** Unconstrained recipes incur broad capability losses. SST variants stay much closer to the base model across benchmarks and domains.

on average (Table 4). Figure 5 shows that with R10, retention does not come at the expense of target-task learning. More frequent refresh favors continued adaptation at a small retention cost, Appendix B.6 reports every variant. A non-thinking instruct model, Qwen2.5-7B-Instruct, shows the same learning–retention pattern across three domains including Tool Use (Appendix B.4).

3.2 SST scales well

SST reduces forgetting as parameters increase. Figure 2 shows the effect of scaling model size on the Science dataset, with each model trained for five epochs. Across model sizes, unconstrained SFT forgets substantially, losing 30-40 points of prior capability. By contrast, SST preserves the base model better as scale grows. SST also remains competitive with or improves upon SFT on the new task, so the retention benefit is not obtained by suppressing learning. Training details are in Appendix B.5, and a separate two-epoch sweep through 32B in Table 3 shows the same trend.

SST prevents monotonic decline in capability as the training data scales. Figure 3 shows the impact of scaling data on performance on prior and new tasks. We compare two setting under a 45,000-example budget. In the first, we have a novel stream of data and in the second we look at the multi-epoch setting: a 9k example subset repeated for five epochs. In both settings, SFT moves progressively farther from the base model as training proceeds, while SST maintains prior capability and continues improving on the target task. SST therefore remains effective in both the single-epoch novel-data setting and the multi-epoch repeated-data setting. The corresponding table can be found in Appendix B.5.

3.3 SST stays close to the base model

SST keeps the model close to its base distribution. Figure 7 measures divergence from the base model on WikiText-2 [32], a general-domain corpus. Across the training-data (Figure 7) and model-scale studies (Figure 14, Appendix B.7), SST remains substantially closer to the base distribution than the corresponding unconstrained model. The right panel of Figure 7 traces new-task accuracy against drift as training proceeds, and shows that SST improves accuracy with roughly 3× less movement away from the base policy. Across model sizes and training recipes, the reduction is 2-4× (Appendix B.7).

A natural alternative is to enforce closeness directly, by adding a reverse-KL penalty toward the base model during SFT. Appendix B.10 (Table 6, Figure 16) shows that this does not recover SST's trade-off. Weak regularization fails to mitigate forgetting, and stronger regularization suppresses task learning without preserving prior capability.

SST reduces hallucinations induced by self-distillation. Finally, we stress-test a difficult case for a reasoning model: finetuning with non-reasoning data. This setting explores a challenge where supervised finetuning can substantially degrade their reasoning capabilities. We fine-tune Olmo-3-7B-Think [35] on answer-only

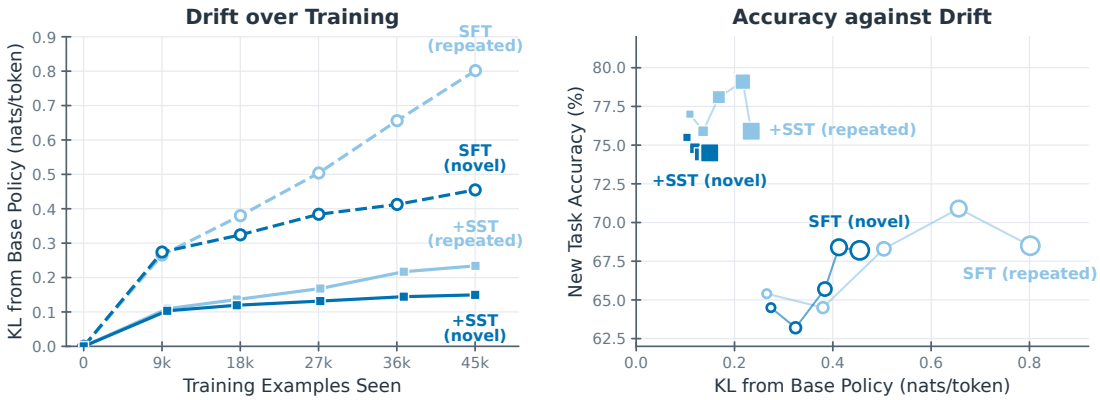

Figure 7: **Policy drift from the base model**, $\mathrm{KL}(\pi_{\text{base}} \,\|\, \pi_\theta)$ on WikiText-2, a general-domain dataset containing no task data. Left: drift over training. Right: the same drift against the in-domain accuracy, marker size growing with data seen.

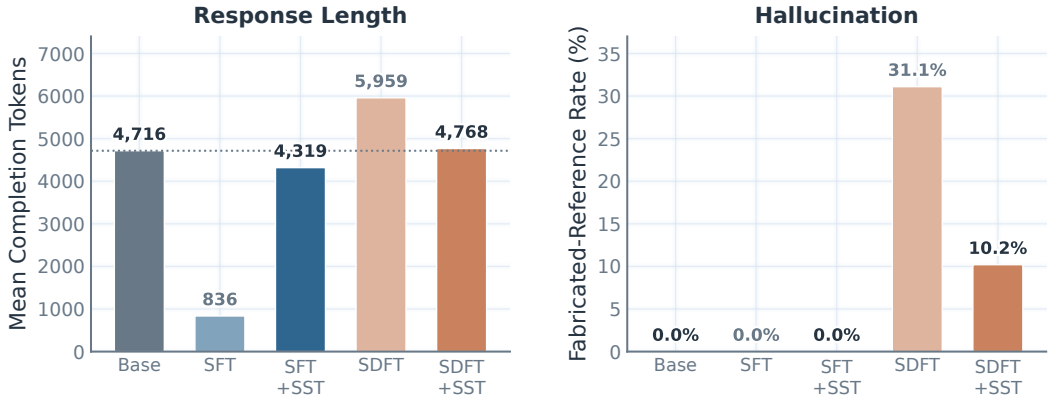

Figure 8: **Each recipe damages the base model in its own way, and SST pulls it back.** Olmo-3-7B-Think fine-tuned on answer-only medical data (+R10 arms shown); the gray bar in each panel is the base model. Left: Mean response length, with the dotted rule marking the base model's response length. Right: Rate of fabricating a "reference solution" the prompt never contained.

medical supervision from HuatuoGPT-o1 [10] and summarize the results in Figure 8. SFT improves answer accuracy but sharply shortens the model's reasoning, while SST keeps its response length close to the base model. SDFT maintains longer reasoning but can overshoot and fabricate references to a solution that was never provided. Adding SST's subspace constraint substantially reduces this behavior. Together, these results show that staying close to the base model preserves qualitative properties not captured by task accuracy alone. The plotted fabrication rate is a conservative phrase-based lower bound. Appendix B.14 gives its exact definition and the grading protocol.

3.4 SST preserves gains across successive fine-tuning stages

A stronger test is whether SST lets a model acquire several skills in sequence without each stage erasing the last. We fine-tune Qwen3-8B on two consecutive tasks ($A \to B$), tracking retention on task A and learning on task B. We also evaluate three-stage chains ($A \to B \to C$) and track each domain throughout.

SST learns more and forgets less. Figure 9 shows multiple $A \to B$ sequences across domain pairs. We use R10 refresh in these runs. Across all the tested combinations, SST outperforms SFT in both retention and new-task learning. Appendix B.2 adds SDFT to three of these chains. SDFT learns each target well but its retention is chain-dependent, for example losing 20 points on FinQA $\to$ Science where SST+R10 loses none.

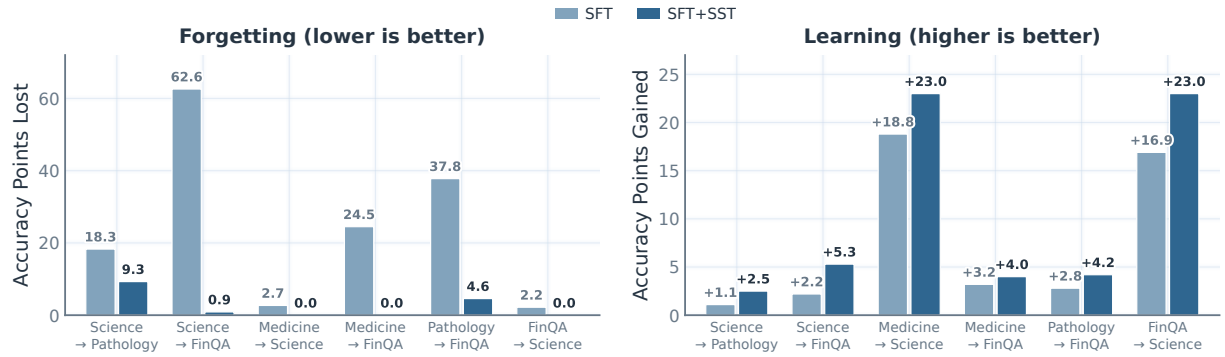

Figure 9: **Two-stage continual fine-tuning** ($A \to B$) **on Qwen3-8B.** Left: accuracy lost on domain A relative to its stage-1 level while stage 2 trains on B. Right: accuracy gained on B over the base model. SST arms use R10 refresh. SST improves *both* quantities in every tested chain.

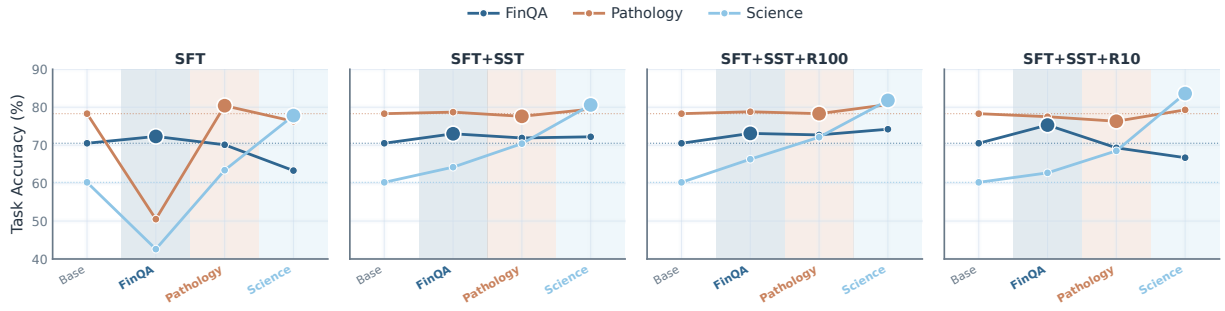

Figure 10: **Three-stage continual fine-tuning on Qwen3-8B.** The SFT family is trained on FinQA → Pathology → Science. Lines track task accuracy on all three domains, dotted lines mark each domain's base-model accuracy, each stage's background is tinted with the domain trained at that stage, and the enlarged marker is the domain just trained. A line should rise in its own stage and hold its level afterward. Under SFT the dormant domains drop sharply and recover only when trained, while the SST variants hold all three domains much closer to their trained levels.

SST accumulates task knowledge. In the FinQA → Pathology → Science chain of Figure 10, each domain improves when it is trained and remains near that level as later stages are learned, so the final model retains all of the acquired skills. Under plain SFT the same chain is not cumulative. Accuracy on the untrained Pathology and Science drops by 28 and 18 points while FinQA trains, and FinQA itself gives back 9 points over the later stages. The complementary Medicine → Pathology → Science chain in Appendix B.3 shows the same pattern for the SDFT family. The refresh variants expose a useful trade-off in this sequential setting. More frequent refresh improves accuracy on the current task at a minor cost to retention on previously learned tasks.

As a complementary composition test, Appendix B.1 shows that independently trained SST experts can be merged with far less interference than SFT experts. Figure 11 reports the exact task-level changes relative to each method's own experts.

4 Related Work

Classical continual learning restricts where updates go using stored data or gradients from earlier tasks [9, 18, 23, 27, 40, 51]. These constraints require data that is unavailable at LLM scale and task boundaries that deployed fine-tuning does not have, whereas SST needs neither. A complementary line reduces forgetting through the data instead of the update. Self-distillation rewrites supervision into the model's own distribution

[44], and on-policy RL forgets less because its updates stay close to the base policy [45, 60]. Those methods induce a favorable update geometry through on-policy data. SST, by contrast, constrains weight update, and as a result can compose with on-policy methods (Table 4). We defer discussion of many other related works to Appendix A.

4.1 Comparison with subspace-protecting methods

Two recent methods share our starting point of protecting the top singular directions of pretrained weights while adapting elsewhere. Both differ from SST in the same structural way. They rewrite the update as a product of trained factors and constrain the factors, whereas SST trains the weights directly and constrains the one quantity that reaches them, the applied update.

Sculpting [34] freezes the top singular vectors of each weight matrix and projects the gradient before every AdamW step. Its projection removes only block (1) of Eq. 2, the component that reads from and writes onto protected directions at once. The cross blocks (2) and (3), which touch a protected direction on one side, are never removed. The projection moreover acts on the gradient, rather than the update as suggested by our controlled study. Even removing all three blocks from the gradient leaves 69% of the applied update in protected directions (Figure 4).

OPLoRA [54] freezes the weights and trains a LoRA adapter, the product BA of two low-rank trainable matrices, sandwiched between fixed projectors, $\Delta W = Q_U\, BA\, Q_V$ in our notation. This is the projection Π of Eq. 3 applied to a low-rank product rather than to the optimizer's step, and it confines the update fully to the free subspace. The amount protected is set by the projector rank k, and their experiments protect only the top 16-128 directions while adapting at low rank in everything that remains, nearly the inverted regime from SST, which protects 99% of the spectrum and adapts densely in the tail. By the chain rule through the fixed projectors, the gradient components along protected directions never reach the trainable factors, and the applied update is projected again by construction. This corresponds to scheme (3) in Sec. 2.2, which our controlled study found mis-scales the step (Figure 4), with Adam's step sizes moreover set in factor space rather than for the weight update.

Another difference is that neither method considers refreshing the protected subspace. SST introduces refresh as part of the method, re-anchoring the basis every K steps, and studies it as an explicit retention-adaptation control - swept from no refresh to R10 throughout our experiments (Sec. 2.3).

We compare both methods on the four-domain setup under their official settings (Table 5). Both retain more capability than SFT but lose 18-26 points of prior capability on average where no-refresh SST loses about one, at comparable in-domain accuracy. Giving OPLoRA our protection level ($k{=}0.99d$) recovers much of the gap, yet it still trails SST on three of the four domains (Table 4). Under data scaling across Medicine, Pathology, and FinQA with official settings, both OPLoRA and Sculpting forget substantially more as training continues, while SST stays much closer to the base model (Appendix B.9, Figure 15). We therefore keep SFT and SDFT, the dominant fine-tuning recipes in practice, as the primary baselines throughout the paper, while comparing against Sculpting and OPLoRA in this main setting and on data scaling.

5 Conclusion

Catastrophic forgetting is usually treated as the price of learning new tasks. Our work studies it instead as a property of *where* the update lands, directly controls it, and mitigates forgetting. Confining the applied optimizer step to the complement of each weight matrix's top singular subspace (Eq. 3) turns recipes that destroy prior capability into ones that retain it. The same operator works across recipes (SFT, SDFT, and Muon), from 1.7B to 32B parameters, and across data scale, where SFT degrades monotonically and SST stays flat. It also turns sequential fine-tuning from destructive into cumulative and keeps the model 2–$4\times$ closer to its base distribution in KL, thus preserving behavior that accuracy alone does not capture, such as a thinking model's deliberation length and self-distillation's fabricated reference solutions.

Limitations and Future Work. SST changes the feasible update geometry, its optimal learning rate need not equal that of unconstrained SFT. Establishing fully tuned learning–retention frontiers remains important future work. Scaling SST to models with hundreds of billions of parameters and to longer task sequences is an important next test. Additionally, using information about the update for each batch to dynamically update the protection rather than refreshing at fixed intervals is a next step to make the algorithm more efficient.

References

[1] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023. 2

[2] Armen Aghajanyan, Sonal Gupta, and Luke Zettlemoyer. Intrinsic dimensionality explains the effectiveness of language model fine-tuning. In *Proceedings of the 59th annual meeting of the association for computational linguistics and the 11th international joint conference on natural language processing (volume 1: long papers)*, pages 7319–7328, 2021. 2

[3] Anthropic. Claude Sonnet 4.6 system card. Technical report, Anthropic, February 2026. URL https://www.anthropic.com/claude-sonnet-4-6-system-card. Revised March 6, 2026. 6

[4] Gal Berkooz, Philip Holmes, and John L Lumley. The proper orthogonal decomposition in the analysis of turbulent flows. *Annual review of fluid mechanics*, 25(1):539–575, 1993. 3

[5] Jeremy Bernstein and Laker Newhouse. Old optimizer, new norm: An anthology. *arXiv preprint arXiv:2409.20325*, 2024. 6

[6] Rishi Bommasani, Drew A Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, et al. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021. 2

[7] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020. 2

[8] Aaron Chatterji, Thomas Cunningham, David J Deming, Zoe Hitzig, Christopher Ong, Carl Yan Shan, and Kevin Wadman. How people use chatgpt. Technical report, National Bureau of Economic Research, 2025. 2

[9] Arslan Chaudhry, Marc'Aurelio Ranzato, Marcus Rohrbach, and Mohamed Elhoseiny. Efficient lifelong learning with a-gem. *arXiv preprint arXiv:1812.00420*, 2018. 9, 17

[10] Junying Chen, Zhenyang Cai, Ke Ji, Xidong Wang, Wanlong Liu, Rongsheng Wang, and Benyou Wang. Towards medical complex reasoning with llms through medical verifiable problems. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 14552–14573, 2025. 8, 19, 31

[11] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde De Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021. 6

[12] Zhiyu Chen, Wenhu Chen, Charese Smiley, Sameena Shah, Iana Borova, Dylan Langdon, Reema Moussa, Matt Beane, Ting-Hao Huang, Bryan R Routledge, et al. Finqa: A dataset of numerical reasoning over financial data. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 3697–3711, 2021. 6

[13] Zhiyu Chen, Shiyang Li, Charese Smiley, Zhiqiang Ma, Sameena Shah, and William Yang Wang. Convfinqa: Exploring the chain of numerical reasoning in conversational finance question answering. In *Proceedings of the 2022 conference on empirical methods in natural language processing*, pages 6279–6292, 2022. 6

[14] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. 6

[15] Cursor Research, Aaron Chan, Ahmed Shalaby, Alexander Wettig, Aman Sanger, Andrew Zhai, Anurag Ajay, Ashvin Nair, Charlie Snell, Chen Lu, Chen Shen, Emily Jia, Federico Cassano, Hanpeng Liu, Haoyu Chen, Henry Wildermuth, Jacob Jackson, Janet Li, Jediah Katz, Jiajun Yao, Joey Hejna, Josh Warner, Julius Vering, Kevin Frans, Lee Danilek, Less Wright, Lujing Cen, Luke Melas-Kyriazi, Michael Truell, Michiel de Jong, Naman Jain, Nate Schmidt, Nathan Wang, Niklas Muennighoff, Oleg Rybkin, Paul Loh, Phillip Kravtsov, Rishabh Yadav, Sahil Shah, Sam Kottler, Alexander M. Rush, Shengtong Zhang, Shomil Jain, Sriram Sankar, Stefan Heule, Stuart H. Sul, Sualeh Asif, Victor Rong, Wanqi Zhu, William Lin, Yuchen Wu, Yuri Volkov, Yury Zemlyanskiy, Zack Holbrook, and Zhiyuan Zhang. Composer 2 technical report, March 2026. URL https://arxiv.org/abs/2603.24477. 2

[16] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 conference of the North American chapter of the association for computational linguistics: human language technologies, volume 1 (long and short papers)*, pages 4171–4186, 2019. 2

[17] Hugging Face. Open r1: A fully open reproduction of deepseek-r1, January 2025. URL https://github.com/huggingface/open-r1. 6

[18] Mehrdad Farajtabar, Navid Azizan, Alex Mott, and Ang Li. Orthogonal gradient descent for continual learning. In *International conference on artificial intelligence and statistics*, pages 3762–3773. PMLR, 2020. 9, 17

[19] Kehua Feng, Xinyi Shen, Weijie Wang, Xiang Zhuang, Yuqi Tang, Qiang Zhang, and Keyan Ding. Sciknoweval: Evaluating multi-level scientific knowledge of large language models. *arXiv preprint arXiv:2406.09098*, 2024. 19, 31

[20] Kunal Handa, Alex Tamkin, Miles McCain, Saffron Huang, Esin Durmus, Sarah Heck, Jared Mueller, Jerry Hong, Stuart Ritchie, Tim Belonax, et al. Which economic tasks are performed with ai? evidence from millions of claude conversations. *arXiv preprint arXiv:2503.04761*, 2025. 2

[21] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2020. 19

[22] Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021. 2

[23] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences*, 114(13):3521–3526, 2017. 2, 9, 17

[24] Suhas Kotha, Jacob Springer, and Aditi Raghunathan. Understanding catastrophic forgetting in language models via implicit inference. In *International Conference on Learning Representations*, volume 2024, pages 24110–24139, 2024. 2, 17

[25] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving

with pagedattention. In *Proceedings of the 29th symposium on operating systems principles*, pages 611–626, 2023. 31

[26] Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human falsehoods. In *Proceedings of the 60th annual meeting of the association for computational linguistics (volume 1: long papers)*, pages 3214–3252, 2022. 19

[27] David Lopez-Paz and Marc'Aurelio Ranzato. Gradient episodic memory for continual learning. *Advances in neural information processing systems*, 30, 2017. 9, 17

[28] Kevin Lu and Thinking Machines Lab. On-policy distillation. *Thinking Machines Lab: Connectionism*, 2025. doi: 10.64434/tml.20251026. https://thinkingmachines.ai/blog/on-policy-distillation. 2

[29] Yun Luo, Zhen Yang, Fandong Meng, Yafu Li, Jie Zhou, and Yue Zhang. An empirical study of catastrophic forgetting in large language models during continual fine-tuning. *IEEE Transactions on Audio, Speech and Language Processing*, 2025. 2, 17

[30] Charles H Martin and Michael W Mahoney. Implicit self-regularization in deep neural networks: Evidence from random matrix theory and implications for learning. *Journal of Machine Learning Research*, 22(165):1–73, 2021. 2, 3

[31] Fanxu Meng, Zhaohui Wang, and Muhan Zhang. Pissa: Principal singular values and singular vectors adaptation of large language models. *Advances in Neural Information Processing Systems*, 37:121038–121072, 2024. 17

[32] Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models. *arXiv preprint arXiv:1609.07843*, 2016. 7

[33] Sagnik Mukherjee, Lifan Yuan, Dilek Hakkani-Tur, and Hao Peng. Reinforcement learning finetunes small subnetworks in large language models. *Advances in Neural Information Processing Systems*, 38: 132119–132138, 2026. 2

[34] Nikhil Shivakumar Nayak, Krishnateja Killamsetty, Ligong Han, Abhishek Bhandwaldar, Prateek Chanda, Kai Xu, Oleg Silkin, Mustafa Eyceoz, Hao Wang, Aldo Pareja, et al. Sculpting subspaces: Constrained full fine-tuning in llms for continual learning. In *International Conference on Learning Representations*, volume 2026, pages 20212–20243, 2026. 10, 24

[35] Team Olmo, Allyson Ettinger, Amanda Bertsch, Bailey Kuehl, David Graham, David Heineman, Dirk Groeneveld, Faeze Brahman, Finbarr Timbers, Hamish Ivison, et al. Olmo 3. *arXiv preprint arXiv:2512.13961*, 2025. 7

[36] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022. 2

[37] Ankit Pal, Logesh Kumar Umapathi, and Malaikannan Sankarasubbu. Medmcqa: A large-scale multi-subject multi-choice dataset for medical domain question answering. In *Conference on health, inference, and learning*, pages 248–260. PMLR, 2022. 6, 31

[38] Jeff Rasley, Samyam Rajbhandari, Olatunji Ruwase, and Yuxiong He. Deepspeed: System optimizations enable training deep learning models with over 100 billion parameters. In *Proceedings of the 26th ACM SIGKDD international conference on knowledge discovery & data mining*, pages 3505–3506, 2020. 31

[39] David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R Bowman. Gpqa: A graduate-level google-proof q&a benchmark. *arXiv preprint arXiv:2311.12022*, 2023. 6

[40] Gobinda Saha, Isha Garg, and Kaushik Roy. Gradient projection memory for continual learning. *arXiv preprint arXiv:2103.09762*, 2021. 9, 17

[41] Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. Winogrande: An adversarial winograd schema challenge at scale. *Communications of the ACM*, 64(9):99–106, 2021. 19

[42] Peter J Schmid. Dynamic mode decomposition of numerical and experimental data. *Journal of fluid mechanics*, 656:5–28, 2010. 3

[43] Pratyusha Sharma, Jordan Ash, and Dipendra Kumar Misra. The truth is in there: Improving reasoning in language models with layer-selective rank reduction. In *International Conference on Learning Representations*, volume 2024, pages 17632–17651, 2024. 2, 3

[44] Idan Shenfeld, Mehul Damani, Jonas Hübotter, and Pulkit Agrawal. Self-distillation enables continual learning. *arXiv preprint arXiv:2601.19897*, 2026. 2, 6, 10, 17, 19

[45] Idan Shenfeld, Jyothish Pari, and Pulkit Agrawal. Rl's razor: Why online reinforcement learning forgets less. In *International Conference on Learning Representations*, volume 2026, pages 59839–59864, 2026. 2, 10, 17

[46] Karan Singhal, Tao Tu, Juraj Gottweis, Rory Sayres, Ellery Wulczyn, Mohamed Amin, Le Hou, Kevin Clark, Stephen R Pfohl, Heather Cole-Lewis, et al. Toward expert-level medical question answering with large language models. *Nature medicine*, 31(3):943–950, 2025. 2

[47] Qiaoyu Tang, Ziliang Deng, Hongyu Lin, Xianpei Han, Qiao Liang, Boxi Cao, and Le Sun. Toolalpaca: Generalized tool learning for language models with 3000 simulated cases. *arXiv preprint arXiv:2306.05301*, 2023. 19, 31

[48] Rishabh Tiwari, Kusha Sareen, Lakshya A Agrawal, Joseph E Gonzalez, Matei Zaharia, Kurt Keutzer, Inderjit S Dhillon, Rishabh Agarwal, and Devvrit Khatri. Learning, fast and slow: Towards llms that adapt continually. *arXiv preprint arXiv:2605.12484*, 2026. 17

[49] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023. 2

[50] Hanqing Wang, Yixia Li, Shuo Wang, Guanhua Chen, and Yun Chen. Milora: Harnessing minor singular components for parameter-efficient llm finetuning. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 4823–4836, 2025. 17

[51] Shipeng Wang, Xiaorong Li, Jian Sun, and Zongben Xu. Training networks in null space of feature covariance for continual learning. In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 184–193. IEEE, 2021. 9, 17

[52] Xiao Wang, Tianze Chen, Qiming Ge, Han Xia, Rong Bao, Rui Zheng, Qi Zhang, Tao Gui, and Xuan-Jing Huang. Orthogonal subspace learning for language model continual learning. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 10658–10671, 2023. 17

[53] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, et al. Mmlu-pro: A more robust and challenging multi-task language understanding benchmark. *Advances in Neural Information Processing Systems*, 37:95266–95290, 2024. 6

[54] Yifeng Xiong and Xiaohui Xie. Oplora: Orthogonal projection lora prevents catastrophic forgetting during parameter-efficient fine-tuning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 34088–34096, 2026. 10, 24

[55] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024. 19

[56] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025. 2, 6

[57] Yibo Yang, Xiaojie Li, Zhongzhu Zhou, Shuaiwen L Song, Jianlong Wu, Liqiang Nie, and Bernard Ghanem. Corda: Context-oriented decomposition adaptation of large language models for task-aware parameter-efficient fine-tuning. *Advances in neural information processing systems*, 37:71768–71791, 2024. 17

[58] Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. Hellaswag: Can a machine really finish your sentence? In *Proceedings of the 57th annual meeting of the association for computational linguistics*, pages 4791–4800, 2019. 19

[59] Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*, 2023. 6, 19

[60] Hanqing Zhu, Zhenyu Zhang, Hanxian Huang, DiJia Su, Zechun Liu, Jiawei Zhao, Igor Fedorov, Hamed Pirsiavash, Zhizhou Sha, Jinwon Lee, et al. The path not taken: Rlvr provably learns off the principals. *arXiv preprint arXiv:2511.08567*, 2025. 2, 10

Appendix Contents

A Extended Related Work **17**

B Additional Experiments & Details **17**

B.1 Model Merging . . . 17
B.2 Two-stage continual learning with SDFT . . . 18
B.3 Three-stage continual learning with SDFT . . . 18
B.4 Instruct tuning without forgetting . . . 19
B.5 Scaling Experiments . . . 19
B.6 Qwen3-8B thinking: full result tables . . . 19
B.7 Policy drift across model sizes . . . 21
B.8 Toy study: full specification . . . 21
B.9 Comparison with subspace-constrained baselines . . . 24
B.10 KL regularisation does not prevent forgetting . . . 28
B.11 Rank ablation on FinQA . . . 28
B.12 Effect of Optimizer Hyperparameters on SST . . . 29
B.13 Computational Cost . . . 30
B.14 Training and evaluation details . . . 31

A Extended Related Work

Classical continual-learning approaches often restrict the update direction: regularization penalizes movement along Fisher-important directions [23], while projected-gradient methods enforce constraints by replacing gradient g with a projection $\Pi(g)$ built from stored data, past-task gradients [9, 18, 27], activation subspaces [40], the feature-covariance null space [51], or earlier tasks' adapters [52]. Continual learning for LLMs differs from these settings in several ways. These methods estimate subspaces from data, but reconstructing a distribution representative of an LLM's pretraining corpus is infeasible at scale. Truly continual deployment may also lack explicit task boundaries. SST instead preserves a subspace derived from the weight spectrum (Eq. 3), requiring neither stored examples nor task boundaries.

Several parameter-efficient fine-tuning methods also construct adapters from singular directions. PiSSA [31] initializes and trains LoRA factors from the principal singular components, prioritizing high-energy directions rather than protecting them. MiLoRA [50] is closer in motivation: it retains the principal components and adapts minor singular components, but represents the update through a fixed low-rank factorization. CorDA [57] first forms a context-oriented decomposition using input-activation covariance and initializes low-rank adapters from task- or knowledge-oriented components. In all three methods, the spectral choice is encoded in a parameter-efficient factorization derived at initialization. SST instead projects the dense update actually proposed by the optimizer into the two-sided non-principal complement. It therefore constrains the realized weight change after stateful optimizer preconditioning, does not rely on a trainable adapter factorization or task-specific covariance estimate, and can refresh its protected basis during training.

A complementary line leaves the optimizer alone and reduces the distribution shift driving forgetting: self-distillation rewrites the supervision into the model's own output distribution [44], and online RL forgets less because its updates stay KL-close to the base policy [45, 48], both consistent with forgetting tracking distributional movement [24, 29]. These methods act on the data, whereas SST acts on the update, and the two compose (Table 4).

B Additional Experiments & Details

This appendix contains additional results supporting the main text. In particular, we report extra results on model merging independently fine-tuned experts, the Qwen2.5 instruct-model, complete model- and data-scaling results, full Qwen3 tables, policy drift across model sizes, the toy study and algorithm specification, comparisons with Sculpting and OPLoRA, reverse-KL and protected-rank controls, and full training and evaluation details.

B.1 Model Merging

If a constrained update stays in the low-energy complement of the weight spectrum, then two updates trained independently should interfere less when added back together. We test this with model merging experiments. In particular, we fine-tune a separate single-domain expert for Science, Medicine, and FinQA, then form $W_{\text{merged}} = W_{\text{base}} + \sum_d (W_d - W_{\text{base}})$ over all the 252 weight matrices. We copy every other parameter from the base model. No merge coefficients are tuned and no data is used to fit the combination. Each merged model is then scored in-domain against *its own* single-domain expert under identical decode settings. We report the interference caused by merging rather than the raw accuracy gap, which already favors SST.

SST experts compose, SFT experts collide. Figure 11 shows that independently trained SFT experts interfere strongly when their task vectors are combined, and the interference compounds as more experts are merged. SST experts compose much more cleanly: pairwise merges remain close to their corresponding single-domain experts, and the three-expert merge degrades substantially more slowly. Each SST expert modifies only the low-energy subspace, so adding their updates leaves the protected principal directions unchanged. Unconstrained experts can each alter these directions, causing their changes to accumulate during merging.

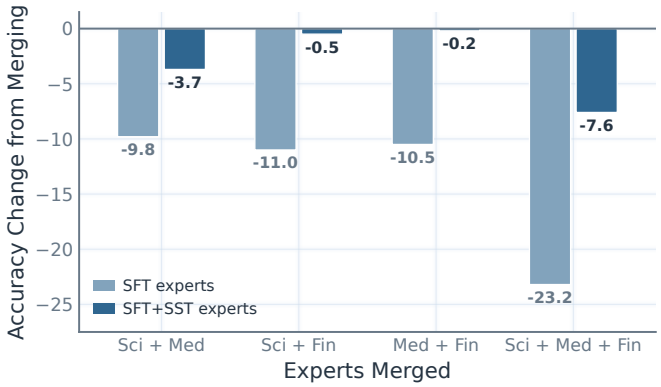

Figure 11: Interference from task-arithmetic merging. Each bar is the mean change in task accuracy relative to the same method's own single-domain experts.

B.2 Two-stage continual learning with SDFT

Figure 12 adds plain SDFT to the same two-stage setup as Figure 9. SDFT improves target learning over SFT in all three chains, but its retention is chain-dependent. It reduces forgetting on Science → Pathology and Pathology → FinQA, yet loses 20.0 points on FinQA → Science. SST+R10 gives the most consistent retention–learning trade-off.

For Pathology → FinQA, we regrade the saved SDFT generations by prioritizing the final boxed numeric answer before the existing `Answer:` fallback. This corrects an evaluator-format mismatch, giving 75.2 target accuracy (+4.7 over base); the source-task score, and therefore the 5.6 points of forgetting, is unchanged.

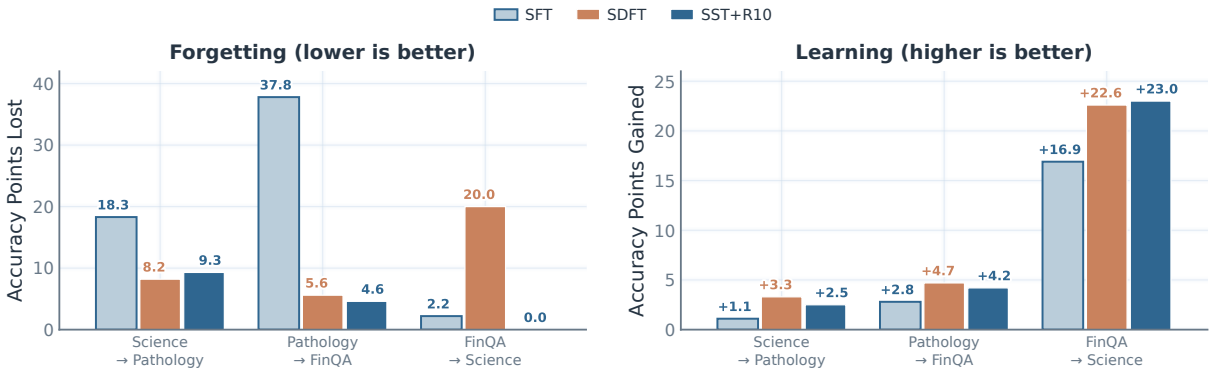

Figure 12: **Two-stage continual fine-tuning with SDFT** ($A \rightarrow B$) **on Qwen3-8B.** Left: accuracy lost on domain A relative to its stage-1 level after training on B. Right: accuracy gained on B over the base model. All arms use the same task chains and stage budgets; SST+R10 denotes SFT with SST and R10 refresh.

B.3 Three-stage continual learning with SDFT

Figure 13 shows the Medicine → Pathology → Science continual learning chain discussed in the main text on SDFT. Across the three stages, plain SDFT consistently degrades performance on Medicine and Science. SST stabilizes these trajectories, especially the R10 variant accumulates learning across all three domains.

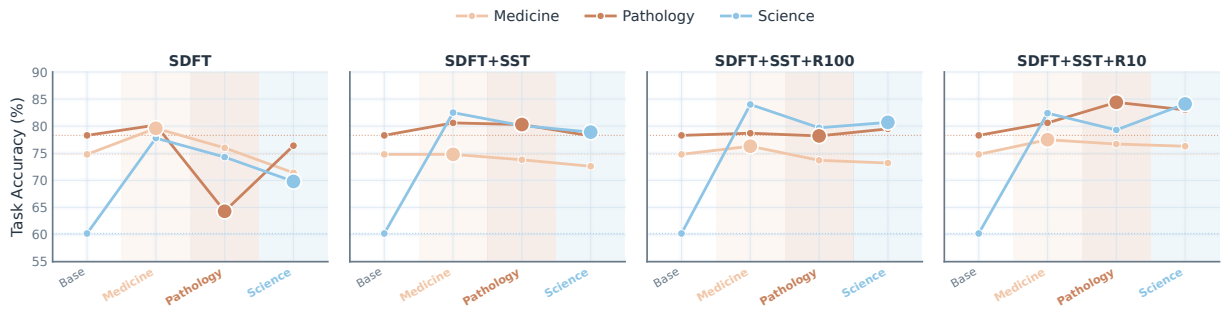

Figure 13: **Three-stage SDFT chain on Qwen3-8B.** The SDFT family is trained on Medicine → Pathology → Science. Lines track all three domains, dotted lines mark base-model accuracy, and the background tint identifies the domain trained at each stage.

B.4 Instruct tuning without forgetting

We also study a non-thinking instruction model with shorter sequences, where forgetting is less severe. Table 1 reports Qwen2.5-7B-Instruct [55] on ToolAlpaca [47], HuatuoGPT-o1 [10], and SciKnowEval [19], matching the SDFT setup [44]. All cells use rank fraction r=0.99 except Science, which uses r=0.90. On this model and domain, the tail left free at r=0.99 was too small to fit the task. This is the only setting in the paper that required a looser protected rank.

SFT learns the task but forgets. Table 1 shows that standard fine-tuning improves the target task but reduces prior capability, with the clearest losses in instruction following and knowledge. SDFT moderates this effect in some settings but is not consistently protective across domains.

SST holds capability at matched fit. Adding SST recovers much of the lost prior capability on Tool Use and Medical while remaining competitive on in-domain accuracy. Science is a narrower trade-off: SST variants can improve task learning, but no single variant dominates both learning and retention. Table 1 gives the exact per-domain and per-benchmark values. The capability average covers HellaSwag [58], TruthfulQA-MC1 [26], MMLU [21], IFEval [59], and WinoGrande [41].

B.5 Scaling Experiments

The complete scaling controls reproduce the main trend. Across both model-scale sweeps, SST preserves substantially more prior capability than unconstrained tuning while remaining competitive on the target task, and the retention advantage strengthens with model size. The matched-budget data comparison shows the same separation under both repeated and novel exposure (Figure 3 in the main text): SFT drifts steadily, whereas SST remains stable and continues learning. Tables 2 and 3 report the exact values.

B.6 Qwen3-8B thinking: full result tables

The full per-domain sweep. Across Science, Finance, Pathology, and Medicine, unconstrained recipes learn the target task but broadly degrade prior capability, while adding SST keeps the model much closer to its baseline while retaining strong target-task performance. Table 4 reports the exact values for every recipe and benchmark.

Table 1: Qwen2.5-7B-Instruct. ID = in-distribution accuracy; LM average over HellaSwag, TruthfulQA-MC1, MMLU, IFEval, and WinoGrande; Δ is relative to the pretrained baseline (68.5). Bold entries mark column maxima among non-base recipes within each domain.

Domain	Method	ID	HellaSwag	TQA-mc1	MMLU	IFEval	Winogr.	LM avg (Δ)
ToolUse	Baseline	—	80.5	47.6	71.8	71.5	71.0	68.5 $(+0.0)$
	SFT	70.1	**74.3**	39.5	57.7	45.7	**68.7**	57.2 (-11.3)
	SFT+SST	70.1	65.5	**47.6**	68.4	70.4	62.6	62.9 (-5.6)
	SFT+SST+R100	69.1	65.3	47.2	68.3	69.5	63.0	62.7 (-5.8)
	SFT+SST+R10	71.1	67.6	47.4	68.5	69.7	63.3	63.3 (-5.2)
	SDFT	70.1	70.8	44.2	**68.9**	65.8	67.7	63.5 (-5.0)
	SDFT+SST	66.0	66.7	43.5	68.5	71.3	64.3	62.9 (-5.6)
	SDFT+SST+R100	64.9	67.0	44.6	68.8	71.7	63.8	63.2 (-5.3)
	SDFT+SST+R10	**75.3**	68.2	44.7	68.8	**74.5**	63.5	**63.9** (-4.6)
Medical	Baseline	—	80.5	47.6	71.8	71.5	71.0	68.5 $(+0.0)$
	SFT	20.5	**72.3**	36.5	68.5	61.9	67.6	61.4 (-7.1)
	SFT+SST	19.5	70.8	41.4	69.6	68.8	68.3	**63.8** (-4.7)
	SFT+SST+R100	19.0	71.1	41.2	69.3	68.2	67.6	63.5 (-5.0)
	SFT+SST+R10	19.2	70.5	41.7	**70.1**	68.9	67.0	63.7 (-4.8)
	SDFT	**29.7**	67.1	39.7	55.7	58.8	**68.4**	57.9 (-10.6)
	SDFT+SST	26.9	68.0	40.5	69.0	**73.0**	68.0	63.7 (-4.8)
	SDFT+SST+R100	29.6	68.1	41.7	69.2	72.1	68.1	**63.8** (-4.7)
	SDFT+SST+R10	28.0	67.3	**41.9**	68.3	**73.0**	67.2	63.5 (-5.0)
Science	Baseline	—	80.5	47.6	71.8	71.5	71.0	68.5 $(+0.0)$
	SFT	68.6	**73.2**	43.0	69.2	61.4	68.2	**63.0** (-5.5)
	SFT+SST	67.7	71.3	42.5	**70.6**	63.8	66.1	62.8 (-5.7)
	SFT+SST+R100	67.5	71.4	42.4	69.7	63.6	66.8	62.8 (-5.7)
	SFT+SST+R10	69.8	**73.2**	42.2	68.2	55.3	67.2	61.2 (-7.3)
	SDFT	69.0	70.1	**43.7**	67.2	61.4	66.5	61.8 (-6.7)
	SDFT+SST	67.1	68.9	43.5	67.4	62.5	66.4	61.7 (-6.8)
	SDFT+SST+R100	**70.8**	69.3	42.8	66.1	**64.1**	68.0	62.1 (-6.4)
	SDFT+SST+R10	70.0	69.7	41.1	65.6	50.3	**69.0**	59.1 (-9.4)

Table 2: **Five-epoch model-scale results underlying Figure 2.** Science fine-tuning with the same configuration at each size. LM avg is the five-benchmark average and Δ is relative to each model's own base. Bold entries mark column maxima among non-base recipes within each model size.

Model	Method	ID	GPQA	MMLU-Pro	GSM8K	HumanEval	IFEval	LM avg	Δ
Qwen3-14B	Base	69.4	62.1	75.5	95.6	93.9	85.8	82.6	$+0.0$
	SFT	78.0	39.9	48.9	72.4	36.6	26.8	44.9	-37.7
	SFT+SST	86.5	**63.6**	75.4	94.5	**95.1**	79.7	**81.7**	$\mathbf{-0.9}$
	SFT+SST+R10	**88.1**	62.6	**76.2**	**94.7**	89.6	**82.1**	81.0	-1.6
Qwen3-8B	Base	60.2	60.1	73.0	94.8	91.5	85.6	81.0	$+0.0$
	SFT	76.1	35.4	48.7	81.0	36.0	34.2	47.1	-33.9
	SFT+SST	**85.4**	**61.1**	72.3	94.4	91.5	**81.1**	**80.1**	$\mathbf{-0.9}$
	SFT+SST+R10	85.3	58.6	**74.1**	**94.5**	**92.7**	80.2	80.0	-1.0
Qwen3-4B	Base	67.9	55.1	70.0	94.0	92.1	82.8	78.8	$+0.0$
	SFT	73.4	32.3	44.2	78.3	29.9	27.0	42.3	-36.5
	SFT+SST	80.4	51.5	67.7	**94.1**	**90.2**	72.5	75.2	-3.6
	SFT+SST+R10	**84.3**	**53.0**	**68.7**	93.9	86.0	**75.6**	**75.4**	$\mathbf{-3.4}$
Qwen3-1.7B	Base	54.1	38.4	58.6	89.3	83.5	72.3	68.4	$+0.0$
	SFT	69.5	27.3	38.6	72.4	23.8	27.9	38.0	-30.4
	SFT+SST	67.1	32.3	54.3	**87.1**	67.7	**60.4**	60.4	-8.0
	SFT+SST+R10	**70.4**	**34.8**	**56.0**	86.6	**71.3**	58.6	**61.5**	$\mathbf{-6.9}$

B.7 Policy drift across model sizes

SST reduces policy drift at every model scale. Figure 14 reports WikiText-2 KL divergence at the final checkpoint, with each model scored against its own pretrained base. Across sizes, SST remains substantially closer to the base distribution than SFT, and the separation persists at the largest scale. The figure gives the exact drift values for each model.

B.8 Toy study: full specification

Applied-update projection is the only constraint that is both exact and stable. The controlled setup below isolates where the projection is applied. All schemes learn the toy task, but only projecting the optimizer's applied update eliminates protected-subspace leakage without substantially changing the held-out capability.

Model, optimizer, targets. Qwen2.5-0.5B, fine-tuned with plain AdamW at learning rate 1×10^{-4} and weight decay 0 (unless stated). Each scheme's projection is applied to the standard target linears—attention q, k, v, o and MLP $gate, up, down$.

Task (8 fixed examples). Next-token cross-entropy on the completion tokens only (prompt tokens masked), on eight short prompt→completion pairs: `("The capital of France is", " Paris.")`, `("2 plus 3 equals", " 5.")`, `("The opposite of hot is", " cold.")`, `("Water is made of hydrogen and", " oxygen.")`, `("The first month of the year is", " January.")`, `("A group of wolves is called a", " pack.")`, `("The chemical symbol for gold is", " Au.")`, `("Seven times eight is", " 56.")`. The task is trivially learnable, so every scheme drives task loss to ≈ 0; the discriminator is the forgetting on out of domain tasks.

Table 3: **Two-epoch model-scale sweep on Science.** Δ is relative to each model's own baseline. Cells marked — were not produced; bold entries mark column maxima among non-base recipes within each model size.

Model	Method	ID	GPQA	MMLU-Pro	GSM8K	HumanEval	IFEval	LM avg	Δ
Qwen3-32B	Baseline	68.2	64.6	73.3	95.2	94.5	85.2	82.6	+0.0
	SFT	83.6	50.5	62.4	93.6	80.5	48.2	67.0	−15.6
	SFT+SST	84.1	62.6	73.5	94.8	**97.0**	79.9	81.6	−1.0
	SFT+SST+R100	85.6	58.6	73.8	**95.2**	95.1	**81.0**	80.7	−1.9
	SFT+SST+R10	**87.3**	**65.7**	**74.3**	95.1	94.5	79.1	**81.7**	**−0.9**
Qwen3-14B	Baseline	69.4	62.1	75.5	95.6	93.9	85.8	82.6	+0.0
	SFT	82.3	50.0	64.7	93.9	64.0	54.5	65.4	−17.2
	SFT+SST	84.0	**66.2**	75.5	**94.8**	93.9	82.6	**82.6**	**+0.0**
	SFT+SST+R100	85.0	59.6	**75.9**	94.5	94.5	**84.8**	81.9	−0.7
	SFT+SST+R10	86.3	60.6	75.6	**94.8**	92.1	81.0	80.8	−1.8
	SDFT	85.2	38.9	60.2	91.0	77.4	50.6	63.6	−19.0
	SDFT+SST	83.7	53.0	71.1	94.4	**95.7**	83.7	79.6	−3.0
	SDFT+SST+R100	87.9	53.5	73.5	94.7	94.5	84.1	80.1	−2.5
	SDFT+SST+R10	**88.1**	55.1	73.9	94.4	94.5	82.8	80.1	−2.5
Qwen3-8B	Baseline	60.2	60.1	73.0	94.8	91.5	85.6	81.0	+0.0
	SFT	77.6	38.4	62.3	93.9	36.6	46.2	55.5	−25.5
	SFT+SST	80.1	55.1	72.0	94.2	92.7	80.4	78.9	−2.1
	SFT+SST+R100	81.7	58.1	71.9	94.8	91.5	82.1	79.7	−1.3
	SFT+SST+R10	84.0	**58.6**	**74.4**	94.9	92.7	80.4	80.2	−0.8
	SDFT	**88.1**	42.4	62.5	92.9	61.6	25.3	56.9	−24.1
	SDFT+SST	82.9	53.0	72.5	94.8	92.1	**85.0**	79.5	−1.5
	SDFT+SST+R100	83.4	57.1	73.1	**95.1**	93.3	84.7	**80.7**	**−0.3**
	SDFT+SST+R10	85.8	56.6	70.2	93.6	**95.7**	84.1	80.0	−1.0
Qwen3-4B	Baseline	67.9	55.1	70.0	94.0	92.1	82.8	78.8	+0.0
	SFT	76.5	38.9	61.1	92.1	24.4	42.1	51.7	−27.1
	SFT+SST	78.1	53.0	68.2	**93.8**	**93.9**	79.9	**77.8**	**−1.0**
	SFT+SST+R100	78.7	**55.1**	67.8	93.6	92.7	79.5	77.7	−1.1
	SFT+SST+R10	80.6	53.5	**68.6**	93.5	88.4	79.5	76.7	−2.1
	SDFT	81.9	34.8	51.5	84.1	58.5	42.3	54.2	−24.6
	SDFT+SST	81.1	42.4	61.7	92.4	89.6	73.8	72.0	−6.8
	SDFT+SST+R100	83.9	44.9	65.0	92.6	92.7	77.8	74.6	−4.2
	SDFT+SST+R10	**86.1**	48.5	67.9	92.7	90.9	**80.2**	76.0	−2.8
Qwen3-1.7B	Baseline	54.1	38.4	58.6	89.3	83.5	72.3	68.4	+0.0
	SFT	66.2	29.3	51.5	85.2	48.2	34.8	49.8	−18.6
	SFT+SST	65.3	31.8	**55.8**	87.6	**75.0**	63.6	62.8	−5.6
	SFT+SST+R100	65.1	33.3	55.1	**88.2**	72.6	**65.4**	**62.9**	**−5.5**
	SFT+SST+R10	63.8	**37.4**	54.3	86.8	69.5	61.0	61.8	−6.6
	SDFT	75.4	27.3	36.0	57.2	41.5	25.7	37.5	−30.9
	SDFT+SST	65.7	23.2	40.4	80.3	69.5	61.7	55.0	−13.4
	SDFT+SST+R100	72.7	27.3	45.5	82.6	72.6	64.1	58.4	−10.0
	SDFT+SST+R10	**76.1**	30.3	48.6	84.2	74.4	61.7	59.8	−8.6

Table 4: Qwen3-8B (thinking mode). Methods are grouped by recipe. Bold entries mark column maxima among non-base recipes within each domain. Δ = LM avg $-$ baseline (81.0).

Domain	Method	ID	GPQA	MMLU-Pro	GSM8K	HumanEval	IFEval	LM avg	Δ
Science	Baseline	60.2	60.1	73.0	94.8	91.5	85.6	81.0	+0.0
	SFT	77.6	38.4	62.3	93.9	36.6	46.2	55.5	−25.5
	SFT+SST	80.1	57.6	73.1	**95.1**	92.1	82.1	80.0	−1.0
	SFT+SST+R100	81.7	57.6	73.1	**95.1**	92.1	82.1	80.0	−1.0
	SFT+SST+R10	84.0	58.6	**74.4**	94.9	92.7	80.4	80.2	−0.8
	SDFT	**88.1**	42.4	62.5	92.9	61.6	25.3	57.0	−24.0
	SDFT+SST	82.9	53.0	72.5	94.8	92.1	**85.0**	79.5	−1.5
	SDFT+SST+R100	83.4	57.1	73.1	**95.1**	93.3	84.7	**80.7**	**−0.3**
	SDFT+SST+R10	85.8	56.6	70.2	93.6	**95.7**	84.1	80.0	−1.0
	SFT+Muon	80.2	51.0	64.4	93.8	42.1	46.3	59.5	−21.5
	SFT+Muon+SST	80.9	**59.1**	72.0	94.5	91.5	81.3	79.7	−1.3
	SFT+Muon+SST+R10	82.1	55.6	73.3	95.0	90.2	82.8	79.4	−1.6
	SFT+ReInvoke	68.4	52.5	64.6	91.1	76.8	66.4	70.3	−10.7
	OPLoRA (k=0.99d)	80.8	53.5	72.2	94.9	90.9	84.7	79.2	−1.8
FinQA	Baseline	70.5	60.1	73.0	94.8	91.5	85.6	81.0	+0.0
	SFT	72.3	10.1	15.5	43.9	15.9	33.1	23.7	−57.3
	SFT+SST	73.0	**62.6**	**74.1**	94.8	**93.9**	86.0	**82.3**	**+1.3**
	SFT+SST+R100	73.1	53.5	73.9	94.7	93.3	84.8	80.1	−0.9
	SFT+SST+R10	**75.3**	40.4	68.3	94.6	89.0	**86.3**	75.7	−5.3
	SDFT	67.9	27.3	34.2	88.1	89.0	41.4	56.0	−25.0
	SDFT+SST	64.1	57.6	73.9	95.1	93.3	86.1	81.2	+0.2
	SDFT+SST+R100	67.6	55.1	72.3	94.5	**93.9**	84.3	80.0	−1.0
	SDFT+SST+R10	74.1	51.0	65.5	95.2	**93.9**	80.8	77.3	−3.7
	SFT+Muon	74.9	16.2	16.9	63.7	42.1	45.8	36.9	−44.1
	SFT+Muon+SST	71.4	60.1	73.4	**95.3**	91.5	86.0	81.2	+0.2
	SFT+Muon+SST+R10	74.3	36.9	67.9	94.1	92.1	85.0	75.2	−5.8
	SFT+ReInvoke	72.2	9.6	15.3	43.7	16.5	27.4	22.5	−58.5
	OPLoRA (k=0.99d)	71.5	32.8	63.4	93.6	89.0	82.8	72.3	−8.7
Pathology	Baseline	78.3	60.1	73.0	94.8	91.5	85.6	81.0	+0.0
	SFT	80.4	28.3	48.6	84.5	0.6	11.6	34.7	−46.3
	SFT+SST	77.3	53.0	67.0	94.4	90.9	81.0	77.2	−3.8
	SFT+SST+R100	77.6	55.1	68.6	94.8	93.3	81.0	78.5	−2.5
	SFT+SST+R10	82.1	39.9	65.6	91.4	92.7	72.1	72.3	−8.7
	SDFT	**84.9**	49.0	67.5	94.2	76.2	22.9	62.0	−19.0
	SDFT+SST	82.0	56.1	72.5	94.6	**95.1**	82.6	80.2	−0.8
	SDFT+SST+R100	80.3	**59.1**	**72.8**	95.1	92.1	**83.4**	**80.5**	**−0.5**
	SDFT+SST+R10	83.2	53.0	72.0	**95.6**	93.9	82.3	79.4	−1.6
	SFT+Muon	79.3	36.9	52.2	89.1	1.8	14.2	38.8	−42.2
	SFT+Muon+SST	77.0	**59.1**	70.0	95.1	94.5	81.5	80.0	−1.0
	SFT+Muon+SST+R10	78.4	51.0	64.0	95.4	93.3	78.7	76.5	−4.5
	SFT+ReInvoke	77.6	49.5	69.5	93.9	85.4	70.2	73.7	−7.3
	OPLoRA (k=0.99d)	76.2	38.9	58.4	93.6	91.5	71.3	70.7	−10.3
Medicine	Baseline	74.8	60.1	73.0	94.8	91.5	85.6	81.0	+0.0
	SFT	72.9	33.3	49.0	87.6	0.0	11.6	36.3	−44.7
	SFT+SST	72.4	56.6	68.8	94.8	93.3	81.1	78.9	−2.1
	SFT+SST+R100	74.3	55.1	68.1	94.7	93.3	80.2	78.3	−2.7
	SFT+SST+R10	76.6	46.0	65.9	86.8	87.2	70.2	71.2	−9.8
	SDFT	**79.6**	59.6	70.1	93.9	82.9	25.1	66.3	−14.7
	SDFT+SST	74.8	58.1	**72.5**	94.8	92.7	85.2	80.7	−0.3
	SDFT+SST+R100	76.3	**60.6**	**72.5**	**95.1**	**94.5**	**85.4**	**81.6**	**+0.6**
	SDFT+SST+R10	77.5	52.5	72.1	94.1	93.9	82.3	79.0	−2.0
	SFT+Muon	71.6	29.8	49.1	86.1	0.6	21.7	37.5	−43.5
	SFT+Muon+SST	72.8	54.5	70.0	94.9	90.9	80.8	78.2	−2.8
	SFT+Muon+SST+R10	75.0	52.5	63.8	94.8	92.1	78.9	76.4	−4.6
	SFT+ReInvoke	71.6	55.6	69.4	94.4	86.6	70.6	75.3	−5.7
	OPLoRA (k=0.99d)	70.8	34.3	59.5	93.0	92.1	71.0	70.0	−11.0

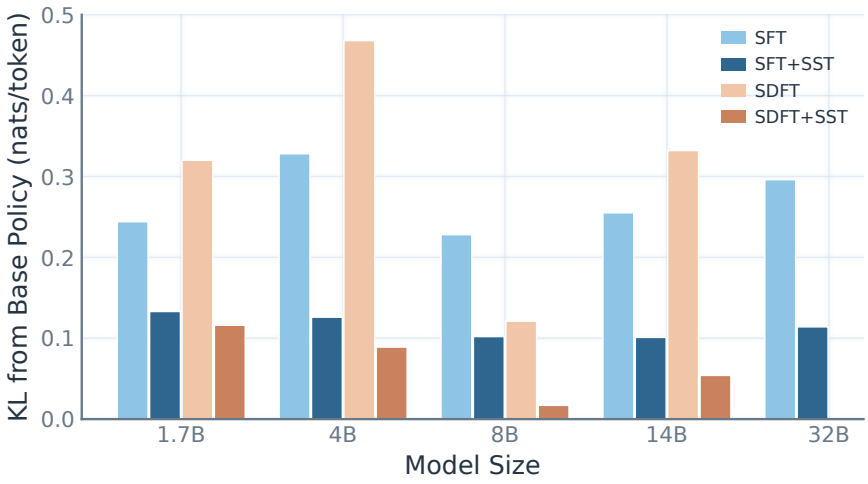

Figure 14: Policy drift at the final fine-tuned checkpoint across Qwen3 sizes. Light bars are the recipe alone, full-color bars their + SST version (no-refresh).

Capability probe (held-out). Mean per-token NLL on four generic sentences the model never trains on; **Probe**$\Delta = \mathrm{NLL}_{\text{after}} - \mathrm{NLL}_{\text{before}}$ (base-model probe NLL $\approx$2.81 nats/token). The probe sentences: *"The history of the Roman Empire spans over a thousand years of political change."*; *"Photosynthesis converts carbon dioxide and water into glucose using sunlight."*; *"In distributed systems, consensus algorithms ensure agreement despite failures."*; *"The novel explores themes of memory, loss, and the passage of time."*

The three enforcement schemes. With projector Π (Eq. 3), Adam moments $\hat{m}, \hat{v}$, and preconditioned raw step $\hat{m}/(\sqrt{\hat{v}} + \epsilon)$:

- **Project gradient**: $g \leftarrow \Pi(g)$; then $\hat{m}, \hat{v} \leftarrow \mathrm{Adam}(g)$; $W \leftarrow W - \mathrm{lr} \cdot \hat{m}/(\sqrt{\hat{v}} + \epsilon)$. Here $\hat{v}$ sees only free-subspace energy, and the applied step $D\,\Pi(\hat{m})$ with $D = \mathrm{diag}(1/(\sqrt{\hat{v}} + \epsilon))$ re-enters the protected space because D and Π do not commute.
- **Project the applied update**: $\hat{m}, \hat{v} \leftarrow \mathrm{Adam}(g)$ on the *full* gradient; $W_{\text{pre}} \leftarrow W$; $W \leftarrow W - \mathrm{lr} \cdot \hat{m}/(\sqrt{\hat{v}} + \epsilon)$; $W \leftarrow W_{\text{pre}} + \Pi(W - W_{\text{pre}})$. The update is confined to the low singular value subspace because projection is last, and the free step keeps the scale induced by the full gradient.
- **Re-project cumulative change**: $g \leftarrow \Pi(g)$; take the Adam step; then set $W \leftarrow W_{\text{ref}} + \Pi(W - W_{\text{ref}})$. Note that $W \leftarrow W_{\text{ref}} + \Pi(W - W_{\text{ref}})$ is equivalent to $W_t + \Pi(\delta_t)$, with δ_t the step just applied, since every previously applied change already lies in the free subspace. The update lies in the correct subspace, but the step is based on projected-gradient moments and is therefore mis-scaled.

The full procedure. Algorithm 1 details SST's implementation, using the second of these three schemes. A superscript ℓ indexes one constrained matrix, so W_t^ℓ is that matrix at step t and plays the role of W_{pre} above. The optimizer itself is never modified. It receives the full gradient, updates its moments, and proposes an ordinary step; SST then replaces the part of that step which lands on a constrained matrix with its free-to-free block. Parameters outside $\mathcal{I}$, such as embeddings and layer norms, are left to the optimizer.

B.9 Comparison with subspace-constrained baselines

The closest baselines are Sculpting [34] and OPLoRA [54]. Both restrict updates to low-energy directions through a trainable factorisation; SST instead projects the applied full-parameter update. We use the official

Algorithm 1 Subspace Tuning (SST). The optimizer runs untouched; SST edits only the update it produces, keeping the free-to-free block of Eq. 3 and discarding the rest.

Require: Constrained matrices $\mathcal{I}$ (attention q, k, v, o and MLP gate, up, down) with pretrained weights W^ℓ of shape $d^\ell_{\text{out}} \times d^\ell_{\text{in}}$; rank fraction r, the share of singular directions held fixed; refresh interval K in optimizer steps; optimizer $\mathcal{O}$ with state s; loss $\mathcal{L}$ with minibatches $\mathcal{B}_t$; step budget T.

Setup: protect the leading singular subspace of all matrices in $\mathcal{I}$.

for all $\ell \in \mathcal{I}$ **do**
  $W^\ell_{\text{ref}} \leftarrow W^\ell$
  $k_\ell \leftarrow \lfloor r \min(d^\ell_{\text{out}}, d^\ell_{\text{in}}) \rfloor$ ▷ number of directions to protect
  $U^\ell_k, V^\ell_k \leftarrow$ top-k_ℓ left and right singular vectors of W^ℓ_{ref}
  $Q^\ell_U \leftarrow I - U^\ell_k (U^\ell_k)^\top, \quad Q^\ell_V \leftarrow I - V^\ell_k (V^\ell_k)^\top$ ▷ the free complements
end for

Training: take the ordinary step, then project the part that reaches the weights.

for $t = 0, \ldots, T-1$ **do**
  $g_t \leftarrow \nabla_W \mathcal{L}(W_t; \mathcal{B}_t)$ ▷ full gradient
  $(\widetilde{W}_{t+1}, s_{t+1}) \leftarrow \mathcal{O}(W_t, g_t, s_t)$ ▷ ordinary step
  $W_{t+1} \leftarrow \widetilde{W}_{t+1}$
  for all $\ell \in \mathcal{I}$ **do**
    $\Delta W^\ell \leftarrow \widetilde{W}^\ell_{t+1} - W^\ell_t$ ▷ the optimizer update
    $W^\ell_{t+1} \leftarrow W^\ell_t + Q^\ell_U \, \Delta W^\ell \, Q^\ell_V$ ▷ apply $\Pi(\Delta W^\ell)$ instead, Eq. 3
  end for
  if $K < \infty$ and $(t+1) \bmod K = 0$ **then**
    for all $\ell \in \mathcal{I}$ **do**
      $W^\ell_{\text{ref}} \leftarrow W^\ell_{t+1}$ ▷ re-compute projection with current weights
      $U^\ell_k, V^\ell_k \leftarrow$ top-k_ℓ singular vectors of W^ℓ_{ref}
      $Q^\ell_U \leftarrow I - U^\ell_k (U^\ell_k)^\top, \quad Q^\ell_V \leftarrow I - V^\ell_k (V^\ell_k)^\top$
    end for
  end if
end for
return W_T

hyperparameters reported in each baseline's paper and otherwise keep the data, training budget, and evaluation harness fixed.

SST preserves more capability than factorised subspace baselines. Table 5 shows that Sculpting and OPLoRA reduce forgetting relative to SFT, but remain below no-refresh SST on prior capability across the four-domain comparison. The gap is modest on Science and becomes clearer on Finance and the medical domains. Refresh shifts SST toward stronger task learning at some cost to retention. The table reports the exact domain-level endpoints.

Data scaling across Medicine, Pathology, and FinQA. Figure 15 combines the repeated-data Medicine and Pathology sweeps with the matched-exposure FinQA comparison. Medicine and Pathology each repeat $9{,}000$ examples for five epochs ($45{,}000$ examples total), with evaluation every $9{,}000$ examples. FinQA compares $2{,}000$ unique examples repeated for three epochs against $6{,}000$ unique examples seen once; both arms therefore contain $6{,}000$ examples of exposure and use the same held-out $1{,}000$ examples.

SST retains capability while continuing to learn. Both SST variants preserve substantially more prior capability than Sculpting and OPLoRA on Medicine and Pathology. On FinQA, no-refresh SST remains closest to the base model, while R10 reaches higher target accuracy in both exposure regimes. SFT is omitted so the constrained trajectories remain legible.

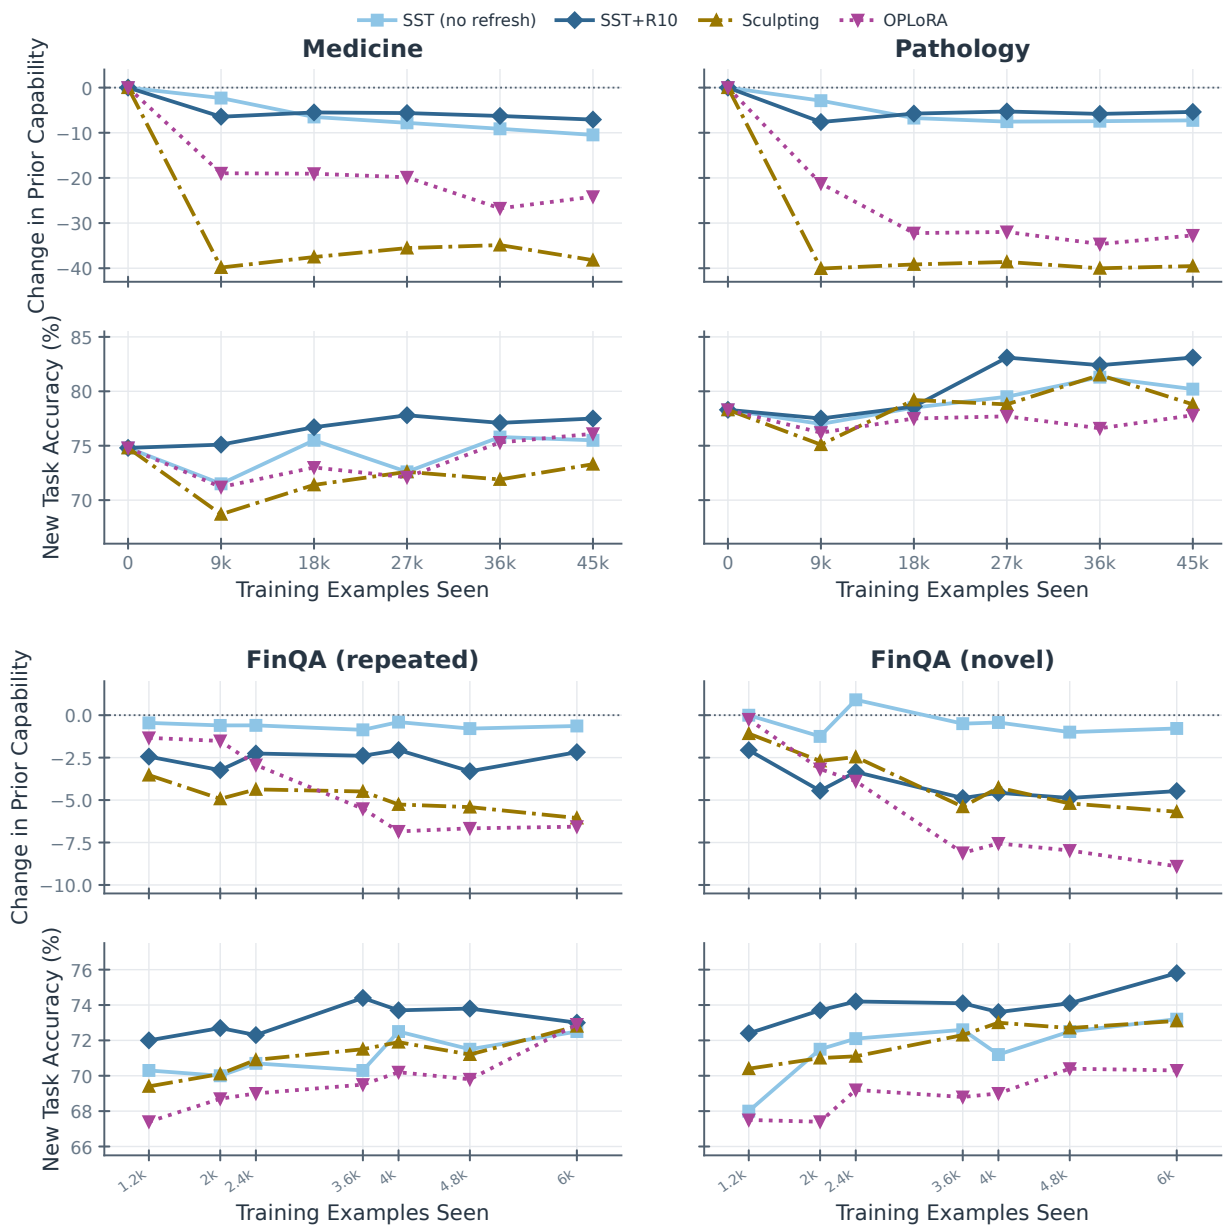

Figure 15: **Data scaling against subspace-constrained baselines on Qwen3-8B.** The top two rows show Medicine (left) and Pathology (right), each using 9,000 unique examples repeated for five epochs. The bottom two rows show FinQA with matched 6,000-example exposure: 2k unique traces repeated for three epochs (left) or 6k unique traces seen once (right). Within each pair, the upper panels show change in the five-benchmark prior-capability average from the base value of 81.0, and the lower panels show target accuracy. Methods share the data and evaluation schedule within each setting. Curves compare no-refresh SST, SST+R10, Sculpting, and OPLoRA; SFT is omitted to keep the constrained trajectories legible.

Table 5: Comparison with subspace-constrained baselines using the official method-specific settings from their papers. All methods use the same datasets and evaluation suite. ID is in-domain accuracy and LM is the five-benchmark capability average (base 81.0). Bold entries mark column maxima among non-base recipes.

	Science		FinQA		Pathology		Medicine		Mean	
Method	ID	LM	ID	LM	ID	LM	ID	LM	ID	LM
Baseline	60.2	81.0	70.5	81.0	78.3	81.0	74.8	81.0	71.0	81.0
SFT	77.6	55.5	72.3	23.7	80.4	34.7	72.9	36.3	75.8	37.6
Sculpting	83.2	77.2	74.5	57.7	79.4	43.5	74.6	41.3	77.9	54.9
OPLoRA	78.5	79.1	72.8	61.7	77.4	51.6	70.0	60.0	74.7	63.1
SFT+SST	80.1	80.0	73.0	**82.3**	77.3	77.2	72.4	**78.9**	75.7	**79.6**
SFT+SST+R100	81.7	80.0	73.1	80.1	77.6	**78.5**	74.3	78.3	76.7	79.2
SFT+SST+R10	**84.0**	**80.2**	**75.3**	75.7	**82.1**	72.3	**76.6**	71.2	**79.5**	74.9

B.10 KL regularisation does not prevent forgetting

Penalising the policy does not mitigate forgetting. The natural alternative to constraining *where* an update may move is to leave the update unconstrained and instead penalise how far the resulting policy drifts. We fine-tune Qwen3-8B on FinQA and on Science with $\mathcal{L} = \mathrm{CE} + \lambda\,\mathrm{KL}(\pi_\theta \,\|\, \pi_{\mathrm{base}})$, a reverse KL over the response tokens against a frozen copy of the base model, at $\lambda \in \{0.001, 0.01, 0.1\}$, plus a λ=0 control on FinQA. Table 6 and Figure 16 show that reverse-KL regularisation does not provide a reliable learning–retention trade-off. Weak regularisation leaves forgetting essentially unchanged, while stronger regularisation can reduce target-task learning without restoring prior capability. The effect also varies across tasks, unlike the consistent retention obtained by constraining the update subspace. The table reports the exact per-task and per-benchmark values.

Table 6: **SFT with reverse-KL regularisation, Qwen3-8B.** Prior capability is the mean of the five held-out benchmarks shown. Bold entries mark column maxima among non-base recipes across the reported rows.

Task	Method	ID	GPQA	MMLU-Pro	GSM8K	HumanEval	IFEval	Prior cap.
—	Base		58.1	59.7	94.4	92.1	85.8	78.0
FinQA	λ=0	73.4	10.1	13.4	41.8	24.4	28.1	23.6
	λ=0.001	73.4	7.6	14.2	43.5	24.4	26.4	23.2
	λ=0.01	71.3	7.1	13.2	42.0	17.7	26.6	21.3
	λ=0.1	66.0	13.6	15.8	14.2	0.0	9.4	10.6
Science	λ=0.001	**81.7**	**49.5**	54.4	93.3	60.4	**55.1**	62.5
	λ=0.01	79.7	46.5	**55.0**	**93.5**	50.6	**55.1**	60.1
	λ=0.1	75.9	47.5	51.2	93.4	**77.4**	51.9	**64.3**

B.11 Rank ablation on FinQA

The constraint maintains prior capability. FinQA is the hardest case in Section B.10: every KL coefficient leaves prior capability at or below plain SFT. We repeat that setting with SST instead, sweeping the protected rank fraction $r \in \{0.1, 0.5, 0.9\}$ with no refresh. Figure 17 reports the sweep against the same base control and the same λ=0 SFT run.

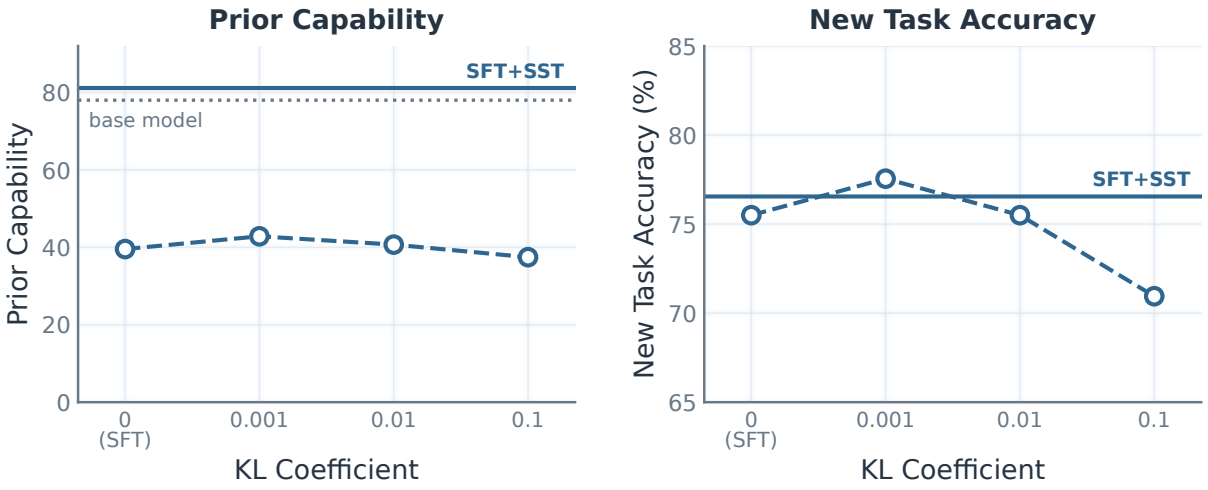

Figure 16: **Reverse-KL regularisation across three coefficients.** Prior capability (left) and in-domain accuracy (right) against SFT with KL regularization, averaged over FinQA and Science. The grey rule is the un-finetuned base model and the solid blue rule is SST at r=0.99 on the same two tasks.

Rank controls retention. Figure 17 shows prior capability improving monotonically as a larger fraction of the weight subspace is protected. The held-out benchmarks move together, indicating a broad recovery rather than an average driven by one task.

Protecting more only impacts in-domain accuracy for high ranks. In-domain accuracy also improves across the sweep until $r = 0.99$, where the new task accuracy declines. Figure 17 shows the trajectory. Although a lower r permits adaptation along more directions, supervision with a dated projection matrix can disrupt general model behavior (e.g., collapse response length), ultimately hindering learning of the target distribution.

B.12 Effect of Optimizer Hyperparameters on SST

We test sensitivity to two optimizer hyperparameters on FinQA: the learning rate and its schedule. We cross three learning rates (10^{-4}, 5×10^{-5}, and 10^{-5}) with flat and cosine schedules, holding the data, training budget, and evaluation fixed. Figure 18 shows all six configurations for SFT, SST without refresh, and SST with refresh every 10 steps (R10), together with their minima, maxima, means, and population variances.

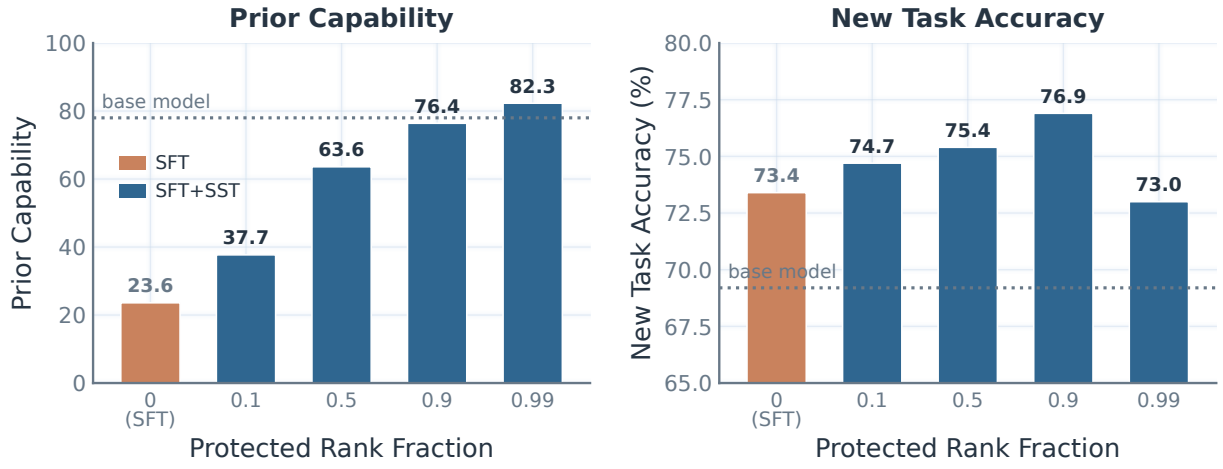

Figure 17: **Protected rank fraction on FinQA.** Prior capability (left) and in-domain accuracy (right) against the fraction of each weight's subspace held fixed, with plain SFT as the $r{=}0$ end of the same sweep. The grey rule is the un-finetuned base model. Where no KL coefficient recovers prior capability (Figure 16), rank recovers it almost completely and in-domain accuracy rises rather than falls.

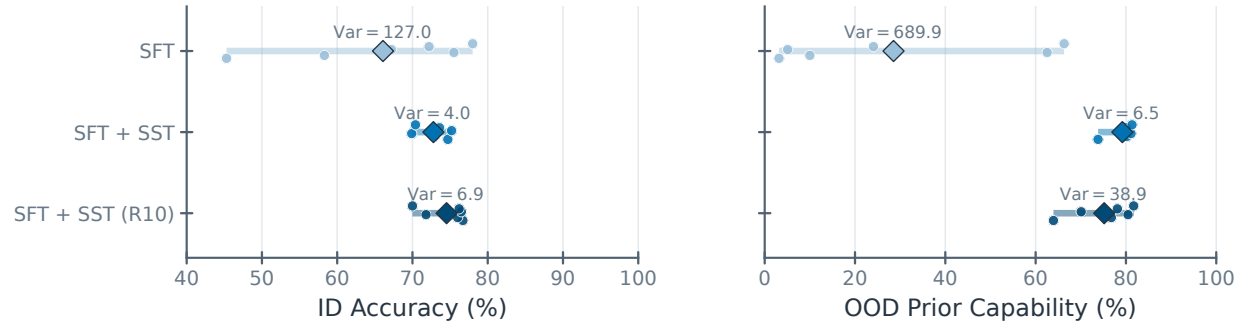

Figure 18: **FinQA sensitivity to learning rate and schedule.** Each dot is one of six configurations, the thick segment spans the minimum and maximum, and the diamond marks the mean. SST greatly reduces variation in both in-domain accuracy and the out-of-domain (OOD) prior-capability average.

SST is substantially more robust to these choices than unconstrained SFT. Both SST variants exceed SFT in-domain accuracy in four of the six matched configurations while retaining prior capabilities. At the lowest learning rate, however, SFT attains the highest individual in-domain scores, so this sweep does not imply universal in-domain dominance. Rather, SST consistently improves retention and generally yields a more stable joint learning–retention outcome; R10 shifts the balance toward higher average in-domain accuracy at some cost to retention robustness.

B.13 Computational Cost

In this section, we consider the computation cost of subspace-constrained finetuning. SST adds no extra trainable parameters or inference-time cost. The method projects each optimizer update and optionally refreshes the protected basis. Refresh uses Gram eigendecompositions, using the smaller Gram matrix

for rectangular weights. The multi-GPU training setup uses DeepSpeed ZeRO-3 [38], and the estimates distribute the independent matrix operations across eight GPUs.

We report the measured SFT step from full-parameter bfloat16 Qwen3-8B training on eight NVIDIA L40S GPUs, with effective batch size 16, and maximum sequence length 4096. The measured SFT step took 1.73 s. Across the model's 36 layers and 252 projected matrices, projection adds 26.2 ms (1.5%) per step in all settings. Refresh adds an amortized 62.4 ms per step for R100 and 623.5 ms for R10, giving total overheads of 5.1% and 37.6%, respectively. Splitting the SVD computation across the 8 GPUs brings the overhead to 0.6% and 4.7% respectively.

B.14 Training and evaluation details

Training. Unless set by an official baseline recipe in Appendix B.9, the Qwen3 and Olmo studies use learning rate 5×10^{-5} and effective batch size 16.. These settings are shared across paired recipe/SST comparisons and were not selected separately for each method or domain. Training is full-parameter except for the factored Sculpting and OPLoRA baselines and uses DeepSpeed ZeRO-3 [38].

Datasets and synthetic traces. The Qwen2.5 instruct study uses ToolAlpaca [47], HuatuoGPT-o1 [10], and SciKnowEval [19]. Qwen3 Pathology and Medicine use MedMCQA questions and answer labels [37]. For each question, the generator receives the answer options followed by `Choose the correct answer (A, B, C, or D) and explain your reasoning.` We use the Bedrock model identifier `us.anthropic.claude-sonnet-4-6`; its response is serialized inside `<think>...</think>` with a boxed answer. The archived sets contain 8,996 Medicine traces and 8,999 Pathology traces, all with a nonempty thinking block and boxed answer.

Evaluation. Evaluation uses vLLM [25]. Prior capability is the average of GPQA-Diamond, MMLU-Pro, GSM8K, HumanEval, and IFEval; in-domain accuracy uses the task's native answer extractor.

For the Olmo-3 medical study, open-ended accuracy is graded with the Bedrock model identifier `us.anthropic.claude-sonnet-4-6`. The judge receives the question, reference answer, and full model response and returns `CORRECT` or `INCORRECT`. The fabrication metric is a deterministic lower bound that flags explicit claims of an absent source, including "reference solution," "provided answer," and "based on the feedback." None of the 1,000 evaluation prompts contains a trigger phrase.