When Is a Relation Worth Trusting? Manipulation Costs, Exposure, and the Trust Trap in Multi-Relational Graph Learning

Vincess Dongmo [*1] **Astride Melvin Fokam Ninyim** [*1 2]

arXiv:2610.37239v1 [cs.LG] 27 Oct 2026

Abstract

Multi-relational graph detectors learn from clean data how much to trust each relation, but an adversary pays a different price to manipulate each relation and each action; taking such costs as given, we ask when using a relation is worse than ignoring it. Our first main result proves that degree gating strictly beats every selection or reweighting of a relation when part of the population is isolated on it and the attacker cannot afford to open the gate, and that this window is sharp; on real graphs, validated gating gains exactly on the relations that meet these conditions and nowhere else. Our second shows that a per-node exposure bound, which extends to message passing, tells a platform which relation to harden: exposure per unit cost, computed without any attack, picks the relation whose hardening raises the exact trap threshold most in 15 of 16 cases, against at most 7 for informativeness, cheapness, or an analysis with equal edit costs. Supporting tools give exact worst cases for linear relational models, certified brackets for R-GCN and tree ensembles, and a finite-sample guarantee for adversarial validation, under which clean-trained graph detectors fall below their graph-free counterparts in 58 of 90 settings on six graphs. A selection protocol built on these results is never significantly worse than ignoring the graph; with at most four relations it mostly abstains, as validation alone would, and its cost-aware proposals help when relations are many, so far shown on synthetic graphs. Without cost information, the safe choice is to abstain.

1. Introduction

Fraud, collusion, and spam detectors increasingly learn from multi-relational graphs: reviews that share a user or a product, firms that share an owner or a supplier, accounts that share a device. Such models learn from clean data how much to trust each relation, through relation-specific weights (Schlichtkrull et al., 2018), semantic attention (Wang et al., 2019), neighbor selection (Dou et al., 2020; Liu et al., 2021), or spectral filters (Tang et al., 2022; Wu et al., 2023). That such detectors can be manipulated is well known: fraudsters camouflage their relations (Dou et al., 2020; Zhang et al., 2024), structural attacks degrade graph models (Zügner et al., 2018; Geisler et al., 2021), gangs coordinate injections (Choi et al., 2025), and cost-aware attacks price the nodes of homogeneous graphs (Liao et al., 2024).

What is missing is a theory of how the *economics* of manipulation interact with the *information* each relation carries. An adversary does not pay the same price for every relation or every action: reviewing a popular product creates an edge almost for free, reusing an aged account is costly, and concealing a tie to an accomplice may be cheap or expensive depending on identity verification. A relation can therefore be statistically valuable and economically unsafe. Neighbor filters ask which neighbors to keep within a relation, cost-aware defenses ask which nodes are at risk, and robust optimization asks how to train under attack. We ask a prior question: *should a relation be used at all, given how informative and how manipulable it is?*

Contributions. The paper has two main results; the rest are tools and supporting results (Table 7 compares each with the workshop version and the closest prior work). **(i) Degree gating strictly beats selection and reweighting** of a relation on which part of the population is isolated, as long as the attacker cannot open the gate, and the window is sharp (Theorem 1); on real graphs, validated gating gains on exactly the relations that meet these conditions (Table 18). **(ii) Which relations are dangerous, and which to harden.** A per-node exposure bound that extends to message passing (Propositions 1 and 2) and the budget beyond which no detector can use the graph (Theorem 2) say when a relation's signal is outweighed by its manipulation cost.

[*]Equal contribution [1]SYNTHI-AI, Cameroon [2]Université du Québec à Chicoutimi (UQAC), Canada. Correspondence to: Vincess Dongmo <vincess.dongmo@synthi-ai.com>, Astride Melvin Fokam Ninyim <amfninyim@etu.uqac.ca>.

Preprint. October 28, 2026.

Combined with exact thresholds, they answer a question that robustness analyses with equal edit costs cannot pose: which relation a platform should make harder to manipulate. The first-order exposure per unit cost, computed without any attack, picks the relation whose hardening raises the exact trap threshold most in 15 of 16 cases, against at most 7 for informativeness, cheapness, weight size, or an exact analysis under equal costs (Table 25). **Tools: exact, certified, and validated evaluation.** Exact worst cases and thresholds for linear relational models under heterogeneous costs (Proposition 3), certified brackets for R-GCN and tree ensembles that we label tight, informative, or loose (Proposition 4), and a finite-sample guarantee for adversarial validation whose radius grows with the number of candidates checked (Proposition 5). *Supporting results*: the threshold's structure (Lemma 1), robust planning, maximin choice, and hardening (Corollary 1). *Protocol*: CARS-D lets the theory propose three candidates and adversarial validation pick one.

What the evidence supports. (1) Clean-trained relational detectors are fragile under cheap manipulations: exactly for linear models, through evaluated upper bounds for nonlinear ones. (2) Gating helps exactly where the theorem says, and nowhere else. (3) On real graphs with one to four relations, CARS-D is never significantly worse than ignoring the graph and mostly abstains, as validation alone would; its cost-aware proposals help with many relations, so far on synthetic graphs (Table 23). All of this is conditional on specified costs: the paper studies selection under given cost scenarios, not real manipulation costs (Section C). A workshop paper (Vincess & Melvin, 2026) introduced the threat model and studied the phenomenon on the CSBM, YelpChi, and Amazon; the exposure bound (Proposition 1), the linear-response part of Theorem 2, and the cost-monotonicity part of Corollary 1 appeared there. Section C (Table 7) lists what is new.

2. Problem Setup

Let $G = (V, \{E_r\}_{r=1}^R, X)$ be a graph with R relations, node features X, and labels $y_v \in \{0, 1\}$ (1 for fraud); $\mathcal{N}_r(v)$ denotes the neighbors of v under r. A detector f is trained on the clean graph. At test time, each fraudulent node v edits its own edges (evasion): on relation r it may inject $b_{v,r}$ edges at unit cost c_r^+ or delete $k_{v,r}$ edges to accomplices at unit cost c_r^-. Writing $\theta_v = (b_v, k_v) \in \mathbb{R}^{2R}_{\geq 0}$ and $c = (c^+, c^-) \in \mathbb{R}^{2R}_{>0}$, the per-node feasible set is

$$\Theta_B(c) = \big\{\theta \geq 0 :\ c^\top \theta \leq B,\ k_r \leq K_r\big\}, \qquad (1)$$

with K_r the number of accomplices on r. A *gang* instead shares a total budget $\bar{B}|T|$ over its members T. When relations are projections of entities (two reviews share a user), we also consider *realizable* actions that reassign an entity attribute and rebuild all induced edges at once (Section 5). Costs encode relative manipulation difficulty rather than money; since $\Theta_{\alpha B}(\alpha c) = \Theta_B(c)$, only cost-to-budget ratios matter.

Definition 1 (Trust trap)**.** Let f_0 be the graph-free counterpart of f, the same learner on node features only. With Δ_{clean} the clean performance gap of f over f_0 and $\Delta_{\text{rob}}(B, c)$ the gap under the worst attack in $\Theta_B(c)$, f is in a *trust trap* at (B, c) if $\Delta_{\text{clean}} > 0 \geq \Delta_{\text{rob}}(B, c)$. Its *trap threshold* is $B_{\text{trap}}(f; c) = \inf\{B : \Delta_{\text{rob}}(B, c) \leq 0\}$.

3. Theory

Levels of guarantee. We separate three kinds of statements. *Exact* results hold without approximation inside a stated model: Lemma 1 for any detector, Theorem 2(a) in the CSBM, and Proposition 3 for linear relational models. *Bounds* hold for broader classes but may be loose: Propositions 1, 2 and 7, the certified AUC of Proposition 4, and the statistical guarantee of Proposition 5. *Evaluated* worst cases, reported for nonlinear detectors in Section 5, come from the strongest attacks we know and upper-bound the true worst case; Proposition 4 brackets it from below.

3.1. The structure of the trap threshold

The first result is elementary and needs no model of the data; we state it because the experiments and the robust-planning results rely on its exact form, including where it does *not* give concavity. An attack policy π assigns each fraudulent node v an edit vector $\theta_v(\pi) \in \mathbb{Z}^{2R}_{\geq 0}$ (counts of injections and deletions per relation, with $k \leq K$); performance is measured after all edits are applied. Let $\mathcal{T}(f)$ be the set of *trap policies*, under which f performs at most as well as f_0. A policy is affordable under per-node budgets if $\max_v c^\top \theta_v(\pi) \leq B$, and under a gang budget if $\frac{1}{|T|}\sum_v c^\top \theta_v(\pi) \leq \bar{B}$.

Lemma 1 (Structure)**.** *For every detector f and performance measure, (a) $B_{\text{trap}}(f; c) = \min_{\pi \in \mathcal{T}(f)} \max_v c^\top \theta_v(\pi)$ and the gang threshold is $\bar{B}_{\text{trap}}(f; c) = \min_{\pi \in \mathcal{T}(f)} \frac{1}{|T|}\sum_v c^\top \theta_v(\pi) \leq B_{\text{trap}}(f; c)$, both minima being attained. (b) Both thresholds are nondecreasing and positively homogeneous of degree one in c. (c) $\bar{B}_{\text{trap}}$ is concave and piecewise linear in c, and so is B_{trap} whenever performance depends on a policy only through a common edit vector, as in the CSBM; in general B_{trap} need not be concave.*

The gang threshold and the CSBM threshold are lower envelopes of lines, one per attack strategy: they grow linearly in a relation's cost while the cheapest trap attack uses that relation and saturate once another strategy becomes cheaper. Under heterogeneous per-node budgets, a max over nodes enters, and the threshold can be convex (two fraud nodes

that can only be trapped through different relations give $B_{\text{trap}} = \max(c_1, c_2)$). Monotonicity and homogeneity hold in every case, and they suffice for robust planning.

Corollary 1 (Robust planning, maximin choice, and hardening)**.** *(a) If $c' \geq c$ componentwise, then $\Theta_B(c') \subseteq \Theta_B(c)$, so every detector's worst-case performance under c' is at least its worst-case performance under c. Since $\Theta_{\rho B}(c) = \Theta_B(c/\rho)$, a detector whose worst case is A at planning costs c/ρ and budget B keeps worst case at least A for every true cost vector $c' \geq c/\rho$ and, at costs c, for every budget up to ρB. For a gang, the smallest threshold over any polytope of cost vectors is attained at one of its vertices. (b) If the true costs are only known to lie in a set U with a componentwise-minimal element $\underline{c}$ (a box, for instance), then $\min_{c \in U} A_B(f; c) = A_B(f; \underline{c})$ for every detector f, so the maximin choice over U among any set of detectors is the one with the best worst case at $\underline{c}$. (c) If some cheapest trap policy at c uses no action on relation r, raising the costs of r leaves $B_{\text{trap}}(f; c)$ unchanged: hardening a relation raises the threshold only if every cheapest trap policy uses it.*

Overestimating costs can only overestimate robustness; planning at the lower corner of a cost box certifies the whole box, and the price of robustness is the performance lost at the nominal costs.

A multi-relational CSBM. To make the envelope explicit, we extend the contextual stochastic block model (Deshpande et al., 2018; Baranwal et al., 2021; Ma et al., 2022). A node's projected feature is $x_v \sim \mathcal{N}(\pm m, 1)$ (+ for fraud). Under relation r, a fraudulent node has d_r neighbors, a fraction p_r of them fraudulent, and a benign node has d_r^0 neighbors, a fraction $q_r \leq p_r$ fraudulent, independently across relations. A linear relational detector scores $z_v = w_0 x_v + \sum_r w_r a_{v,r}$, with $a_{v,r}$ the mean of x_u over $\mathcal{N}_r(v)$; injected neighbors have mean $-(m + \delta)$, $\delta \geq 0$. Under an attack θ, the class moments are explicit (Section A), and with $\Delta(\theta) = \mu_1(\theta) - \mu_0$,

$$\text{AUC}(w; \theta) = \Phi(\zeta(w; \theta)), \quad \zeta = \frac{w^\top \Delta(\theta)}{\sqrt{w^\top (\Sigma_1(\theta) + \Sigma_0) w}}. \tag{2}$$

The clean-optimal (Fisher) weights are $w^\star \propto (\Sigma_1(0) + \Sigma_0)^{-1} \Delta(0)$, and the graph-free detector e_0 has $\zeta_0 = \sqrt{2}m$ under every attack.

In the base configuration of Section 5 (an identity, a behavior, and a dense relation), the cheapest trap attack against $w^\star$ deletes all $p_{\text{id}} d_{\text{id}} = 3.2$ identity accomplices and injects 0.72 behavior edges, so that $B_{\text{trap}} = 3.2\, c_{\text{id}} + 1.45$ until a camouflage of the behavior and dense relations becomes cheaper, at 30.4 (Figure 1a). The trap is bought by severing ties: when the absence of ties is the signal, deletion is the cheapest manipulation.

3.2. Exposure: a lower bound

Proposition 1 (Exposure)**.** *Let $L_j(B) = \sup_{\theta \in \Theta_B(c)} (-\partial \zeta / \partial \theta_j)_+$. Then $B_{\text{trap}}(w; c) \geq \sup\{B : B \max_j L_j(B)/c_j < \zeta(w; 0) - \zeta_0\}$. At $\theta = 0$, the injection exposure in relation r is $w_r(2p_r m + \delta)/(s(w)\, c_r^+ d_r) + O(w_r^2 d_r^{-2})$, with $s(w)^2 = w^\top (\Sigma_1(0) + \Sigma_0) w$.*

To first order, the adversary spends its whole budget on the action of largest exposure, reliance w_r over effective cost $c_r^+ d_r$; a cheap relation is safe only if it is dense, and a dense relation is useful only if it is informative, so danger peaks at intermediate degrees. The bound extends beyond linear detectors.

Proposition 2 (Exposure of message passing)**.** *Let $h_v = \sigma(W_0 x_v + \sum_r W_r \bar{x}_{v,r})$ with $\bar{x}_{v,r}$ the mean of x_u over $\mathcal{N}_r(v)$, and let the score $s_v = g(h_v)$ be Λ-Lipschitz in the pre-activation. An attack with b_r injections and k_r deletions changes s_v by at most $\Lambda \sum_r \|W_r\|_2 \Delta_{v,r}$, where $\Delta_{v,r} = (b_r + k_r) D_{v,r} / (d_r(v) - k_r + b_r)$ if r keeps at least one neighbor of v, $D_{v,r}$ bounding $\|x_u - \bar{x}_{v,r}\|$ over the edited endpoints (with $\bar{x}_{v,r} = 0$ for an isolated node), and $\Delta_{v,r} = \|\bar{x}_{v,r}\|$ if the attack empties r (the empty aggregate is 0). For one unit of budget, the exposure of relation r is at most $\Lambda \|W_r\|_2 D_{v,r} / (c_r^+ (d_r(v) + 1))$ for injections, and $\Lambda \|W_r\|_2 D_{v,r} / (c_r^- (d_r(v) - 1))$ for deletions if $d_r(v) \geq 2$, or $\Lambda \|W_r\|_2 \|\bar{x}_{v,r}\| / c_r^-$ if $d_r(v) = 1$.*

Reliance becomes the relation's operator norm and degree enters exactly as in the linear model; stacking layers multiplies the Lipschitz constants. The bound is per node, and low-degree nodes are the weak spot: a node without neighbors on r has exposure $\Lambda \|W_r\|_2 D_{v,r} / c_r$, since one injection replaces its empty aggregate. Gating a relation by degree removes this weak spot.

A degree gate lets relation r enter h_v only if $d_r(v) \geq \tau_r$; it replaces $d_r(v)$ by at least τ_r in the bound of Proposition 2 for every open node, and an attacker must pay $c_r^-(d_r(v) - \tau_r + 1)$ to close a gate and $c_r^+(\tau_r - d_r(v))$ to open one (Proposition 7 in Section A). Gating trades coverage for exposure; the next result shows that, when part of the population is isolated on a relation, which is the regime of FDCompCN and YelpChi, gating is strictly better than anything relation selection or reweighting can do.

Theorem 1 (Gating strictly beats selection and reweighting)**.** *In the CSBM with one relation r, let a fraction $\pi_0 \in (0, 1)$ of the nodes of each class be isolated on r and the others have degree d, with $p > q$. Let the attacker afford $b = \lfloor B / c^+ \rfloor$ injections per node, with $1 \leq b < \bar{b}_d := d\,(\mu_f - \mu_b)/(\mu_b + m + \delta)$, where $\mu_f = (2p - 1)m$ and $\mu_b = (2q - 1)m$, and no gate-closing deletion ($B < c^-(d - \tau + 1)$). Consider linear detectors $s_v = w_0 x_v + w\, \bar{x}_{v,r} \mathbb{1}[d_r(v) \geq \tau]$ with $w \geq 0$. For every gate τ with $b < \tau \leq d$: (a) for every $w > 0$, the*

worst-case AUC of the gated detector is strictly larger than that of the ungated one ($\tau = 0$) with the same weights; (b) the derivative of the gated worst-case AUC in w at $w = 0$ is proportional to $(1 - \pi_0)(\mu'_d - \mu_b) > 0$, with $\mu'_d = (d\mu_f - b(m + \delta))/(d + b)$, so some $w > 0$ makes the gated detector strictly better than dropping the relation ($w = 0$). Hence $\sup_w A_{\text{gated}} > \max\{\sup_{w\geq 0} A_{\text{ungated}}, A_{\text{drop}}\}$. (c) The window is sharp: if $b \geq \tau$, the gated and ungated detectors with the same weights have the same worst-case AUC, and if $\tau > d$, the gated detector is never better than dropping the relation.

Selection keeps or drops a relation for every node, and reweighting scales it for every node; neither can separate isolated nodes, on which one injection replaces an empty aggregate, from connected nodes, on which it is diluted by d neighbors. A gate that the attacker cannot open within budget does. The closed-form check in Table 16 finds the best ungated weight to be $w = 0$ (dropping the relation) for every isolated fraction $\pi_0 \geq 0.5$ in our grid, while gating keeps a strict gain; exact attacks on CSBM graphs confirm it for RLR, and only where $b < \tau$ (Table 17); by part (c), the window cannot be widened. On real graphs, validated gating beats validated reweighting and dropping on exactly the relations where the conditions hold, informative relations on which part of the fraud is isolated, and only while the gate cannot be opened: YelpChi's R-U-R (82% of fraud isolated; up to +.021 AUC) and S-FFSD's T-S-T (26%; up to +.061), 2 of the 17 relations tested; none of its 14 wins occurs once the gate can be opened, and validation errors cost up to −.022 elsewhere (Table 18). The theorem explains a mechanism; it does not promise a gain where no relation is both informative and partly isolated. Section 4 chooses τ_r from data.

3.3. Abandoning the graph and robust weights

Let $\bar{b}_r = \frac{2m(p_r - q_r)d_r}{2q_r m + \delta}$ and $\bar{k}_r = \frac{(p_r - q_r)d_r}{1 - q_r}$ be the numbers of injections and deletions that neutralize relation r ($\Delta_r = 0$), $\kappa_r = \min(c_r^+ \bar{b}_r, c_r^- \bar{k}_r)$ its *neutralization cost*, $\zeta_r = \Delta_r(0)/\sigma_r$ its clean signal-to-noise ratio, and $\zeta_\star^2 = \zeta_0^2 + \sum_r \zeta_r^2$.

Theorem 2 (Abandonment and robust weights)**.** *(a) In the CSBM, no detector beats the graph-free one under attack once $B \geq \sum_r \kappa_r$: $B_{\text{ab}} := \sup_w B_{\text{trap}}(w; c) \leq \sum_r \kappa_r$. (b) If each action lowers Δ_r at the constant rate that neutralizes r and leaves variances unchanged (linear response), then $B_{\text{ab}} = \sum_r \kappa_r$, $B_{\text{trap}}(w^\star) = \zeta_\star(\zeta_\star - \zeta_0) \min_{r:p_r > q_r} \kappa_r/\zeta_r^2$, and the minimax detector clips the Fisher weights at caps proportional to neutralization cost: $w_r = \min(\lambda\Delta_r/\sigma_r^2,\ t\kappa_r/\Delta_r)$ for some $\lambda, t \geq 0$.*

Part (a) is exact: the neutralization set of each relation is a half-space in (b_r, k_r), so its cheapest point is a vertex. The gap between the trap of the clean-optimal detector, governed by its weakest link, and abandonment, which requires neutralizing every relation, is the price of naive trust: in between, structure still has value, but only for a detector that discounts each relation by its exposure. Numerically, the abandonment threshold lies within 9% of $\sum_r \kappa_r$ at every cost (Figure 1a). Gangs that share a budget drive the AUC down to the lower convex envelope of the per-node attack curve (Proposition 6 in Section A).

3.4. Exact worst case of linear relational models

Proposition 3 (Exact attack)**.** *For relational logistic regression $s_v = w_0^\top x_v + \sum_r w_r^\top \bar{x}_{v,r}$ with costs on a grid of step ϵ, the target-side worst case under* (1) *is computed exactly in $O(R\,(B/\epsilon)^2 + \sum_r b_r^{\max} k_r^{\max})$ per target: for fixed counts, the optimal edits inject the candidates with the smallest projections $w_r^\top x_u$ and delete the accomplices with the largest, and a dynamic program allocates the budget across relations.*

Because the dynamic program returns the worst score of every target for every budget on the grid at once, the separable worst-case AUC curve $B \mapsto A(B)$ of a linear relational model, gated or not, and hence its trap threshold, are computed exactly in one pass per target; relation subsets give the per-relation and joint neutralization budgets on real graphs.

3.5. Certified evaluation and validated selection

For nonlinear detectors, an attack only shows that the true worst case is *at most* what it achieves. A lower bound needs per-node certificates, and a guarantee for any selection rule needs to account for the validation sample.

Proposition 4 (Certified AUC)**.** *Under the target-side model (each fraudulent node edits its own ties; benign scores fixed), let $\ell_v \leq \min_{\theta\in\Theta_B(c)} s_v(\theta)$ for every fraudulent test node. Then the true worst-case AUC satisfies $A_{\text{cert}} \leq A_{\text{true}}(B) \leq A_{\text{eval}}$, where $A_{\text{cert}} = \frac{1}{n_+ n_-} \sum_{v,w} \mathbb{1}[\ell_v > s_w]$ and A_{eval} is the AUC under any evaluated attack. In particular, if $|s_v(\theta) - s_v(0)| \leq \Delta_v$ on $\Theta_B(c)$ (Proposition 2), the AUC drops by at most the fraction of fraud–benign pairs whose clean margin $s_v - s_w$ is at most Δ_v. For mean aggregation, each coordinate of $W_r\bar{x}'_{v,r}$ after b injections and k deletions lies in $[(S_j - \text{top}_k(\text{acc})_j + \text{bot}_b(\mathcal{C})_j)/(d - k + b),\ (S_j - \text{bot}_k(\text{acc})_j + \text{top}_b(\mathcal{C})_j)/(d - k + b)]$, with $S = \sum_{u\in\mathcal{N}_r(v)} W_r x_u$ and the top and bottom sums taken per coordinate over accomplices and candidates; if the attack empties the relation ($d - k + b = 0$), the aggregate is exactly 0; these intervals propagate through R-GCN layers and through tree ensembles.*

For linear relational models $A_{\text{cert}} = A_{\text{true}}$ by Proposition 3; for R-GCN and gradient-boosted trees the bracket can be wide, and its width measures how much of a robustness claim rests on the attacks.

Proposition 5 (Validated selection)**.** *Let a rule pick, among K fixed candidate detectors, the one with the highest evaluated worst-case AUC on n_+ fraudulent and n_- benign validation nodes drawn i.i.d. from the evaluated population, under a target-side evaluation in which each node's attacked score depends only on that node. With $\varepsilon = \sqrt{(1/n_+ + 1/n_-)\log(2K/\delta)/2}$, with probability at least $1-\delta$ the picked detector's population evaluated worst-case AUC is at least that of every candidate minus 2ε; in particular, at least that of the graph-free candidate minus 2ε.*

The guarantee is relative to the evaluated attacks, or to the true worst case if certified scores replace evaluated ones, and it degrades with $\log K$: a selection rule that checks few, well-chosen candidates pays less for validation than one that searches many configurations. This is the role of the plug-in in Section 4.

4. Cost-Aware Relation Selection

Theorem 2(b) says what a robust linear detector does with a relation: keep it at its clean weight if its exposure is low, clip it otherwise, and drop it when it is cheap to neutralize relative to its signal; Proposition 7 adds that exposure is a property of each node's degree. CARS turns these results into a training-free defense (Algorithm 1 in Section B).

(1) A moment plug-in of the threat. From training labels, CARS fits a class-balanced relational logistic regression and summarizes every node by its ego projection and, for each relation r, by the mean of its neighbors' projections on the relation's own weight direction. It clusters training fraud and benign nodes by their log-degree vectors and estimates, per group, the class-conditional means and variances of these summaries, the mean projection α_r of accomplices, the projection z_r of the most benign-looking candidates, and the fraud-neighbor rate p_r, which caps deletions. An attack with b injections and k deletions moves a fraud node's mean on r to $(d\mu - k\alpha_r + bz_r)/(d-k+b)$, with edit counts capped by $\lfloor B/c \rfloor$, so a relation on which not one edit is affordable is untouchable. On the multi-relational CSBM these moments reduce to the closed forms of Section 3. Because attacks are per node, the worst-case AUC is the mixture over fraud groups of per-group worst cases, each a small constrained program.

(2) Selection and gates. CARS evaluates every subset of informative relations with its clean-optimal linear detector and keeps the subset of highest plug-in worst-case AUC. CARS-D also chooses a degree threshold τ_r per relation from a grid, by coordinate ascent from the CARS solution; in the plug-in, the attacker may edit a relation while keeping its gate open, close it by deleting accomplices, or open it by injecting candidates. Plug-in models can misjudge real graphs, so CARS-D finally keeps, among the graph-free counterpart itself (LR, the MLP, or HGB-X), the CARS configuration, and the gated configuration, the one with the highest evaluated worst-case AUC on validation nodes under the full attack suite, exhaustive-greedy attack included. The plug-in's role is to propose few candidates: by Proposition 5, the validation guarantee degrades with the number of candidates checked, and each check costs a full attack evaluation; since the counterpart is a candidate, the picked detector is, with high probability, at most 2ε below it.

(3) Robust planning and adversarial training. With uncertain costs, CARS-R plans at costs c/ρ, which by Corollary 1 covers every true cost vector above c/ρ. CARS+AT trains a cost-aware adversarially trained model and keeps it instead of CARS-D only if its validation worst case is higher. CARS and CARS-D apply to any detector, trees included, take seconds to minutes, and generate no attack on training nodes.

5. Experiments

Setup. We use the base multi-relational CSBM of Section 3 (N=6000; an identity relation with $(p, d) = (0.8, 4)$, a behavior relation $(0.55, 10)$, and a dense relation $(0.25, 30)$) and five real graphs: YelpChi and Amazon reviews (Rayana & Akoglu, 2015; McAuley & Leskovec, 2013; Dou et al., 2020); FDCompCN, financial-statement fraud among listed companies linked by investment, supplier, and customer ties (Wu et al., 2023); S-FFSD, simulated card transactions linked by shared card, merchant, location, and type (Xiang et al., 2023); and Tolokers, banned crowd workers linked by common tasks (Platonov et al., 2023). Profile P1 makes the identity-like relation (R-U-R, U-V-U, C-I-C, T-S-T) cheap to manipulate and P2 makes it expensive (Table 3); For each test fraud node, the *evaluated worst case* is the minimum score over a query search, a cost-aware projected-gradient attack (Geisler et al., 2021), transfers, the exact attack for linear models (Proposition 3), and a new *exhaustive-greedy* attack for R-GCN and tree models, which scores every candidate edit exactly at each step, including moves that open or close degree gates (Section B). Edits of all targets are applied jointly and every node is rescored; certified AUCs follow Proposition 4. We compare graph-free baselines (MLP, gradient-boosted trees HGB-X, logistic regression LR), clean-trained graph models (RLR, R-GCN, HGB-Graph (Tang et al., 2023), CARE-GNN (Dou et al., 2020), GNNGuard- and BWGNN-style models (Zhang & Zitnik, 2020; Tang et al., 2022)), models adversarially trained at the evaluated costs (CA-AT, AT-RLR, CA-HGB), and CARS, CARS-D, and CARS+AT, over five seeds with stratified 40/20/40 splits. The confirmatory analysis is fixed in advance: 45 comparisons of CARS-D with its graph-free counterpart (three families, five graphs, three cells), hierarchical bootstrap over seeds and test nodes,

and Holm correction; everything else is exploratory. Costs are scenarios, not estimates (Section C).

The theory predicts where the trap opens. In the CSBM, the numerical trap threshold coincides with the lower envelope of the trap-attack lines to within 0.1% (Figure 1a), and the exposure bound recovers between 0.29 and 0.69 of it. On real graphs, the exact cost maps of linear models show the same weakest-link structure (Figure 1c, Figure 2): the clean-trained RLR falls below LR in 132 of 160 (dataset, cost scale, budget) cells, and the trap closes only when every cheap action becomes expensive. On the CSBM, a factorial study over the identity relation's signal, its isolation rate, its cost, and the architecture explains 0.83 of the variance of the trap depth with cost as the dominant factor; isolating 85% of the nodes on the relation deepens the trap from $-.074$ to $-.253$ on average (Table 38), which is the mechanism behind FDCompCN.

The trust trap is the rule. Clean-trained graph models fall below their graph-free counterparts in 58 of 90 (dataset, model, cell) combinations and lose up to .906 AUC (Table 1). FDCompCN is a negative control: no relation carries clean signal, so no model is trapped, yet clean-trained graph models collapse, because 70–91% of fraudulent firms have no neighbors in a given relation and one injection replaces their empty aggregate (Proposition 2). Adversarial training at the evaluated costs does not remove the trap: at P1, $B{=}2$, CA-AT stays below the MLP on every graph, and CA-HGB, the strongest baseline under the earlier attacks, falls below HGB-X on YelpChi, S-FFSD, and Tolokers once attacks are exhaustive.

How close are the attacks to the worst case? The exhaustive-greedy attack lowers the evaluated worst case of clean-trained nonlinear models by a median of .108 AUC over the earlier suite, and by up to .299 (HGB-Graph on FDCompCN, P2, $B{=}20$; Table 10): it finds a lower score than the suite for 69% of sampled targets (Table 11). Exhaustive search within a pool of candidates still improves on it for 50% of targets on average, and the certified lower bound of Proposition 4 is far below for clean-trained models (Figure 1b): their reported worst cases are optimistic, and the traps understated. We therefore label every nonlinear bracket: 17 of the 20 brackets of clean-trained nonlinear models are loose (width above 0.05), so for them we claim only an upper bound, whereas 15 of the 20 brackets of CARS-D detectors are tight (width at most 0.01) and 4 more informative (Table 10). A single realized joint attack differs from the per-node minimum by $-.001$ AUC on average and by at most .051, so the reported worst cases are achievable up to small cross effects.

Which relation to harden. A platform can raise the cost of manipulating a relation; which one should it harden? Analyses with equal edit costs cannot pose this question. By Corollary 1(c), hardening a relation that some cheapest trap policy avoids changes nothing, and to first order the cheapest policy spends on the action of largest exposure per unit cost (Propositions 1 and 2). Making one relation at a time unmanipulable and computing the exact threshold of the clean-trained RLR on every graph where it has a clean gain and is trapped (16 cases: the CSBM, YelpChi, and CSBMs with 4 to 12 relations), 85% of single-relation hardenings leave the threshold unchanged, and the relation of largest first-order exposure per unit cost, computed from one edit per node without any attack, is the best one to harden in 15 of 16 cases; the most informative relation, the cheapest one, the largest weight, and an exact analysis under equal costs are right in at most 7; for HGB-Graph, the same pick is best in only 2 of 5 graphs (Table 25).

What cost-aware selection adds. CARS-D is at or above its counterpart in 45 of 45 combinations and never more than .002 below; in the confirmatory analysis, no Holm-adjusted comparison is significantly negative, 6 are significantly positive, and non-inferiority at a margin of 0.01 AUC holds in 44 of 45 (Tables 27 and 29; seed standard deviations up to .009). On these graphs, which have one to four relations, its value is abstention: it returns the graph-free detector in 82% of 225 selections, 32 comparisons are exact ties, and validation alone would do as well: validating every relation subset matches it in 50 of 50 cases with about twice the attack evaluations, and the gated candidate adds $+.000$ (Tables 13 and 20). Its gains, where a relation is costly to manipulate, are certified on YelpChi at P2 (Table 10). Planning matters when relations are many (Table 2). On CSBMs with 4 to 12 relations, CARS-D gains $+.112$ over the graph-free detector with at most three validated candidates, whereas the same planner without relation-specific costs gains $+.076$, and $\{$graph-free, all relations$\}$ and random searches with three or eight candidates gain $+.065$, $+.037$, and $+.067$. The advantage grows with the number of relations: with 4 relations all rules are close (CARS-D $+.050$, random search $+.047$), with 12 it is $+.172$ against $+.125$ without costs and $+.087$ for random search. Relation-specific costs are what the plug-in contributes: with identical candidates and validation, cost-blind planning is worse in 20 and better in 7 of 52 cases. Validating every configuration evaluated for any rule (VAL-POOL) does better still, $+.121$, but it validates 13 candidates per cell (up to 24), and exhaustive subset search needs 2^R. On real graphs augmented with six attribute-sharing relations, no rule gains (CARS-D $-.001$), because the added relations carry little signal beyond the features, and the plug-in's gated choice used without validation loses $-.203$ on average: validation is what makes the protocol safe. With 5% of the validation labels, CARS-D

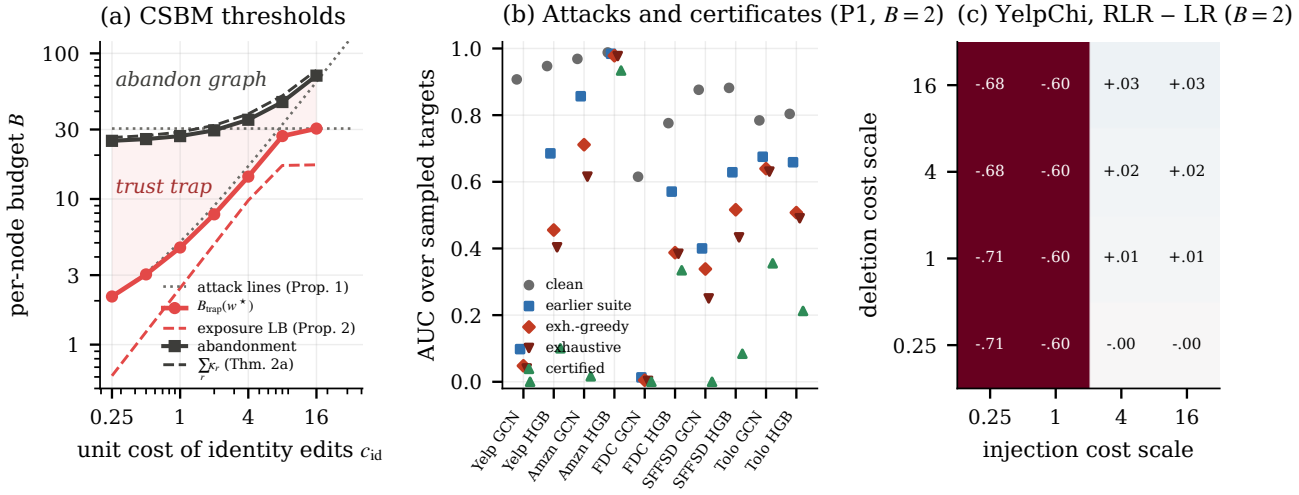

Figure 1. (a) CSBM: the numerical trap threshold of the clean-optimal detector, the lower envelope of the trap-attack lines (Lemma 1), the exposure lower bound (Proposition 1), and the abandonment threshold with its bound $\sum_r \kappa_r$ (Theorem 2). (b) AUC over 150 sampled test fraud nodes per graph (P1, B=2, seed 0) for clean-trained R-GCN and HGB-Graph: clean, under the earlier attack suite, under the exhaustive-greedy attack, under exhaustive search within a pool of candidates, and the certified lower bound; the true worst case lies between the last two. (c) YelpChi: exact worst-case AUC of RLR minus LR at B=2 when injection and deletion costs are scaled relative to P1.

falls below the graph-free detector in 9% of draws, by up to −.099 (Table 22).

Decisions under cost uncertainty. With exact attacks over 16 cost scales, CARS-D planned at the true costs is at or above LR in 152 of 160 cells and above it in 26 (Figure 2). Costs are rarely known: let the defender have an estimate off by a log-normal factor (σ=1) and a box of log half-width w around it (Table 26). Planning at the estimate falls below LR in 11 of 40 draws, by up to −.673, because an error that makes a cheap relation look expensive keeps that relation. The maximin plan, at the box's lower corner (Corollary 1(b)), falls below LR in 1 of 40 draws at w=2 and 0 at w=3, where the box covers the truth in 90% and 100% of draws, but its mean gain is .000, against +.003 with known costs (gains in 2 of 10 cells): safety is bought by giving up the graph, and the gain a defender with exact costs keeps is the value of cost information; minimax regret behaves the same once the box covers the truth. A defender who knows only a lower bound on costs never falls below LR but never gains either (Table 30). In 480 cases with random cost vectors from an independent and a behavior-structured distribution, the trap persists in 95% of cases with a clean gain, and CARS-D falls below LR in 46, by at most −.023, mostly on FDCompCN, whose relations carry no signal (Table 28). Entity-level actions, gangs, and rolling temporal windows on S-FFSD give the same picture (Tables 33, 37 and 50).

6. Related Work

Graph fraud detection. Multi-relational detectors weigh relations through attention or relation-specific weights (Schlichtkrull et al., 2018; Wang et al., 2019), select neighbors against camouflage (Dou et al., 2020; Liu et al., 2021; Zhang et al., 2024), or filter spectra under heterophily (Tang et al., 2022; Gao et al., 2023; Wu et al., 2023); benchmarks show that tree ensembles on neighborhood aggregates are strong baselines (Tang et al., 2023). All of them choose relations for their clean signal; we ask how that choice should change once relations have prices. **Structural attacks and defenses.** Gradient and combinatorial attacks perturb edges under uniform budgets (Zügner et al., 2018; Geisler et al., 2021), low-degree nodes are known to be easier targets (Zügner et al., 2018), and many defenses are broken by adaptive attacks (Mujkanovic et al., 2022; Tramèr et al., 2020); purification (Zhang & Zitnik, 2020; Zhang et al., 2022) reweights or prunes neighbors, and certificates bound the effect of a number of flips (Bojchevski & Günnemann, 2019; Zügner & Günnemann, 2019); our certificates are instead per-node and cost-aware (interval propagation (Gowal et al., 2018), tree bounds (Chen et al., 2019)). Fraud gangs coordinate node injections against several targets (Choi et al., 2025). Cost-aware attacks and defenses price nodes in homogeneous graphs (Liao et al., 2024), and adversaries with action costs go back to adversarial classification (Dalvi et al., 2004; Lowd & Meek, 2005), with cost-sensitive robustness studied for images (Zhang & Evans, 2019) and tabular data (Kireev et al., 2023). We price relations and actions separately. **Strategic classification**

Table 1. Clean and evaluated worst-case test AUC (mean over five seeds; on YelpChi, Amazon, and S-FFSD, CARE-GNN, GNNGuard, BWGNN, CA-AT, AT-RLR, and CARS+AT (R-GCN) use three). Underlined: trust trap ($\Delta_{\text{clean}} > 0.002$ and $\Delta_{\text{rob}} < -0.002$ against the graph-free counterpart: LR for linear models, MLP for neural models, HGB-X for trees). Bold: ours at or above the counterpart (within 0.002). Adversarially trained models and ours are trained or planned at the evaluated cell; their clean column is that of the P1, B=2 plan. Linear models (RLR, AT-RLR, RLR + CARS-D) report exact target-side worst cases; the other graph models report evaluated worst cases, upper bounds on the true worst case, whose certified brackets are labeled in Table 10. All cells are in Table 9.

	YelpChi			Amazon			FDCompCN			S-FFSD			Tolokers		
Model	clean	P1,2	P2,20	clean	P1,2	P2,20	clean	P1,2	P2,20	clean	P1,2	P2,20	clean	P1,2	P2,20
MLP	.836	.836	.836	.975	.975	.975	.715	.715	.715	.788	.788	.788	.740	.740	.740
HGB-X	.935	.935	.935	.977	.977	.977	.789	.789	.789	.845	.845	.845	.734	.734	.734
LR	.771	.771	.771	.969	.969	.969	.680	.680	.680	.644	.644	.644	.722	.722	.722
Clean-trained, graph															
RLR	.798	.181	.546	.971	.865	.674	.656	.001	.001	.660	.355	.061	.768	.529	.529
R-GCN	.915	.043	.639	.976	.739	.594	.631	.003	.002	.878	.325	.143	.811	.649	.648
HGB-Graph	.960	.442	.924	.986	.971	.951	.785	.473	.250	.869	.551	.385	.818	.524	.524
CARE-GNN	.846	.773	.808	.974	.967	.963	.697	.026	.014	.701	.654	.624	.747	.705	.705
GNNGuard-style	.917	.169	.850	.976	.912	.880	.652	.014	.011	.868	.567	.529	.808	.637	.636
BWGNN-style	.902	.593	.868	.978	.967	.946	.688	.093	.082	.867	.715	.654	.831	.729	.726
Adversarially trained															
CA-AT (R-GCN)	.866	.679	.834	.976	.892	.881	.644	.229	.120	.853	.686	.628	.819	.734	.734
AT-RLR (exact)	.773	.366	.597	.967	.935	.901	.619	.014	.012	.645	.562	.529	.737	.698	.697
CA-HGB	.962	.919	.944	.987	.979	.964	.784	.748	.540	.864	.702	.639	.818	.680	.680
Ours															
RLR + CARS	.771	**.771**	**.791**	.972	.951	**.969**	.680	.674	**.680**	.644	.504	**.664**	.722	**.722**	**.722**
R-GCN + CARS	.836	**.836**	**.897**	.976	.943	**.975**	.671	.669	**.715**	.843	.328	**.858**	.740	**.740**	**.740**
HGB + CARS	.935	**.935**	**.952**	.982	**.976**	**.977**	.788	.787	**.789**	.852	.449	**.863**	.734	**.734**	**.734**
RLR + CARS-D	.771	**.771**	**.792**	.969	**.969**	**.969**	.682	**.678**	**.680**	.644	**.644**	**.664**	.722	**.722**	**.722**
R-GCN + CARS-D	.836	**.836**	**.897**	.975	**.975**	**.975**	.715	**.715**	**.715**	.788	**.788**	**.858**	.746	**.741**	**.741**
HGB + CARS-D	.935	**.935**	**.952**	.979	**.975**	**.977**	.790	**.789**	**.789**	.845	**.845**	**.863**	.734	**.734**	**.734**
CARS+AT (R-GCN)	.837	**.837**	**.896**	.973	**.973**	**.973**	.715	**.715**	**.715**	.782	.782	**.856**	.746	**.741**	**.741**
CARS+AT (HGB)	.935	**.935**	**.952**	.979	**.976**	**.977**	.790	**.789**	**.789**	.845	**.845**	**.863**	.734	**.734**	**.734**

Table 2. Selection rules at equal evaluation budget (K validated candidates, median): mean test worst-case AUC difference with the graph-free detector on real graphs, real graphs with six attribute-sharing relations added (Aug.), and CSBMs with 4–12 relations (Rich), the last also with 10% of the validation labels. Counts and regrets in Table 23.

Rule	K	Real	Aug.	Rich	Rich, 10%
CARS-D	2	.000	−.001	+.112	+.110
cost-blind planning	2	.000	−.003	+.076	+.074
no validation (subset)	1	−.005	−.078	+.101	+.101
no validation (gates)	1	−.044	−.203	+.076	+.076
no gated candidate	2	.000	−.001	+.112	+.110
no graph-free candidate	2	−.002	−.078	+.103	+.102
VAL-FA	2	.000	.000	+.065	+.063
VAL-RAND-3	3	−.001	.000	+.037	+.033
VAL-RAND-8	8	−.002	.000	+.067	+.061
VAL-SUB	16	.000	–	+.060	+.054
VAL-POOL	13	−.002	−.004	+.121	+.116

and robust optimization. Strategic classification lets agents alter features at a cost (Hardt et al., 2016; Eilat et al., 2023), and adversarial training solves a min-max problem (Madry et al., 2018). Lemma 1 and Theorem 2 are statements about this min-max problem in a specific structured form: they characterize when an entire input channel, a relation, should be discounted or abandoned, and how the answer depends on the cost vector. **CSBM analyses.** Contextual stochastic block models explain when graph convolution helps (Deshpande et al., 2018; Baranwal et al., 2021; Ma et al., 2022); we extend them to several relations, priced manipulations, and a robust rather than a clean criterion.

7. Conclusion and Limitations

Taking manipulation prices as given, degree gates strictly beat selection and reweighting of a relation on which part of the population is isolated, within a sharp window that real graphs confirm, and exposure per unit cost says which relation to harden; linear models admit exact worst cases, and nonlinear ones certified brackets. Empirically, the trap is the rule, and validated cost-aware selection is never significantly worse than ignoring the graph, mostly abstains with few relations, and earns its planning step with many.

Limitations. Our guarantees are of four kinds (Table 6): exact results are unconditional within their model; certified bounds assume the target-side model and are loose for clean-trained nonlinear detectors, whose reported worst cases are therefore upper bounds; and the validation guarantee is relative to the evaluated attacks and assumes i.i.d. validation nodes. The study is conditional on cost scenarios: no cost was estimated from operational data, errors that make a cheap relation look expensive are harmful unless planning covers them, and without cost information the safe choice is to abstain (Table 26). CARS-D needs labeled validation fraud (Table 22); the evidence that planning helps with many relations comes from synthetic graphs, and augmenting real graphs with attribute relations showed no gain (Table 23); S-

FFSD, simulated, is the only dated graph; the phase diagram and the sensitivity study use the earlier attack suite (Table 8). Eliciting real costs is the next step.

Impact Statement

This paper studies how adversaries can manipulate multi-relational fraud detectors and how defenders can decide which relations to trust or harden. The thresholds, exposure measures, and exact attacks could in principle help an adversary find the cheapest manipulation; we judge the defensive value to dominate, since the attacks are standard white-box evaluations of the kind defenders already run, the exact attack applies only to linear models, and the main practical outputs are decisions to abstain from a relation, to gate it, or to make it harder to manipulate. Our cost profiles are scenarios rather than measurements, and our results should not be read as estimates of the vulnerability of any deployed system. The datasets are public benchmarks; S-FFSD is simulated, and no personal data was collected.

References

Baranwal, A., Fountoulakis, K., and Jagannath, A. Graph convolution for semi-supervised classification: Improved linear separability and out-of-distribution generalization. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, pp. 684–693, 2021.

Bojchevski, A. and Günnemann, S. Certifiable robustness to graph perturbations. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

Chen, H., Zhang, H., Si, S., Li, Y., Boning, D., and Hsieh, C.-J. Robustness verification of tree-based models. In *Advances in Neural Information Processing Systems*, 2019.

Choi, J., Kim, H., and Whang, J. J. Unveiling the threat of fraud gangs to graph neural networks: Multi-target graph injection attacks against GNN-based fraud detectors. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 16028–16036, 2025.

Dalvi, N., Domingos, P., Mausam, Sanghai, S., and Verma, D. Adversarial classification. In *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 99–108, 2004.

Deshpande, Y., Sen, S., Montanari, A., and Mossel, E. Contextual stochastic block models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

Dou, Y., Liu, Z., Sun, L., Deng, Y., Peng, H., and Yu, P. S. Enhancing graph neural network-based fraud detectors against camouflaged fraudsters. In *Proceedings of the 29th ACM International Conference on Information and Knowledge Management (CIKM)*, pp. 315–324, 2020.

Eilat, I., Finkelshtein, B., Baskin, C., and Rosenfeld, N. Strategic classification with graph neural networks. In *International Conference on Learning Representations (ICLR)*, 2023.

Gao, Y., Wang, X., He, X., Liu, Z., Feng, H., and Zhang, Y. Addressing heterophily in graph anomaly detection: A perspective of graph spectrum. In *Proceedings of the ACM Web Conference (WWW)*, pp. 1528–1538, 2023.

Geisler, S., Schmidt, T., Şirin, H., Zügner, D., Bojchevski, A., and Günnemann, S. Robustness of graph neural networks at scale. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.

Gowal, S., Dvijotham, K., Stanforth, R., Bunel, R., Qin, C., Uesato, J., Arandjelovic, R., Mann, T., and Kohli, P. On the effectiveness of interval bound propagation for training verifiably robust models. *arXiv preprint arXiv:1810.12715*, 2018.

Hardt, M., Megiddo, N., Papadimitriou, C., and Wootters, M. Strategic classification. In *Proceedings of the 2016 ACM Conference on Innovations in Theoretical Computer Science (ITCS)*, pp. 111–122, 2016.

Kireev, K., Kulynych, B., and Troncoso, C. Adversarial robustness for tabular data through cost and utility awareness. In *Network and Distributed System Security Symposium (NDSS)*, 2023.

Liao, J., Fu, W., Wang, C., Wei, Z., and Xu, J. Value at adversarial risk: A graph defense strategy against cost-aware attacks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pp. 13763–13771, 2024.

Liu, Y., Ao, X., Qin, Z., Chi, J., Feng, J., Yang, H., and He, Q. Pick and choose: A GNN-based imbalanced learning approach for fraud detection. In *Proceedings of the Web Conference (WWW)*, pp. 3168–3177, 2021.

Lowd, D. and Meek, C. Adversarial learning. In *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 641–647, 2005.

Ma, Y., Liu, X., Shah, N., and Tang, J. Is homophily a necessity for graph neural networks? In *International Conference on Learning Representations (ICLR)*, 2022.

Madry, A., Makelov, A., Schmidt, L., Tsipras, D., and Vladu, A. Towards deep learning models resistant to adversarial attacks. In *International Conference on Learning Representations (ICLR)*, 2018.

McAuley, J. J. and Leskovec, J. From amateurs to connoisseurs: Modeling the evolution of user expertise through

online reviews. In *Proceedings of the 22nd International Conference on World Wide Web (WWW)*, pp. 897–908, 2013.

Mujkanovic, F., Geisler, S., Günnemann, S., and Bojchevski, A. Are defenses for graph neural networks robust? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Platonov, O., Kuznedelev, D., Diskin, M., Babenko, A., and Prokhorenkova, L. A critical look at the evaluation of GNNs under heterophily: Are we really making progress? In *International Conference on Learning Representations (ICLR)*, 2023.

Rayana, S. and Akoglu, L. Collective opinion spam detection: Bridging review networks and metadata. In *Proceedings of the 21st ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 985–994, 2015.

Schlichtkrull, M., Kipf, T. N., Bloem, P., van den Berg, R., Titov, I., and Welling, M. Modeling relational data with graph convolutional networks. In *The Semantic Web: 15th International Conference (ESWC)*, pp. 593–607, 2018.

Tang, J., Li, J., Gao, Z., and Li, J. Rethinking graph neural networks for anomaly detection. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*, pp. 21076–21089, 2022.

Tang, J., Hua, F., Gao, Z., Zhao, P., and Li, J. GADBench: Revisiting and benchmarking supervised graph anomaly detection. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.

Tramèr, F., Carlini, N., Brendel, W., and Madry, A. On adaptive attacks to adversarial example defenses. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Vincess, D. and Melvin, F. Not all edges cost the same: Relation-specific manipulation costs and the trust trap in fraud graphs. In *GraphML Workshop at the Asian Conference on Machine Learning (ACML)*, 2026.

Wang, X., Ji, H., Shi, C., Wang, B., Ye, Y., Cui, P., and Yu, P. S. Heterogeneous graph attention network. In *Proceedings of the World Wide Web Conference (WWW)*, pp. 2022–2032, 2019.

Wu, B., Yao, X., Zhang, B., Chao, K.-M., and Li, Y. SplitGNN: Spectral graph neural network for fraud detection against heterophily. In *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management (CIKM)*, 2023.

Xiang, S., Zhu, M., Cheng, D., Li, E., Zhao, R., Ouyang, Y., Chen, L., and Zheng, Y. Semi-supervised credit card fraud detection via attribute-driven graph representation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pp. 14557–14565, 2023.

Zhang, J., Xu, Z., Lv, D., Shi, Z., Shen, D., Jin, J., and Dong, F. DiG-In-GNN: Discriminative feature guided GNN-based fraud detector against inconsistencies in multi-relation fraud graph. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pp. 9323–9331, 2024.

Zhang, M., Wang, X., Zhu, M., Shi, C., Zhang, Z., and Zhou, J. Robust heterogeneous graph neural networks against adversarial attacks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pp. 4363–4370, 2022.

Zhang, X. and Evans, D. Cost-sensitive robustness against adversarial examples. In *International Conference on Learning Representations*, 2019.

Zhang, X. and Zitnik, M. GNNGuard: Defending graph neural networks against adversarial attacks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Zügner, D. and Günnemann, S. Certifiable robustness and robust training for graph convolutional networks. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, 2019.

Zügner, D., Akbarnejad, A., and Günnemann, S. Adversarial attacks on neural networks for graph data. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pp. 2847–2856, 2018.

A. Proofs

CSBM moments. Let $n_r^f = p_r d_r - k_r$, $n_r^b = (1-p_r)d_r + b_r$, and $n_r = n_r^f + n_r^b$. For a fraudulent node under attack (b, k),

$$\mu_{1,r} = \frac{m\, n_r^f - m(1-p_r)d_r - (m+\delta)\, b_r}{n_r}, \qquad \Sigma_{1,rr} = \frac{1}{n_r} + \frac{4m^2 d_r p_r (1-p_r)}{n_r^2},$$

where the second term accounts for the binomial composition of the neighborhood. For a benign node, $\mu_{0,r} = (2q_r - 1)m$ and $\Sigma_{0,rr} = (1 + 4m^2 q_r(1-q_r))/d_r^0$. For the own feature, $\mu_{y,0} = \pm m$ and $\Sigma_{y,00} = 1$. Relations are independent, so the covariances are diagonal, z_v is Gaussian given the class, and $\mathrm{AUC} = \Pr(z_{\text{fraud}} > z_{\text{benign}}) = \Phi(\zeta)$ with ζ as in (2). We write $\sigma_r^2 = \Sigma_{1,rr}(0) + \Sigma_{0,rr}$.

Proof of Lemma 1. (a) Under per-node budgets, f is trapped at (B, c) if and only if some affordable policy is a trap policy, that is, if and only if $B \geq \max_v c^\top \theta_v(\pi)$ for some $\pi \in \mathcal{T}(f)$. The set of budgets at which f is trapped is therefore $\{B \geq \inf_{\pi \in \mathcal{T}(f)} \max_v c^\top \theta_v(\pi)\}$. The infimum is attained: edit counts are nonnegative integers and $c > 0$, so only finitely many edit vectors cost at most any given level, and the set of values $\max_v c^\top \theta_v(\pi)$ below any level is finite. The gang case is identical with the mean in place of the maximum, and $\frac{1}{|T|}\sum_v \leq \max_v$ gives $\bar{B}_{\text{trap}} \leq B_{\text{trap}}$. (b) For fixed π, $c \mapsto \max_v c^\top \theta_v(\pi)$ and $c \mapsto \frac{1}{|T|}\sum_v c^\top \theta_v(\pi)$ are nondecreasing and positively homogeneous because $\theta_v \geq 0$; both properties survive a pointwise minimum. (c) For fixed π, $c \mapsto \frac{1}{|T|}\sum_v c^\top \theta_v(\pi) = c^\top \bar{\theta}(\pi)$ is linear, and a pointwise minimum of linear functions is concave; it is piecewise linear because, on any bounded set of costs, only finitely many mean edit vectors can attain the minimum. If performance depends on π only through a common edit vector θ ($\theta_v = \theta$ for all v), the per-node cost is $c^\top \theta$, linear, and the same argument applies. Without that property, $\max_v$ is convex in c; the two-node example in the text, where node 1 can be trapped only by one unit of action 1 and node 2 only by one unit of action 2, gives $B_{\text{trap}} = \max(c_1, c_2)$, which is convex and not concave. No property of f, of the data, or of the performance measure was used. The vertex statement of Corollary 1 follows because a concave function attains its minimum over a polytope at a vertex. □

Proof of Proposition 4. Under the target-side model, the attacked score of each fraudulent test node is at least ℓ_v and benign scores do not move, so every fraud–benign pair ordered correctly by (ℓ_v, s_w) is ordered correctly under every feasible attack; hence $A_{\text{true}} \geq A_{\text{cert}}$. Any evaluated attack is feasible, so $A_{\text{true}} \leq A_{\text{eval}}$. The margin statement takes $\ell_v = s_v - \Delta_v$. For the interval, write $\bar{x}'_{v,r} = (S' - \sum_{u \in D} x_u + \sum_{u \in I} x_u)/(d - k + b)$ for a deletion set D of k accomplices and an injection set I of b candidates. Coordinate j of $W_r \bar{x}'_{v,r}$ is $(S_j - \sum_{u \in D}(W_r x_u)_j + \sum_{u \in I}(W_r x_u)_j)/(d - k + b)$, which is smallest when D holds the k accomplices with the largest $(W_r x_u)_j$ and I the b candidates with the smallest, and largest in the opposite case; the bounds are attained coordinatewise, though not jointly. If $d - k + b = 0$, the relation is empty and the aggregate is exactly 0, which the certificate includes as a separate case. Interval arithmetic through ReLU (monotone) and linear maps gives valid bounds for deeper layers; for neighbors and candidates whose own aggregation changes because the target joins or leaves it, we bound their states over every subset of relations in which this happens. For a tree ensemble, every leaf reached by some input in the box of reachable features is a candidate, and the sum over trees of the smallest reachable leaf value lower-bounds the score. □

Proof of Proposition 5. Fix a candidate. Its evaluated worst-case validation AUC $\hat{A}$ is a two-sample U-statistic of i.i.d. attacked scores, so $\mathbb{E}\hat{A} = A$, the population value. Replacing one fraudulent node changes $\hat{A}$ by at most $1/n_+$ and one benign node by at most $1/n_-$; by McDiarmid's inequality, $\Pr(|\hat{A} - A| \geq \varepsilon) \leq 2\exp(-2\varepsilon^2/(1/n_+ + 1/n_-))$. A union bound over the K candidates gives $|\hat{A}_k - A_k| < \varepsilon$ for all k with probability at least $1 - \delta$ at the stated ε. On that event, the picked candidate $\hat{k}$ satisfies $A_{\hat{k}} > \hat{A}_{\hat{k}} - \varepsilon \geq \hat{A}_k - \varepsilon > A_k - 2\varepsilon$ for every k. □

Gangs with cross effects. Let A_{sep} be the AUC computed from separable attacked fraud scores and clean benign scores, and A_{real} the AUC after all edits. The two score vectors differ only on $E_+ \cup E_-$. A fraudulent node in E_+ enters n_- pair indicators and a benign node in E_- enters n_+, so $|A_{\text{real}} - A_{\text{sep}}| \leq (|E_+|n_- + |E_-|n_+)/(n_+ n_-) = |E_+|/n_+ + |E_-|/n_-$. □

Proof of Proposition 1. The set $\Theta_B(c)$ is convex and contains 0, so for $\theta \in \Theta_B(c)$ the segment $t\theta$, $t \in [0, 1]$, stays in it. Since ζ is continuously differentiable on the feasible set (all neighborhood sizes stay positive or the relation empties, where

the moments extend continuously),

$$\zeta(\theta) - \zeta(0) = \int_0^1 \nabla\zeta(t\theta)^\top \theta \, dt \;\geq\; -\sum_j L_j(B)\theta_j \;\geq\; -\max_j \frac{L_j(B)}{c_j}\, c^\top\theta \;\geq\; -B \max_j \frac{L_j(B)}{c_j}.$$

If the right-hand side exceeds $\zeta_0 - \zeta(0)$, every feasible attack leaves $\zeta > \zeta_0$, so $B_{\mathrm{trap}} \geq B$. For the first-order exposure, $\partial\mu_{1,r}/\partial b_r|_0 = -(2p_r m + \delta)/d_r$ and $\partial\Sigma_{1,rr}/\partial b_r|_0 = -(1 + 8m^2 p_r(1-p_r))/d_r^2$; the variance term enters the derivative of ζ with the factor $\zeta w_r^2/(2s^2)$, hence the remainder $O(w_r^2 d_r^{-2})$. Over small budgets the caps are inactive and the feasible set is the simplex with vertices 0 and $(B/c_j)e_j$, so the attacker's first-order optimum is a vertex: the whole budget goes to the action of largest exposure. □

Proof of Proposition 2. Let $d = d_r(v)$, $\bar{x} = \bar{x}_{v,r}$ (with $\bar{x} = 0$ if $d = 0$), and $S = \sum_{u\in\mathcal{N}_r(v)} x_u = d\bar{x}$. After injecting the endpoints I ($|I| = b$) and deleting the accomplices D ($|D| = k$), the relation has $n = d - k + b$ neighbors. *Case* $n \geq 1$. The new mean is $\bar{x}' = (S + \sum_{u\in I} x_u - \sum_{u\in D} x_u)/n$, and since $S = d\bar{x}$ and $n = d - k + b$,

$$\bar{x}' - \bar{x} = \frac{\sum_{u\in I}(x_u - \bar{x}) - \sum_{u\in D}(x_u - \bar{x})}{n}, \qquad \|\bar{x}' - \bar{x}\| \leq \frac{(b+k)D_{v,r}}{n}.$$

This identity also holds for $d = 0$, where $\bar{x} = 0$ and $D_{v,r}$ bounds the norms of the injected endpoints. *Case* $n = 0$. The attack deletes every neighbor ($k = d$, $b = 0$), the aggregate becomes 0 by convention, and $\|\bar{x}' - \bar{x}\| = \|\bar{x}\|$; this change is not controlled by deviations from the mean (two neighbors equal to $\bar{x}$ give $D_{v,r} = 0$), which is why the statement bounds it separately. In both cases the pre-activation changes by $\sum_r W_r(\bar{x}'_{v,r} - \bar{x}_{v,r})$, whose norm is at most $\sum_r \|W_r\|_2 \Delta_{v,r}$, and s_v is Λ-Lipschitz in it. For one unit of budget: one injection gives $n = d + 1 \geq 1$ and a change of at most $D_{v,r}/(d+1)$ at cost c_r^+; one deletion gives $n = d - 1$, a change of at most $D_{v,r}/(d-1)$ if $d \geq 2$ and exactly $\|\bar{x}\|$ if $d = 1$, at cost c_r^-. The node-level bounds of Table 43 use these two cases. □

Proof of Theorem 2. (a) For relation r, $\Delta_r(b_r, k_r) \leq 0$ is equivalent to $m(p_r d_r - k_r) - m(1-p_r)d_r - (m+\delta)b_r \leq \mu_{0,r}(d_r - k_r + b_r)$, because the denominator is positive; this is a half-space in (b_r, k_r), so its cheapest point under $c_r^+ b_r + c_r^- k_r$ is a vertex, $(\bar{b}_r, 0)$ or $(0, \bar{k}_r)$, and $\bar{k}_r \leq p_r d_r$ respects the cap. Hence κ_r is the exact cost of neutralizing r. Let w be any detector and suppose $B \geq \sum_r \kappa_r$. The attacker neutralizes every relation with $w_r > 0$ and leaves the others untouched. Then $w_r\Delta_r(\theta) \leq 0$ for $w_r > 0$, and $w_r\Delta_r(0) \leq 0$ for $w_r \leq 0$ because $\Delta_r(0) \geq 0$ when $p_r \geq q_r$. The numerator of ζ is at most $w_0\Delta_0 = 2mw_0$. If $w_0 \leq 0$, $\zeta \leq 0 < \zeta_0$. If $w_0 > 0$, the denominator is at least $w_0\sqrt{\Sigma_{1,00} + \Sigma_{0,00}} = \sqrt{2}w_0$, so $\zeta \leq \sqrt{2}m = \zeta_0$. Thus no detector is above the graph-free one, and $B_{\mathrm{ab}} \leq \sum_r \kappa_r$.

(b) Under linear response, action j on relation $r(j)$ lowers $\Delta_{r(j)}$ by $\Delta_{r(j)}/n_j$ per unit, $n_j \in \{\bar{b}_{r(j)}, \bar{k}_{r(j)}\}$, and variances are fixed, so $\zeta(w;\theta) = (w^\top\Delta - \sum_j \theta_j w_{r(j)}\Delta_{r(j)}/n_j)/\|w\|$ with $\|w\|^2 = \sum_{r\geq 0} w_r^2\sigma_r^2$. The attacker minimizes a linear function over a simplex, so it spends its budget on one action, and the damage is $B\max_r(w_r)_+\Delta_r/\kappa_r$. For $w_r^\star = \Delta_r/\sigma_r^2$, $w^{\star\top}\Delta = \|w^\star\|^2 = \zeta_\star^2$ and the damage is $B\max_r \zeta_r^2/\kappa_r$, which gives $B_{\mathrm{trap}}(w^\star)$. If $B \geq \sum_r \kappa_r$, part (a) applies. If $B < \sum_r \kappa_r$, the direction $v_r = \kappa_r/\Delta_r$ gives $\zeta(e_0 + \varepsilon v) = \zeta_0 + \varepsilon(\sum_r \kappa_r - B)/\sigma_0 + O(\varepsilon^2) > \zeta_0$ for small $\varepsilon > 0$, so $B_{\mathrm{ab}} = \sum_r \kappa_r$. Finally, the robust objective has a concave, positively homogeneous numerator, so maximizing it over $\{\|w\| \leq 1\}$ is a convex program whose KKT conditions read $\Delta - B\sum_r \mu_r(\Delta_r/\kappa_r)e_r = 2\nu\,\mathrm{diag}(\sigma^2)w$, with μ a probability vector on the relations that attain the maximum t. Off its support, $w_r = \lambda\Delta_r/\sigma_r^2$ with $\lambda = 1/(2\nu)$, as for w_0; on its support, $w_r\Delta_r/\kappa_r = t$ and $w_r \leq \lambda\Delta_r/\sigma_r^2$. □

Coordinated adversaries.

Proposition 6 (Gangs convexify the attack curve)**.** *Let $A(B)$ be the AUC under the per-node worst case at budget B, with benign scores fixed and edits acting on their own targets. A gang with average budget $\bar{B}$ can drive the AUC down to the lower convex envelope $\breve{A}(\bar{B})$ of A (up to rounding of $\phi|T|$). In particular, if $A(0) > A_{\mathrm{free}} > A(B_1)$, the gang traps the detector with average budget $\bar{B} \leq f^\star B_1$, where $f^\star = (A(0) - A_{\mathrm{free}})/(A(0) - A(B_1))$.*

A detector whose clean advantage is small relative to its per-node damage is trapped by a gang that spends a few units on a few members. When edits also move other nodes' scores (shared endpoints, two-hop aggregation), the realized AUC differs from the envelope by at most $|E_+|/n_+ + |E_-|/n_-$, where $E_\pm$ are the fraudulent and benign test nodes whose scores are affected by edits other than their own; we measure both.

Proof of Proposition 6. Let $F_v(b)$ be the fraction of benign nodes scored above target v when v receives budget b and plays its per-node worst case; with benign scores fixed and edits acting on their own targets, the AUC of an allocation (b_v) is $\frac{1}{|T|}\sum_v F_v(b_v)$, and $A(b) = \frac{1}{|T|}\sum_v F_v(b)$. For budgets b_1, b_2 and a fraction $\phi \in [0,1]$ with $\phi b_1 + (1-\phi)b_2 = \bar{B}$, assign b_1 to a uniformly random subset of $\phi|T|$ targets and b_2 to the others. The expected AUC is $\phi A(b_1) + (1-\phi)A(b_2)$, and some deterministic subset does at least as well. Minimizing over (ϕ, b_1, b_2) gives the lower convex envelope of A on the real line, since two points suffice in one dimension (Carathéodory). The second claim takes $b_2 = 0$, $b_1 = B_1$, and $\phi = f^\star$. A gang that targets the most vulnerable members does strictly better whenever the F_v differ. □

Proof of Proposition 3. The score of target v depends on its edits only through the means $\bar{x}_{v,r}$, and $w_r^\top \bar{x}'_{v,r} = (\sum_{u\in\mathcal{N}_r(v)} \pi_r(u) + \sum_{u\in I} \pi_r(u) - \sum_{u\in D} \pi_r(u))/(d-k+b)$ with $\pi_r(u) = w_r^\top x_u$. For fixed counts (b,k) the denominator is fixed, so the minimum takes the b candidates with the smallest π_r and the k accomplices with the largest. Prefix sums give the minimum for every (b,k) in $O(b_r^{\max} k_r^{\max})$; a running minimum over cost units gives the best contribution of relation r for every budget; and the relations are combined by a dynamic program over cost units in $O(R(B/\epsilon)^2)$. The edits of other targets do not enter v's score, which is why the result is exact on the target side; the AUC is then evaluated with all edits applied jointly. □

Degree gating bounds.

Proposition 7 (Degree gating)**.** *In the setting of Proposition 2, let relation r enter h_v only if the current degree satisfies $d_r(v) \geq \tau_r \geq 1$. (a) If the gate is open before and after the attack, the attack changes s_v through r by at most $\Lambda\|W_r\|_2(b_r+k_r)D_{v,r}/\tau_r$. (b) Closing an open gate requires $d_r(v) - \tau_r + 1$ deletions, hence as many accomplices and a cost of $c_r^-(d_r(v)-\tau_r+1)$, and changes s_v by at most $\Lambda\|W_r\|_2\|\bar{x}_{v,r}\|$. (c) Opening a closed gate requires $\tau_r - d_r(v)$ injections, at cost $c_r^+(\tau_r - d_r(v))$. Hence, under a per-node budget B, nodes with $d_r(v) > \tau_r - 1 + B/c_r^-$ cannot be switched off, nodes with $d_r(v) < \tau_r - B/c_r^+$ cannot be switched on, and the per-unit exposure of every open node is at most $\Lambda\|W_r\|_2 D_{v,r}/(\tau_r \min(c_r^+, c_r^-))$.*

Proof of Proposition 7. (a) Let $n = d_r(v) - k_r + b_r$ be the degree after the attack, with $n \geq \tau_r$ because the gate stays open. The new mean is $\bar{x}' = (d_r(v)\bar{x}_{v,r} + \sum_{u\in I} x_u - \sum_{a\in K} x_a)/n$, so $\bar{x}' - \bar{x}_{v,r} = (\sum_{u\in I}(x_u - \bar{x}_{v,r}) - \sum_{a\in K}(x_a - \bar{x}_{v,r}))/n$ and $\|\bar{x}' - \bar{x}_{v,r}\| \leq (b_r+k_r)D_{v,r}/\tau_r$. The pre-activation moves by at most $\|W_r\|_2$ times this, and s_v by at most Λ times that. (b) After the attack the gate is closed, so $d_r(v) - k_r + b_r \leq \tau_r - 1$ and, since $b_r \geq 0$, $k_r \geq d_r(v) - \tau_r + 1$; only accomplices can be deleted, at unit cost c_r^-. A closed gate replaces $\bar{x}_{v,r}$ by 0, which moves the pre-activation by at most $\|W_r\|_2\|\bar{x}_{v,r}\|$. (c) Symmetrically, $d_r(v) + b_r - k_r \geq \tau_r$ forces $b_r \geq \tau_r - d_r(v)$. The last statement combines (a)–(c): with budget B, at most B/c_r^- deletions and B/c_r^+ injections are affordable, and $b_r + k_r \leq B/\min(c_r^+, c_r^-)$. □

Proof of Corollary 1. If $c' \geq c$, every $\theta \geq 0$ with $c'^\top\theta \leq B$ satisfies $c^\top\theta \leq c'^\top\theta \leq B$, and the deletion caps do not depend on costs, so $\Theta_B(c') \subseteq \Theta_B(c)$; a minimum over a smaller set is larger. Finally $c^\top\theta \leq \rho B$ if and only if $(c/\rho)^\top\theta \leq B$. (b) Every $c \in U$ satisfies $c \geq \underline{c}$, so $A_B(f;c) \geq A_B(f;\underline{c})$ by (a), with equality at $c = \underline{c} \in U$; the maximin value of each detector is therefore its worst case at $\underline{c}$. (c) Let π be a cheapest trap policy at c with no action on r, and let c' differ from c only on the coordinates of r, with $c' \geq c$. The cost of π is the same under c', so $B_{\mathrm{trap}}(f;c') \leq B_{\mathrm{trap}}(f;c)$ by Lemma 1(a), and $B_{\mathrm{trap}}(f;c') \geq B_{\mathrm{trap}}(f;c)$ by Lemma 1(b). □

Proof of Theorem 1. Under the Gaussian moments of the CSBM, each (class, degree) component has a Gaussian score, and the AUC is $\sum_{g,h} \pi_g^f \pi_h^b \Phi(\eta_{gh})$ over fraud components g and benign components h, with η_{gh} the standardized mean difference of the pair. Isolated nodes have an empty aggregate, set to 0; injected candidates have projection $-(m+\delta)$. (a) Fix $w > 0$ and $b < \tau \leq d$. Under the gated detector, an isolated fraud node cannot open the gate ($0 + b < \tau$), and closing the gate of a connected fraud node is unaffordable, so the attacker's best response acts on connected fraud nodes only; call it $\pi^\star$. Apply $\pi^\star$ to the connected fraud nodes of the ungated detector, and inject b candidates into every isolated fraud node. Connected nodes then have the same scores under both detectors (their gate is open before and after the attack, since $d \geq \tau$ and injections only raise degrees), benign isolated nodes have the same score $w_0 x_v$, and isolated fraud nodes have score $w_0 x_v - w(m+\delta)$ instead of $w_0 x_v$. Every pair involving an isolated fraud node is therefore ordered correctly with strictly smaller probability and every other pair with the same probability, so the ungated AUC under this attack is strictly below the gated worst case, and the ungated worst case is no larger than this AUC. (b) At $w = 0$, every pair has $\eta_{gh} = 2mw_0/(\sqrt{2}w_0) = \sqrt{2}m$, so $\partial_w \Phi(\eta_{gh})|_{w=0} = \phi(\sqrt{2}m)\,\mathbb{E}[\Delta a_{gh}]/(\sqrt{2}w_0)$ with Δa_{gh} the difference of the relation terms of the pair. Under the attack, the relation term of a connected fraud node has mean μ'_d, that of a connected benign

node μ_b, and isolated nodes contribute 0 (gate closed). Summing over pairs with weights $(1-\pi_0)^2$, $(1-\pi_0)\pi_0$, $\pi_0(1-\pi_0)$, π_0^2 gives a derivative proportional to $(1-\pi_0)^2(\mu'_d-\mu_b)+(1-\pi_0)\pi_0\mu'_d-\pi_0(1-\pi_0)\mu_b=(1-\pi_0)(\mu'_d-\mu_b)$, which is positive iff $b<\bar{b}_d$. The attacker's best response may depend on w, but at $w=0$ every response gives the same AUC, and the worst case is a minimum of differentiable functions over a finite set of injection counts, so its right derivative at 0 is the smallest of their derivatives, each equal to the expression above with the corresponding count $b'\le b$, hence positive. Therefore $A_{\text{gated}}(w)>A_{\text{gated}}(0)=A_{\text{drop}}$ for small $w>0$. Combining (a) and (b), $\sup_w A_{\text{gated}}\ge A_{\text{gated}}(w_1)>A_{\text{drop}}$ for some $w_1>0$, and $\sup_w A_{\text{gated}}\ge A_{\text{gated}}(w)>A_{\text{ungated}}(w)$ for every $w>0$; the supremum over ungated $w\ge 0$ is attained on a compact range of normalized weights, which gives the strict inequality. (c) Let $b\ge\tau$. An isolated fraud node can open its gate with $\tau\le b$ injections of the most benign-looking candidates, and its aggregate is then $-(m+\delta)$, exactly the value it reaches in the ungated detector with any number $b'\ge 1$ of injections; it can also leave its aggregate at 0 in both detectors. Its set of reachable relation terms, $\{0,-(m+\delta)\}$, is therefore the same with and without the gate. A connected fraud node has degree $d\ge\tau$, injections only raise its degree, and closing is unaffordable, so its gate stays open and its reachable terms are the same as without the gate. Benign nodes are not attacked and have the same scores in both detectors (isolated ones contribute 0 either way). Every node thus has the same set of attacked scores, the best responses coincide, and so do the worst-case AUCs. If $\tau>d$, no clean node passes the gate, so the clean scores equal those of the detector without the relation; an attack can only open gates, and with $w\ge 0$ it does so only when this lowers a fraud score, so the gated worst case is at most A_{drop}. □

B. Experimental Details

Datasets. **YelpChi** (Rayana & Akoglu, 2015; Dou et al., 2020): 45,954 reviews (14.5% fraud), 32 features; R-U-R (same user), R-T-R (same product, same month), R-S-R (same product, same rating). Fraud-node mean degrees are 0.29, 25.1, and 142.6, and 82% of fraudulent reviews have no R-U-R neighbor. **Amazon** (McAuley & Leskovec, 2013; Dou et al., 2020): 11,944 users (6.9% fraud), 25 features; U-P-U, U-S-U, U-V-U; following the CARE-GNN release, the first 3,305 users stay in the graph but are not used for training or evaluation. Degrees are heavy-tailed: the median U-P-U degree of fraud nodes is 6 and the mean 49. **FDCompCN** (Wu et al., 2023): 5,317 Chinese listed companies (10.5% financial-statement fraud), 57 features; C-I-C (investment), C-P-C (supplier), C-S-C (customer). Its relations are sparse (fraud-node degrees 0.79, 0.12, 0.17; 70–91% of fraud nodes are isolated) and barely homophilous. **S-FFSD** (Xiang et al., 2023): 77,881 card transactions, of which 29,643 are labeled (17.7% fraud); features are the log amount and the time stamp. We link each transaction to the three previous transactions with the same source (T-S-T), target merchant (T-M-T), location (T-L-T), and type (T-Y-T); unlabeled transactions stay in the graph but are neither trained on, evaluated, nor used as attack endpoints. The relations are strongly homophilous (mean fraud share of neighbors 0.25–0.48 for fraud against 0.06–0.12 for benign transactions), and S-FFSD is simulated, which we keep in mind when interpreting it. **Tolokers** (Platonov et al., 2023): 11,758 crowd workers (21.8% banned), 10 features, one relation (W-T-W: worked on a common task); fraud-node degrees are high (median 38, mean 136). Features are standardized, and we use stratified 40/20/40 splits of the evaluated nodes (five seeds for every real graph, seeds 3 and 4 of YelpChi, Amazon, and S-FFSD without CARE-GNN, GNNGuard, BWGNN, and adversarial training; three for the synthetic graph).

Relation-rich and augmented graphs. The relation-rich CSBMs have N=6000 nodes (10% fraud), m=0.35, and $R\in\{4,8,12\}$ relations; each relation is informative with probability 0.6 ($p\sim U[0.5,0.85]$, $q\sim U[0.02,0.1]$) and uninformative otherwise (p=q=0.1), has degree $d\in\{2,4,10,30\}$, and with probability 0.3 a fraction 0.5 or 0.8 of the nodes has no edge on it. Injection and deletion costs are drawn independently and log-uniformly in $[0.25,16]$ and rounded to multiples of 0.25, so that a relation's cost carries no information about its signal. Augmented real graphs add to FDCompCN, Tolokers, and Amazon six attribute-sharing relations: for a random feature, nodes are sorted by it and cut into bins of about 50, and each node is linked to three random nodes of its bin (symmetrized); their costs are drawn in the same way, and the original relations keep their profile costs.

Models. MLP and HGB-X (histogram gradient boosting, 300 iterations, learning rate 0.05, balanced class weights) use features only; LR is logistic regression on features (RLR without relations). RLR is a logistic regression on $[x_v,\bar{x}_{v,1},\dots,\bar{x}_{v,R}]$. R-GCN has two layers with 32 units, mean aggregation per relation, and a separate ego weight. HGB-Graph is gradient boosting on features and relation means (Tang et al., 2023); a gated relation is a missing value where the gate is closed, which trees handle natively. CARE-GNN is our vectorized re-implementation of the official code (label-aware similarity, top-p neighbor selection with the reinforcement rule, official hyperparameters). GNNGuard-style: per relation and layer, edges are weighted by the cosine similarity of the endpoints' representations, edges below 0.1 are pruned, and weights are

Table 3. Cost profiles: unit costs (c^+, c^-) per relation.

Dataset	Profile	Costs
Synthetic	c_{id}	identity $(c_{\mathrm{id}}, c_{\mathrm{id}})$, behavior $(2, 4)$, dense $(1, 4)$; P1: c_{id}=0.5, P2: c_{id}=16
YelpChi	P1 (cheap identities)	R-U-R $(2, 0.5)$, R-T-R $(1, 4)$, R-S-R $(1, 4)$
	P2 (costly identities)	R-U-R $(25, 25)$, R-T-R $(1, 4)$, R-S-R $(1, 4)$
	entity Q1 / Q2	aged account 2 / 25, fresh account 0.5 / 25, rating 1, timing 1
Amazon	P1	U-P-U $(1, 4)$, U-S-U $(1, 2)$, U-V-U $(1, 0.5)$
	P2	U-P-U $(1, 4)$, U-S-U $(1, 2)$, U-V-U $(8, 8)$
FDCompCN	P1 (concealable ownership)	C-I-C $(4, 0.5)$, C-P-C $(1, 2)$, C-S-C $(1, 2)$
	P2 (transparent registries)	C-I-C $(8, 8)$, C-P-C $(1, 2)$, C-S-C $(1, 2)$
S-FFSD	P1 (cheap cards)	T-S-T $(2, 0.5)$, T-M-T, T-L-T, T-Y-T $(1, 4)$
	P2 (verified cards)	T-S-T $(25, 25)$, T-M-T, T-L-T, T-Y-T $(1, 4)$
Tolokers	P1 / P2	W-T-W $(1, 1)$ / $(8, 8)$

renormalized (Zhang & Zitnik, 2020). BWGNN-style: an MLP encoder, the three Beta-wavelet filters of order 2 per relation, a maximum over relations, and an MLP head (Tang et al., 2022). CA-AT is an R-GCN trained against the stronger of the gradient and query attacks at the evaluated profile and budget (Madry et al., 2018), from epoch 40 (one third of training) every 12 epochs, with checkpoints selected on the mean of clean and attacked validation AUC; AT-RLR is RLR trained the same way with the exact attack of Proposition 3 as inner attack, so its training and its evaluation are exact; CA-HGB refits trees on the aggregates of attacked training fraud nodes. Neural models are trained full-batch with Adam (learning rate 0.01, weight decay 5×10^{-4}) and class-weighted cross-entropy for 120 epochs on YelpChi and Amazon and 200 elsewhere (300 for RLR), selecting checkpoints on validation AUC.

Algorithm 1 Cost-aware relation selection (CARS, CARS-D)

Require: graph, training and validation labels, planning costs (c^+, c^-), budget B, robustness factor ρ, detector family $\mathcal{F}$

1: fit a relational logistic regression; project ego features and each relation on its direction
2: cluster training nodes by log-degree vectors; estimate group moments, α_r, z_r, p_r
3: $V(\tau) \leftarrow$ plug-in worst-case AUC at costs $(c^+, c^-)/\rho$ of the clean-optimal linear detector using relation r on nodes of degree $\geq \tau_r$ (mixture of per-group SLSQP best responses)
4: $\tau^{\mathrm{CARS}} \leftarrow \arg\max V$ over $\tau \in \{0, \infty\}^R$ (subset selection)
5: $\tau^{\mathrm{D}} \leftarrow$ coordinate ascent of V over a grid of thresholds from τ^{CARS}
6: train detectors of $\mathcal{F}$ for τ^{CARS} and τ^{D}
7: **return** CARS: the detector for τ^{CARS}; CARS-D: among the graph-free counterpart and the two trained detectors, the one with the best evaluated worst-case validation AUC

CARS details. The plug-in uses $K = 8$ fraud and 6 benign degree groups (k-means on $\log(1 + d)$ vectors, group degree = median), a candidate pool of the 300 most benign-looking labeled benign nodes, and the threshold grid $\{0, 1, 2, 3, 5, 8, 13, 21, 34, 55, \infty\}$. Per-unit costs of an action that neutralizes a relation with fewer than one edit are inflated accordingly. A relation is informative if its fraud-group mean exceeds its benign-group mean. The per-group best response is a multi-start SLSQP over the edit counts of the relations whose gate is open, enumerating the gate actions (keep, close, open) that are affordable. The validation check of CARS-D uses the same attacks as the evaluation (query, gradient for neural models, exact for RLR, exhaustive-greedy for R-GCN and trees, and transfers) at the planning costs; its graph-free candidate is the counterpart itself (LR, the MLP, or HGB-X). A full CARS-D planning takes 2–90 seconds on one CPU core.

Attacks. Every test fraud node is attacked with per-node budget B, and we report the per-node minimum score over the applicable attacks, an upper bound on the true worst case. The *query* attack evaluates every allocation of the budget to one action, two actions, or all actions uniformly, with three endpoint rules for injections (the most benign-looking candidates, candidates scored like the target, and typical benign nodes from the interquartile range of benign scores, which defeat detectors that learn to distrust extremely benign neighbors), and keeps per target the most damaging configuration. The *gradient* attack relaxes the candidate edits of each target to $[0, 1]$, runs projected descent on the target logits with a per-target projection onto the cost polytope, and rounds randomly under the exact budget (Geisler et al., 2021); degree

gates are smoothed during the descent so that it sees activation and deactivation. Two *transfer* attacks apply the gradient edits computed on the clean R-GCN and the exact edits computed on the clean RLR. For RLR models, gated or not, the *exact* attack of Proposition 3 is added; it covers opening and closing gates, since a relation contributes zero whenever the post-attack degree is below the threshold. Hard-edge models (trees, CARE-GNN, GNNGuard, BWGNN) face query and transfer attacks. Fraudsters know their accomplices; with the label oracle removed (Table 49), injection endpoints are ranked by the detector's scores among all non-target nodes. The black-box attack transfers gradient edits from a surrogate R-GCN trained on a different labeled sample. Entity-level attacks on YelpChi reassign a review's user (aged or fresh account), rating group, or timing and rebuild every induced edge. Gangs share $\bar{B}|T|$ and allocate per-node budget levels with a greedy multiple-choice knapsack on the label-free estimate $\sum_v \hat{F}(s_v)$ of the AUC. The *exhaustive-greedy* attack, used for R-GCN and tree models in every evaluation and in every validation check, scores each candidate move of a target exactly: injecting any node of a candidate pool (3,000 nodes for R-GCN, 1,000 for trees: a third with the lowest scores, a third with the most extreme features, a third random) into any affordable relation, deleting any accomplice tie, and, for gated relations, injecting the best candidates until a gate opens or deleting accomplices until it closes. For R-GCN, exact single-target scoring recomputes the target's two-layer receptive field, including the change in an injected node's own aggregation when the target joins it; for trees, it recomputes the target's aggregate features. The most damaging move is applied while it lowers the score, after the first step among the 100 (R-GCN) or 50 (trees) best candidates per relation, and a target stops once its score is below the 0.5% quantile of benign scores. Transfers of its edits computed on the clean R-GCN and HGB-Graph are added for the other detectors. The *realized* joint attack applies, for every target, the edits of the attack that minimized its own score, all together, and rescores every node. *Certificates* follow Proposition 4: exact per-target worst cases for (gated) RLR; for R-GCN, per-unit exact layer-1 terms, interval bounds on the layer-1 states of neighbors and candidates over every way the target's edits change their neighborhoods, per-unit layer-2 terms, and enumeration of joint budget allocations (or a dynamic program over budget splits when they exceed 4,000); for trees, per-relation boxes of reachable neighbor means for every budget share, a covering set of budget splits, and the sum over trees of the smallest reachable leaf. Exact attacks discretize costs on the largest grid of step $0.5/2^j$ that represents them.

Compute. All experiments run on two CPU cores with 8 GB of memory. The second campaign, with the exhaustive-greedy attack in every evaluation, takes about 100 core-hours, most of it in the main table on YelpChi, Amazon, and S-FFSD (2–4.5 hours per seed); the third campaign (decomposition, hardening, and cost-uncertainty studies) takes about 35 core-hours, most of it in the plug-in planning of the 12-relation CSBMs (one to five hours per graph), whose per-group search enumerates gate actions and grows exponentially with the number of gated relations; the accompanying notebook regenerates every table and figure from cached records in minutes.

C. Threat Model, Protocol, and Scope

What an attacker edits. Relations in these datasets are generated by behavior, not edited directly: an edge appears when two reviews share a user, two firms disclose a supplier tie, or two transactions use the same card. An attacker therefore acts on future events that will generate edges. Table 4 lists, for every relation, the behavior behind an injection and a deletion and the rationale of the two cost profiles. Our edge-level action model abstracts these behaviors as counts with unit costs; on YelpChi, where the entities behind the relations are recoverable, the entity-level attacks of Table 50 rebuild every induced edge at once and confirm the edge-level conclusions. Costs are relative difficulties, not prices; only their ratio to the budget matters ($\Theta_{\alpha B}(\alpha c) = \Theta_B(c)$), and Figure 2 reports every conclusion over a continuum of cost scales rather than at two profiles.

Information used by each component. Table 5 lists the labels and structure each component sees. The graph structure and features are those of the full graph (transductive setting); no component uses test labels except for computing metrics. The attacker is strong by design: it knows the detector, its accomplices, and which labeled nodes are benign; Table 49 removes the label oracle and the white-box access.

Provenance of costs. No cost in this paper is estimated from operational data. We distinguish three kinds: *scenario* costs (P1, P2, Q1, Q2), chosen to represent cheap or verified identities and justified only by the behaviors of Table 4; *synthetic* costs (cost maps, random and noisy costs), used to test whether a conclusion survives other cost vectors; and *derived* quantities, such as the break-even cost of each relation (Table 32), which are computed from the data and tell a practitioner which cost would make a relation safe, without claiming what the cost is. Every conclusion that depends on costs is therefore a statement about a scenario or a cost region.

Table 4. Behavior behind each action and rationale of the profiles (P1 / P2).

Dataset	Relation	Edge means	Inject / delete	Profiles
YelpChi	R-U-R	same user	post from an account with history / use a fresh account per review	P1: fresh accounts are cheap; P2: identity verification makes both costly
	R-T-R	same product, month	review a product in a given month / shift timing	cheap to create, costly to avoid
	R-S-R	same product, rating	choose a rating / change it against the campaign's goal	cheap to create, costly to avoid
Amazon	U-P-U	common product	review a product / avoid products reviewed by accomplices	avoiding is costly
	U-S-U	same rating, same week	rate in a week / spread over time	moderate
	U-V-U	similar review text	write similar text / paraphrase	P1: paraphrasing is cheap; P2: text checks make both costly
FDCompCN	C-I-C	investment tie	disclose a holding / hold through undisclosed intermediaries	P1: concealable ownership; P2: transparent registries
	C-P-C, C-S-C	supplier, customer tie	disclose a trade tie / route trade through third parties	moderate
S-FFSD	T-S-T	same card	reuse a shared card / use a new card	P1: cards are cheap; P2: verified cards
	T-M-T, T-L-T, T-Y-T	merchant, location, type	transact at a merchant / abandon it	cheap to create, costly to avoid
Tolokers	W-T-W	common task	join a task / avoid tasks with accomplices	P1: unit costs; P2: costly participation control

Table 5. Information available to each component (tr, va, te: training, validation, test labels).

Component	Labels	Structure and features
Detector training	tr	clean full graph
CARS plug-in moments	tr	degrees and features of the full graph
CARS-D validation check, CARS+AT pick	va (attacked validation fraud)	clean full graph; attacks as below
Attack candidates for injections	labeled benign nodes (tr, va, te)	detector scores (white box)
Deletions	accomplices (fraud neighbors) of each target	–
Test metrics	te	attacked graph
Temporal split (S-FFSD)	earliest 40% / next 20% / last 40%	full graph, time feature removed

Scope of each claim. Table 6 separates the four kinds of statements the paper makes, and what each does not cover.

Table 6. Four kinds of claims and their scope.

Kind	What it guarantees	What it does not cover
Exact	worst case of linear relational models and CSBM thresholds under the formalized actions and costs (Proposition 3 and Theorems 1 and 2)	other detectors, other actions, misspecified costs
Certified	a lower bound on the true worst-case AUC of R-GCN and tree ensembles under the target-side model (Proposition 4)	cross effects between targets, actions outside the model; bounds can be loose
Statistical	the selected detector is at most 2ε below every candidate under the evaluated attacks, for i.i.d. validation nodes and separable evaluation (Proposition 5)	attacks not evaluated, distribution shift, dependence between targets, cost errors
Empirical	AUC under the strongest attacks we implemented (exhaustive-greedy, gradient, query, transfers, exact)	the true worst case of nonlinear detectors, which pool search shows can be lower
Operational	not claimed	real adversaries, calibrated costs, dated labels

Relation to the workshop version. Table 7 lists, for each contribution, the prior result, what is new, and what it makes possible; the workshop paper (Vincess & Melvin, 2026) studied the phenomenon on the CSBM, YelpChi, and Amazon. Table 8 lists the protocol under which every experiment was run.

D. Additional Results

All cells. Table 9 adds the P1, $B{=}20$ cell to Table 1. At this budget every relation of every real graph can be neutralized, and CARS-D falls back to graph-free features everywhere except where the check accepts a gate.

Attack strength and certified brackets. Table 10 reports, for clean-trained and CARS-D detectors at P1, $B{=}2$ and P2, $B{=}20$, the strongest single attack of the earlier suite (query, gradient, transfers), the exhaustive-greedy attack, the evaluated worst case over all attacks (per-node minimum), the AUC of a single realized joint attack, and the certified lower

Table 7. Contributions relative to the workshop version and to the closest prior work: what was known, what is new, and what the new result changes in the conclusions one can draw. The two main results are marked ⋆; the others are tools or supporting results.

Result	Workshop	Closest prior work	New here	What changes
Threat model, trust trap	introduced	cost-sensitive adversaries (Dalvi et al., 2004; Lowd & Meek, 2005; Zhang & Evans, 2019), strategic classification (Hardt et al., 2016), priced node attacks (Liao et al., 2024)	unchanged	–
Structure (Lemma 1), robust planning, maximin, hardening (Corollary 1)	cost monotonicity	edit-count budgets with equal costs (Zügner et al., 2018; Bojchevski & Günnemann, 2019)	per-node vs gang thresholds, nonconcavity, maximin choice under cost sets, hardening corollary	a budget-only analysis cannot rank relations for hardening; ours does, exactly for linear models (Table 25)
⋆ Exposure (Propositions 1 and 2) and abandonment (Theorem 2)	CSBM, first order; linear response	degree and robustness of message passing (Zügner & Günnemann, 2019)	per-node exposure of message passing; abandonment for every detector	exposure per unit cost predicts which relation to harden without any attack (15 of 16 cases)
⋆ Gating (Proposition 7, Theorem 1)	–	neighbor filtering and relation reweighting (Dou et al., 2020; Zhang et al., 2022; Zhang & Zitnik, 2020)	strict dominance of degree gates over every selection or reweighting under isolation	in the CSBM with isolated nodes, no selection or reweighting matches a degree gate, within a sharp window; on real graphs gating gains exactly where the conditions hold (Table 18)
Exact attack and certificates (Propositions 3 and 4)	–	certified graph robustness (Bojchevski & Günnemann, 2019; Zügner & Günnemann, 2019), tree verification (Chen et al., 2019)	exact thresholds for linear relational models under heterogeneous costs; brackets for R-GCN and trees	reported worst cases of nonlinear models are labeled tight, informative, or loose (Table 10)
Validated selection (Proposition 5)	–	evaluation against adaptive attacks (Tramèr et al., 2020)	finite-sample guarantee with radius growing in $\log K$	few theory-proposed candidates beat cost-blind or random candidates at equal budget on relation-rich graphs (Table 23)
Gangs (Proposition 6)	experiments only	coordinated injections (Choi et al., 2025)	convex envelope with a cross-effect bound	coordinated threat analysis
Attacks	query, gradient, transfer	PR-BCD (Geisler et al., 2021)	exhaustive-greedy with exact single-target scoring, gate moves, pool search	much stronger evaluation
Defense	cost-aware adversarial training	adversarial training (Madry et al., 2018)	plug-in, CARS, CARS-D, CARS-R, CARS+AT	a selection protocol that abstains safely

bound of Proposition 4. Table 11 samples 150 test fraud nodes per graph (P1, B=2, seed 0) and compares, per node, the earlier suite, the exhaustive-greedy attack, exhaustive search over every edit set within a pool of 24 candidates (a third most benign-looking, a third the best single injections, a third random) and all accomplice deletions, and the certificate; entries are each target's rank among benign test scores, averaged over targets (an AUC over the sampled targets).

Selection statistics and the validation guarantee. Table 12 counts the configurations kept by CARS-D over seeds, families, and cells, and compares the radius ε of Proposition 5 (K=3, δ=0.05) with the observed change between validation and test worst-case AUC of every candidate. Table 13 compares CARS-D with selection by adversarial validation alone: all relation subsets (VAL-SUB), {graph-free, all relations} (VAL-FA), and k random configurations of subsets and gates including the graph-free one (VAL-RAND-k), with the test-optimal configuration among all those evaluated (oracle, small graphs).

Degree groups and operating metrics. Table 14 splits test fraud nodes by total degree and reports the evaluated worst-case AUC of each group against all benign test nodes (P1, B=2); Table 15 reports average precision and true-positive rates at 1% and 5% false-positive rate under attack.

Gating against selection and reweighting. Table 16 evaluates the closed-form worst-case AUC of Theorem 1 (Gaussian CSBM moments, m and δ of the base configuration, p=0.8, q=0.02): the best gated detector against the best ungated one over $w \geq 0$, on a grid of weights. For $\pi_0 \geq 0.5$, the best ungated weight is $w = 0$ in 36 of 36 settings and the gated detector beats the relation-free one in 36. Table 17 repeats the comparison on CSBM graphs where a fraction of the nodes is isolated on the identity relation, with RLR and the exact attack (B=2, deletions on the identity relation made unaffordable). Clean-trained weights rely on the relation too heavily for any configuration to beat the relation-free detector; with the relation's weight scaled by a factor chosen on validation under the exact attack, the sup over w of Theorem 1, the gated detector (τ=3) beats both the best ungated one and the relation-free one by more than 0.002 in 3 of the four settings with c_{id}=1, where the attacker affords b=2 $< \tau$ injections as the theorem requires, and is within 0.002 in the remaining one; with c_{id}=0.5, b=4 $\geq \tau$, the gate can be opened, and gating is no better than dropping the relation.

Degree gating on real graphs. Tables 18 and 19 test Theorem 1 outside the CSBM. For every relation of every real graph and the CSBM, alone with the ego features, RLR is evaluated exactly for budgets $B \in \{1, 2, 5, 10, 20\}$ at P1 and P2 (separable worst-case AUC, the dynamic program of Proposition 3, which covers opening a gate by injection and closing it by deletion): graph-free LR, the ungated detector with its relation weight block scaled by $\lambda \in \{0, 0.1, 0.25, 0.5, 0.75, 1\}$

Table 8. Protocol of every experiment: full suite (with the exhaustive-greedy attack in evaluation and validation) or earlier suite.

Experiment	Protocol	Seeds / scope
Main table, brackets, intervals, selection, degree groups, metrics	full	5 seeds per graph (YelpChi, Amazon, S-FFSD: seeds 3–4 without CARE-GNN, GNNGuard, BWGNN, and adversarial training)
Validation-only selection	full	3 seeds (small graphs), 1 seed (large graphs)
Ablations, validation labels	full	3 seeds per graph
Decomposition at equal evaluation budget	full (tree attack pool 300)	real and augmented graphs: seed 0; relation-rich CSBMs: 5 seeds (R=4, 8), 3 seeds (R=12)
Which relation to harden	exact (RLR); full (HGB-Graph)	seed 0; relation-rich CSBMs: 5 seeds
Degree gating on real graphs, decisions under cost uncertainty, random cost distributions	exact attack	seed 0
Attack regret, certificates, exact-attack scaling	full	seed 0, 150 targets
Cost maps, noisy and cost-agnostic planning, misspecification	exact attack	seed 0 (misspecification: YelpChi)
Gangs, factorial study, gating check, adversarial training at matched compute	full	seed 0 (factorial: 2 seeds)
Temporal split and rolling windows (S-FFSD)	full	3 seeds; 4 windows
Sybil, poisoning, features, label scarcity, attacker knowledge	full	seed 0; YelpChi, Tolokers, FDCompCN
Entity-level attacks (YelpChi)	entity attack; CARS-D validated under the full suite	3 seeds
Budget sweeps, random cost vectors	full	seed 0
Phase diagram, gate-aware attack, stronger AT evaluation, sensitivity	earlier	Section E

(reweighting, λ=0 is dropping), and the gated detector with $\tau \in \{1, 2, 3, 5, 8\}$ and the same scaling; λ and τ are chosen on validation nodes at each budget. A win is a gated detector more than 0.002 above both the best reweighting and LR on test. The gate can be opened when $B \geq c_r^+ \tau$.

What each part of CARS-D contributes. Table 20 restricts the candidate set at equal attack budget (each validation check costs one attack evaluation per candidate): the plug-in's choice used without validation, validation between two of the three candidates, and CARS-D (five graphs, three seeds, two cells, three families). Table 21 measures whether CARS-D keeps a relation when one is worth keeping: a configuration *should* be kept when some evaluated configuration beats the graph-free one on test by more than 0.002 (Table 13). Table 22 repeats the validation check on random fractions of the validation labels (20 draws per fraction).

Decomposition at equal evaluation budget. Table 23 isolates each ingredient of CARS-D. Every candidate configuration that any rule may check is attacked once on all validation nodes and on the test set (full suite; for trees the exhaustive-greedy pool has 300 candidates); each rule then checks its candidates on the validation nodes and keeps the one of highest evaluated worst-case AUC, ties going to the simpler candidate. Rules: CARS-D {graph-free, plug-in subset, plug-in gates}; the same planner with every cost set to the geometric mean of the true costs (cost-blind planning); the plug-in's subset or gates without validation; CARS-D without the gated or without the graph-free candidate; VAL-FA {graph-free, all relations}; VAL-RAND-k, the graph-free configuration and $k-1$ random configurations of subsets and gates; VAL-SUB, every relation subset ($R \leq 4$); and VAL-POOL, every configuration evaluated for any rule. Graphs: the CSBM, FDCompCN, Tolokers, and Amazon (real; YelpChi and S-FFSD are covered by Table 13), FDCompCN, Tolokers, and Amazon with six attribute-sharing relations added (each node linked to three random nodes in its bin of about 50 nodes along a random feature, as when accounts share a device or an address; costs drawn log-uniformly in $[0.25, 16]$; augmented real), and CSBMs with $R \in \{4, 8, 12\}$ relations of random signal, degree, and isolation whose costs are drawn independently of their signal (relation-rich, five seeds for $R \in \{4, 8\}$ and three for R=12; Section B). Families RLR (exact attack) and HGB-Graph; two cells per graph. The planner searches subsets exhaustively up to six eligible relations and by local search beyond, so its number of plug-in evaluations grows linearly in R while VAL-SUB needs 2^R attack evaluations.

Table 9. Clean and evaluated worst-case test AUC at all cells (mean over seeds), with the conventions of Table 1.

	YelpChi				Amazon				FDCompCN				S-FFSD				Tolokers			
Model	cl.	P1,2	P1,20	P2,20	cl.	P1,2	P1,20	P2,20	cl.	P1,2	P1,20	P2,20	cl.	P1,2	P1,20	P2,20	cl.	P1,2	P1,20	P2,20
MLP	.836	.836	.836	.836	.975	.975	.975	.975	.715	.715	.715	.715	.788	.788	.788	.788	.740	.740	.740	.740
HGB-X	.935	.935	.935	.935	.977	.977	.977	.977	.789	.789	.789	.789	.845	.845	.845	.845	.734	.734	.734	.734
LR	.771	.771	.771	.771	.969	.969	.969	.969	.680	.680	.680	.680	.644	.644	.644	.644	.722	.722	.722	.722
Clean-trained, graph																				
RLR	.798	.181	.033	.546	.971	.865	.620	.674	.656	.001	.001	.001	.660	.355	.005	.061	.768	.529	.244	.529
R-GCN	.915	.043	.009	.639	.976	.739	.588	.594	.631	.003	.002	.002	.878	.325	.043	.143	.811	.649	.416	.648
HGB-Graph	.960	.442	.200	.924	.986	.971	.945	.951	.785	.473	.245	.250	.869	.551	.258	.385	.818	.524	.380	.524
CARE-GNN	.846	.773	.720	.808	.974	.967	.961	.963	.697	.026	.015	.014	.701	.654	.593	.624	.747	.705	.622	.705
GNNGuard-style	.917	.169	.080	.850	.976	.912	.867	.880	.652	.014	.011	.011	.868	.567	.294	.529	.808	.637	.459	.636
BWGNN-style	.902	.593	.425	.868	.978	.967	.947	.946	.688	.093	.082	.082	.867	.715	.402	.654	.831	.729	.661	.726
Adversarially trained																				
CA-AT (R-GCN)	.866	.679	.591	.834	.976	.892	.877	.881	.644	.229	.166	.120	.853	.686	.400	.628	.819	.734	.603	.734
AT-RLR (exact)	.773	.366	.172	.597	.967	.935	.887	.901	.619	.014	.012	.012	.645	.562	.476	.529	.737	.698	.614	.697
CA-HGB	.962	.919	.710	.944	.987	.979	.960	.964	.784	.748	.521	.540	.864	.702	.532	.639	.818	.680	.469	.680
Ours																				
RLR + CARS	.771	**.771**	**.771**	**.791**	.972	.951	**.969**	**.969**	.680	.674	**.680**	**.680**	.644	.504	**.644**	**.664**	.722	**.722**	**.722**	**.722**
R-GCN + CARS	.836	**.836**	**.836**	**.897**	.976	.943	**.975**	**.975**	.671	.669	**.715**	**.715**	.843	.328	**.788**	**.858**	.740	**.740**	**.740**	**.740**
HGB + CARS	.935	**.935**	**.935**	**.952**	.982	**.976**	**.977**	**.977**	.788	.787	**.789**	**.789**	.852	.449	**.845**	**.863**	.734	**.734**	**.734**	**.734**
RLR + CARS-D	.771	**.771**	**.771**	**.792**	.969	**.969**	**.969**	**.969**	.682	**.678**	**.680**	**.680**	.644	**.644**	**.644**	**.664**	.722	**.722**	**.722**	**.722**
R-GCN + CARS-D	.836	**.836**	**.836**	**.897**	.975	**.975**	**.975**	**.975**	.715	**.715**	**.715**	**.715**	.788	**.788**	**.788**	**.858**	.746	**.741**	**.740**	**.741**
HGB + CARS-D	.935	**.935**	**.935**	**.952**	.979	**.975**	**.977**	**.977**	.790	**.789**	**.789**	**.789**	.845	**.845**	**.845**	**.863**	.734	**.734**	**.734**	**.734**
CARS+AT (R-GCN)	.837	**.837**	**.837**	**.896**	.973	**.973**	**.973**	**.973**	.715	**.715**	**.715**	**.715**	.782	.782	.782	**.856**	.746	**.741**	**.740**	**.741**
CARS+AT (HGB)	.935	**.935**	**.935**	**.952**	.979	**.976**	**.977**	**.977**	.790	**.789**	**.789**	**.789**	.845	**.845**	**.845**	**.863**	.734	**.734**	**.734**	**.734**

Table 24 reports the same comparisons per graph, as paired differences of CARS-D with each alternative and 95% bootstrap intervals (resampling seeds, or cases when there is one seed; cases are cells and families, so intervals over four cases are only indicative).

Which relation to harden. Table 25 multiplies the injection and deletion costs of one relation at a time by 1000 (unmanipulable within the budget grid) and computes the exact trap threshold of the clean-trained RLR with one dynamic program per relation (Proposition 3), on every graph, seed, and profile where RLR has a clean gain over LR and is trapped within $B \leq 64$: the CSBM (P1, P2), YelpChi (P2), and CSBMs with 4, 8, and 12 relations and random costs (five seeds each). A rule is right if it picks a relation whose hardening raises the threshold most. Rules: the first-order exposure per unit cost of Propositions 1 and 2, measured as the largest score drop of one injection or one deletion per unit cost, averaged over fraudulent test nodes (no attack is run); the exact analysis under equal costs (the geometric mean), as a budget-only robustness analysis would do; the most informative relation (largest validation AUC loss when dropped); the largest weight norm; and the cheapest relation. Value: threshold gain of the pick relative to the best. 85% of the 146 single-relation hardenings leave the threshold unchanged, as Corollary 1(c) predicts for relations off the cheapest trap policy. For HGB-Graph (full suite, B=2; the CSBM at P2, YelpChi at P2, and three relation-rich CSBMs), making the RLR-exposure pick unmanipulable is also the best single hardening for the trees in 2 of 5 graphs and raises their evaluated worst-case AUC by +.089 on average: the first-order ranking transfers only partly to nonlinear detectors.

Decisions under cost uncertainty. Table 26 gives the defender a cost estimate $\hat{c} = c\,e^{\epsilon}$ ($\epsilon \sim \mathcal{N}(0,1)$ per coordinate, four draws) and a box of log half-width w around it, on YelpChi, Amazon, FDCompCN, Tolokers, and the CSBM (P1, B=2 and P2, B=20). Naive: RLR + CARS-D planned at $\hat{c}$. Maximin: planned at the lower corner of the box, which is the maximin choice by Corollary 1(b). Minimax regret: among the relation subsets of RLR (graph-free, singletons, all), the one of smallest maximum regret over the lower corner, the centre, and eight sampled corners of the box. Every choice is evaluated exactly at the true costs. Coverage: share of draws in which the box contains the true costs.

Stability under plausible cost distributions. Table 28 replaces the two profiles by 20 random cost vectors from each of two distributions. *Independent*: every injection and deletion cost log-uniform in $[0.25, 16]$. *Structured*, following the behaviors of Section C: the identity-like relation is expensive to enter (a new or verified account, $[4, 32]$) and cheap to leave ($[0.25, 1]$), and the other relations are cheap to enter ($[0.25, 2]$) and harder to leave, since a shared product, merchant, or task cannot be unshared without changing behavior ($[1, 8]$). For each vector and $B \in \{2, 20\}$, the clean-trained RLR and RLR + CARS-D planned at the true costs are evaluated exactly. The questions are whether the paper's qualitative conclusions survive: is the clean model trapped, does CARS-D fall below LR, and when does it gain. The trap persists in 95% of the

Table 10. Evaluated and certified worst-case AUC (mean over seeds): clean AUC, strongest earlier attack, exhaustive-greedy (EG), all attacks (per-node minimum), single realized joint attack, certified lower bound, and the status of the bracket [certified, all]: tight (width ≤ 0.01, the true worst case is known to within 0.01), informative (≤ 0.05), or loose (the true worst case is not pinned down and only the evaluated upper bound is reported).

Dataset	Cell	Model	clean	earlier	EG	all	realized	certified	bracket
YelpChi	P1, B=2	R-GCN	.915	.112	.042	.043	.043	.000	informative
YelpChi	P1, B=2	HGB-Graph	.960	.659	.439	.442	.438	.070	loose
YelpChi	P1, B=2	R-GCN + CARS-D	.836	.836	–	.836	.836	.836	tight
YelpChi	P1, B=2	HGB + CARS-D	.935	.935	.935	.935	.935	.935	tight
YelpChi	P2, B=20	R-GCN	.915	.798	.632	.639	.632	.000	loose
YelpChi	P2, B=20	HGB-Graph	.960	.944	.927	.924	.922	.567	loose
YelpChi	P2, B=20	R-GCN + CARS-D	.897	.897	.897	.897	.897	.836	loose
YelpChi	P2, B=20	HGB + CARS-D	.952	.952	.952	.952	.952	.952	tight
Amazon	P1, B=2	R-GCN	.976	.879	.743	.739	.742	.024	loose
Amazon	P1, B=2	HGB-Graph	.986	.980	.971	.971	.971	.914	loose
Amazon	P1, B=2	R-GCN + CARS-D	.975	.975	–	.975	.975	.975	tight
Amazon	P1, B=2	HGB + CARS-D	.979	.977	.975	.975	.975	.962	informative
Amazon	P2, B=20	R-GCN	.976	.786	.598	.594	.591	.003	loose
Amazon	P2, B=20	HGB-Graph	.986	.973	.950	.951	.949	.840	loose
Amazon	P2, B=20	R-GCN + CARS-D	.975	.975	–	.975	.975	.975	tight
Amazon	P2, B=20	HGB + CARS-D	.977	.977	.977	.977	.977	.977	tight
FDCompCN	P1, B=2	R-GCN	.631	.006	.003	.003	.002	.000	tight
FDCompCN	P1, B=2	HGB-Graph	.785	.621	.469	.473	.469	.387	loose
FDCompCN	P1, B=2	R-GCN + CARS-D	.715	.715	–	.715	.715	.715	tight
FDCompCN	P1, B=2	HGB + CARS-D	.790	.789	.789	.789	.789	.788	tight
FDCompCN	P2, B=20	R-GCN	.631	.002	.002	.002	.001	.000	tight
FDCompCN	P2, B=20	HGB-Graph	.785	.549	.239	.250	.239	.142	loose
FDCompCN	P2, B=20	R-GCN + CARS-D	.715	.715	–	.715	.715	.715	tight
FDCompCN	P2, B=20	HGB + CARS-D	.789	.789	.789	.789	.789	.789	tight
S-FFSD	P1, B=2	R-GCN	.878	.445	.361	.325	.326	.000	loose
S-FFSD	P1, B=2	HGB-Graph	.869	.677	.555	.551	.547	.134	loose
S-FFSD	P1, B=2	R-GCN + CARS-D	.788	.788	–	.788	.788	.788	tight
S-FFSD	P1, B=2	HGB + CARS-D	.845	.845	.845	.845	.845	.845	tight
S-FFSD	P2, B=20	R-GCN	.878	.338	.141	.143	.134	.000	loose
S-FFSD	P2, B=20	HGB-Graph	.869	.414	.466	.385	.366	.058	loose
S-FFSD	P2, B=20	R-GCN + CARS-D	.858	.858	.858	.858	.858	.811	informative
S-FFSD	P2, B=20	HGB + CARS-D	.863	.863	.863	.863	.863	.863	tight
Tolokers	P1, B=2	R-GCN	.811	.707	.646	.649	.644	.349	loose
Tolokers	P1, B=2	HGB-Graph	.818	.632	.529	.524	.521	.222	loose
Tolokers	P1, B=2	R-GCN + CARS-D	.746	.744	.735	.741	.741	.727	informative
Tolokers	P1, B=2	HGB + CARS-D	.734	.734	.734	.734	.734	.734	tight
Tolokers	P2, B=20	R-GCN	.811	.693	.646	.648	.644	.349	loose
Tolokers	P2, B=20	HGB-Graph	.818	.632	.529	.524	.521	.222	loose
Tolokers	P2, B=20	R-GCN + CARS-D	.746	.744	.735	.741	.741	.727	informative
Tolokers	P2, B=20	HGB + CARS-D	.734	.734	.734	.734	.734	.734	tight

320 cases where RLR has a clean gain. CARS-D gains in 90 of 480 cases, most often on YelpChi under the structured distribution, where the identity relation is expensive to enter, and falls below LR in 46, by at most $-.023$; 34 of these are on FDCompCN, whose relations carry no clean signal, so a validation check that keeps one is wrong by a margin within the radius 2ε of Proposition 5 but not zero.

Confirmatory comparisons in full. Table 29 lists the 45 confirmatory comparisons with unrounded mean differences, the standard deviation of the per-seed difference, the hierarchical-bootstrap interval, the Holm-adjusted p-value, non-inferiority at a margin of 0.01 AUC (lower end of the interval above -0.01), and the smallest margin at which non-inferiority holds. Exact ties (difference 0) are cases in which CARS-D returned the graph-free detector in every seed. The margin of 0.01 was not fixed before the analysis; the last column lets a reader apply any margin.

Planning without cost estimates. Table 30 plans RLR + CARS-D at the lower corner of the cost grid (every cost a quarter of P1), the plan of a defender who knows only a lower bound on costs, and evaluates it exactly at every cost scale; it compares it with the plan at the true costs of Figure 2. Table 31 plans at noisy cost estimates $\hat{c} = c\,e^{\epsilon}$ and at the lower corner $\hat{c}e^{-z\sigma}$ of a log box of half-width $z\sigma$, and evaluates at the true costs. Table 32 reports, for each relation with a clean gain, its exact single-relation threshold at P1 and the factor by which its costs would have to grow for it to resist budgets of 2 and 20 (by homogeneity, Lemma 1).

Table 11. Attack regret (P1, B=2, seed 0, 150 sampled targets): AUC over the sampled targets when clean, under the earlier suite, the exhaustive-greedy attack, exhaustive search within the pool (combined with the other attacks), and the certified bound; share of targets where the exhaustive-greedy attack beats the suite and where pool search beats the exhaustive-greedy attack.

Dataset	Model	n	clean	suite	EG	pool	certified	EG<suite	pool<EG
Amazon	HGB-CARS-D	150	.984	.979	.975	.974	.944	71%	29%
Amazon	HGB-Graph	150	.988	.984	.978	.977	.934	92%	68%
Amazon	R-GCN	150	.969	.857	.711	.616	.016	99%	97%
Amazon	R-GCN-CARS-Dp	150	.976	.973	.970	.970	.894	78%	27%
FDCompCN	HGB-CARS-D	150	.783	.783	.784	.783	.777	0%	3%
FDCompCN	HGB-Graph	150	.776	.571	.387	.384	.335	99%	45%
FDCompCN	R-GCN	150	.615	.013	.005	.003	.000	89%	93%
FDCompCN	R-GCN-CARS-Dp	150	.637	.635	.635	.635	.630	0%	11%
S-FFSD	HGB-CARS-Dp	150	.881	.657	.401	.383	.124	97%	27%
S-FFSD	HGB-Graph	150	.882	.629	.516	.433	.084	73%	42%
S-FFSD	R-GCN	150	.876	.400	.338	.250	.000	70%	77%
S-FFSD	R-GCN-CARS-Dp	150	.864	.684	.597	.542	.000	84%	87%
Tolokers	HGB-CARS-Dp	150	.744	.713	.680	.676	.573	39%	15%
Tolokers	HGB-Graph	150	.804	.659	.508	.491	.213	95%	51%
Tolokers	R-GCN	150	.784	.675	.639	.631	.355	94%	69%
Tolokers	R-GCN-CARS-D	150	.741	.725	.715	.714	.650	36%	33%
YelpChi	HGB-CARS-Dp	150	.942	.873	.837	.834	.802	37%	27%
YelpChi	HGB-Graph	150	.947	.685	.455	.404	.100	99%	97%
YelpChi	R-GCN	150	.907	.098	.048	.040	.000	95%	77%
YelpChi	R-GCN-CARS-Dp	150	.870	.755	.730	.723	.628	38%	16%

Table 12. Configurations kept by CARS-D (share of runs) and the validation guarantee: mean ε, 2ε, largest and mean validation-to-test change over candidates.

Dataset	runs	free	CARS	gated	ε	2ε	max $\lvert\Delta\rvert$	mean Δ
YelpChi	45	67%	29%	4%	.046	.092	.019	−.001
Amazon	45	96%	0%	4%	.127	.254	.042	.002
FDCompCN	45	87%	7%	7%	.155	.309	.081	−.028
S-FFSD	45	67%	33%	0%	.053	.105	.048	.001
Tolokers	45	96%	0%	4%	.077	.155	.027	−.007

Rolling temporal windows. Table 33 cuts the time-ordered labeled transactions of S-FFSD into ten blocks and, for windows $w = 0, \ldots, 3$, trains on blocks w to $w{+}3$, validates on block $w{+}4$, and tests on blocks $w{+}5$ and $w{+}6$; it is a case study on simulated data, not evidence of robustness to strategic drift.

Cost maps and noisy costs. Figure 2 scales the P1 injection and deletion costs of every relation by factors in $\{1/4, 1, 4, 16\}$ and reports the exact worst-case AUC (separable and joint attacks with the exact attack) of the clean-trained RLR and of RLR + CARS-D planned at the true costs, minus LR. Table 34 plans RLR + CARS-D at costs multiplied by independent or common log-normal factors and evaluates it exactly at the true P1 costs.

Robust planning. Table 35 plans HGB + CARS-D at either profile with robustness factor ρ and evaluates it under both (seed 0). By Corollary 1, the guarantee of a plan at costs c/ρ covers every true cost vector above c/ρ. On YelpChi at B=20, the P2 plan collapses when identities are in fact cheap, and $\rho \geq 2$ restores graph-free performance under both profiles; at B=2, even ρ=8 leaves an R-U-R edit unaffordable in the plan ($25/8 > 2$) while P1 makes it cheap, so P1 lies outside the covered box and the plan fails, as the corollary allows.

Exact thresholds of relation subsets. Table 36 computes, with one exact dynamic program per relation subset, the smallest budget at which the separable worst-case AUC of RLR restricted to the subset falls to that of LR (seed 0).

Gangs. Table 37 lets a gang share an average budget $\bar{B}$ per member; each attacked member also pays a coordination cost κ, and the gang allocates its budget with the detector's own scores (oracle) or with a surrogate trained on another labeled sample. The realized AUC applies all edits together and rescores every node; its largest deviation from the convex-envelope prediction is .009.

Causal factors. Table 38 regresses the trap depth (evaluated worst-case AUC minus the graph-free counterpart, B=2) on the CSBM factors: the identity relation's homophily $p_{\mathrm{id}} \in \{0.55, 0.7, 0.85\}$, the fraction of nodes isolated on it $\in \{0, 0.5, 0.85\}$, its cost $c_{\mathrm{id}} \in \{0.5, 4\}$, and the architecture (two seeds).

Table 13. Test worst-case AUC of the selected detector: plug-in only (CARS), CARS-D, and validation-only rules; last column: validation evaluations of CARS-D / VAL-SUB.

Dataset	Family	Cell	CARS	CARS-D	VAL-FA	RAND-3	RAND-8	SUB	oracle	evals
Amazon	HGB	P1, B=2	.973	.974	.976	.976	.973	.973	.977	3 / 8
Amazon	HGB	P2, B=20	.976	.976	.976	.976	.976	.976	.976	3 / 8
Amazon	R-GCN	P1, B=2	.897	.971	.971	.971	.971	.971	.972	3 / 8
Amazon	R-GCN	P2, B=20	.971	.971	.971	.971	.971	.971	.971	3 / 8
Amazon	RLR	P1, B=2	.930	.969	.969	.969	.967	.969	.969	3 / 8
Amazon	RLR	P2, B=20	.969	.969	.969	.969	.969	.969	.969	3 / 8
FDCompCN	HGB	P1, B=2	.781	.781	.781	.781	.780	.780	.784	3 / 8
FDCompCN	HGB	P2, B=20	.781	.781	.781	.781	.781	.781	.781	3 / 8
FDCompCN	R-GCN	P1, B=2	.666	.718	.718	.716	.716	.718	.722	3 / 8
FDCompCN	R-GCN	P2, B=20	.718	.718	.718	.718	.718	.718	.718	3 / 8
FDCompCN	RLR	P1, B=2	.668	.674	.679	.679	.677	.674	.680	3 / 8
FDCompCN	RLR	P2, B=20	.679	.679	.679	.679	.679	.679	.679	3 / 8
S-FFSD	HGB	P1, B=2	.470	.842	.842	.842	.842	.842	–	3 / 16
S-FFSD	HGB	P2, B=20	.860	.860	.842	.842	.864	.860	–	3 / 16
S-FFSD	RLR	P1, B=2	.546	.643	.643	.643	.643	.643	–	3 / 16
S-FFSD	RLR	P2, B=20	.661	.661	.643	.643	.699	.661	–	3 / 16
Tolokers	HGB	P1, B=2	.733	.733	.733	.733	.733	.733	.733	3 / 2
Tolokers	HGB	P2, B=20	.733	.733	.733	.733	.733	.733	.733	3 / 2
Tolokers	R-GCN	P1, B=2	.737	.739	.737	.737	.745	.737	.745	3 / 2
Tolokers	R-GCN	P2, B=20	.737	.739	.737	.741	.745	.737	.745	3 / 2
Tolokers	RLR	P1, B=2	.723	.723	.723	.723	.723	.723	.724	3 / 2
Tolokers	RLR	P2, B=20	.723	.723	.723	.723	.723	.723	.724	3 / 2
YelpChi	HGB	P1, B=2	.933	.933	.933	.933	.935	.933	–	3 / 8
YelpChi	HGB	P2, B=20	.948	.948	.933	.933	.950	.948	–	3 / 8
YelpChi	RLR	P1, B=2	.771	.771	.771	.773	.792	.771	–	3 / 8
YelpChi	RLR	P2, B=20	.794	.794	.771	.771	.795	.794	–	3 / 8

Table 14. Evaluated worst-case AUC by total degree of the fraud node (P1, B=2, mean over seeds).

Dataset	Degree	n_+	MLP	R-GCN	R-GCN+D	HGB-X	HGB-Graph	HGB+D
YelpChi	1–2	8	–	.017	.717	–	.407	.834
YelpChi	3–9	47	–	.025	.816	–	.556	.934
YelpChi	10+	2613	–	.044	.837	–	.440	.935
Amazon	10+	327	–	.742	.975	–	.971	.975
FDCompCN	0	130	–	.002	.763	–	.505	.823
FDCompCN	1–2	69	–	.003	.662	–	.434	.749
FDCompCN	3–9	21	–	.005	.573	–	.416	.708
S-FFSD	10+	2101	–	.326	.788	–	.551	.845
Tolokers	1–2	82	–	.103	.734	–	.145	.743
Tolokers	3–9	136	–	.420	.739	–	.263	.751
Tolokers	10+	807	–	.743	.742	–	.607	.730

Adversarial training at matched compute. Table 39 trains CA-AT for a fraction of the standard number of epochs, including the number matched to the wall time of CARS-D, and evaluates it with the full suite.

Ablation of CARS-D. Table 41 reports, per detector family and dataset, the mean margin over the graph-free counterpart of the plug-in subset (CARS), of the plug-in gated proposal, and of the configuration kept by the validation check (CARS-D). The proposals win where the plug-in is accurate (Tolokers, R-GCN; the synthetic graph) and lose where it is not (Amazon and S-FFSD at P1); the check keeps the better configuration in all but a few runs, whose cost is within validation noise.

Temporal split. On S-FFSD, models are trained on the earliest 40% of labeled transactions, validated on the next 20%, and tested on the last 40%, without the time stamp as a feature (three seeds, averaged).

Node-level exposure. For a one-layer relational network, Table 43 compares the node-level bound of Proposition 2, computed with each target's degree and its actual candidate endpoints, with the measured score drop of the best single edit per unit cost (seed 0, up to 400 targets). The bound is never violated (0 violations out of 15590 node–relation–action triples); for injections, unlike the dataset-level bound with mean degrees, it ranks nodes by their measured exposure (Spearman up to 0.95), whereas for deletions, whose damage depends on which accomplices exist, it does not.

Thresholds and weights in the CSBM. Table 44 lists, for every identity cost, the numerical trap threshold, the lower envelope of the trap-attack lines, the exposure lower bound, the linear-response threshold, and the abandonment threshold with its bound. The best response is computed by multi-start SLSQP; a dense random search of 20,000 feasible attacks per

Table 15. Operating metrics under the evaluated worst case (P1, B=2, mean over seeds).

Dataset	Model	AUC	AP	TPR@1%	TPR@5%
YelpChi	MLP	.836	.525	.178	.397
YelpChi	HGB-X	.935	.783	.458	.706
YelpChi	R-GCN	.043	.077	.001	.001
YelpChi	HGB-Graph	.442	.120	.003	.017
YelpChi	R-GCN-CARS-D	.836	.525	.178	.397
YelpChi	HGB-CARS-D	.935	.783	.458	.706
YelpChi	CA-HGB	.919	.723	.346	.628
YelpChi	CARS+AT-HGB	.935	.783	.458	.706
Amazon	MLP	.975	.889	.818	.902
Amazon	HGB-X	.977	.906	.822	.909
Amazon	R-GCN	.739	.566	.426	.550
Amazon	HGB-Graph	.971	.889	.784	.886
Amazon	R-GCN-CARS-D	.975	.889	.818	.902
Amazon	HGB-CARS-D	.975	.902	.810	.906
Amazon	CA-HGB	.979	.897	.812	.908
Amazon	CARS+AT-HGB	.976	.901	.812	.905
FDCompCN	MLP	.715	.282	.074	.250
FDCompCN	HGB-X	.789	.413	.167	.416
FDCompCN	R-GCN	.003	.055	.000	.000
FDCompCN	HGB-Graph	.473	.121	.009	.075
FDCompCN	R-GCN-CARS-D	.715	.282	.074	.250
FDCompCN	HGB-CARS-D	.789	.412	.169	.410
FDCompCN	CA-HGB	.748	.346	.127	.341
FDCompCN	CARS+AT-HGB	.789	.412	.169	.410
S-FFSD	MLP	.788	.582	.314	.417
S-FFSD	HGB-X	.845	.678	.430	.504
S-FFSD	R-GCN	.325	.124	.001	.004
S-FFSD	HGB-Graph	.551	.296	.131	.195
S-FFSD	R-GCN-CARS-D	.788	.582	.314	.417
S-FFSD	HGB-CARS-D	.845	.678	.430	.504
S-FFSD	CA-HGB	.702	.509	.356	.398
S-FFSD	CARS+AT-HGB	.845	.678	.430	.504
Tolokers	MLP	.740	.385	.032	.137
Tolokers	HGB-X	.734	.388	.027	.154
Tolokers	R-GCN	.649	.312	.012	.105
Tolokers	HGB-Graph	.524	.222	.008	.051
Tolokers	R-GCN-CARS-D	.741	.391	.035	.147
Tolokers	HGB-CARS-D	.734	.388	.027	.154
Tolokers	CA-HGB	.680	.323	.021	.107
Tolokers	CARS+AT-HGB	.734	.388	.027	.154

Table 16. Closed-form worst-case AUC, best gated / best ungated linear detector, by degree d of connected nodes, injections b, and isolated fraction π_0; last column: relation-free detector.

d	b	π_0=0	0.25	0.5	0.75	0.9	free
2.0	1.0	.715 / .715	.708 / .690	.702 / .690	.696 / .690	.692 / .690	.690
4.0	1.0	.772 / .772	.752 / .697	.731 / .690	.711 / .690	.698 / .690	.690
8.0	2.0	.824 / .824	.791 / .701	.757 / .690	.724 / .690	.703 / .690	.690
16.0	4.0	.884 / .884	.837 / .706	.789 / .690	.740 / .690	.710 / .690	.690

instance followed by local polishing never improves on it by more than 3×10^{-12} in AUC across 48 instances. Figure 3 shows the minimax weights, which are clipped as Theorem 2(b) predicts, and the cost misspecification of the minimax detector: rows never decrease for true costs above the assumed one, as Corollary 1 implies.

Fidelity of the plug-in. For every relation subset and single-relation gate, Table 45 compares the plug-in worst case of the linear detector with the exact worst case of the gated RLR retrained on the configuration (seed 0): rank correlation with the test worst case, rank correlation between validation and test worst cases, the test regret of picking the plug-in's best configuration or the validation's best, and the mean overestimate of the plug-in.

Statistical reliability and operating points. Table 46 reports, per dataset and cell, the mean paired difference over seeds between our detectors and their graph-free counterparts with its 95% t interval, the number of seeds at or above the counterpart (within 0.002), and the average precision and true-positive rate at 1% false-positive rate under attack.

Compute. Table 47 compares the median time of one CARS-D planning, with the training and validation check of the three detector families, with the training of one CA-AT model at the same cell (one CPU core).

Table 17. CSBM graphs with an isolated fraction of nodes on the identity relation: exact worst-case AUC of RLR ungated and gated at τ=3, with clean-trained weights or with the identity weight scaled on validation (robust), and without the identity relation (B=2, seed 0).

p_{id}	isolated	c_{id}	ungated	ung. robust	dropped	gated	gated robust
0.7	0.5	0.5	.145	.712	.717	.148	.716
0.7	0.5	1	.248	.712	.717	.635	.730
0.7	0.85	0.5	.028	.716	.717	.029	.717
0.7	0.85	1	.060	.716	.717	.692	.722
0.85	0.5	0.5	.132	.712	.717	.133	.716
0.85	0.5	1	.243	.712	.717	.630	.733
0.85	0.85	0.5	.032	.712	.717	.035	.714
0.85	0.85	1	.064	.712	.717	.692	.719

Table 18. Validated degree gating against validated reweighting and dropping, per relation on real graphs and the CSBM (exact attack, five budgets, two profiles): cases, wins, ties (≤ 0.002), losses, and mean and largest gain, by whether the attacker can open the chosen gate.

Regime	cases	wins	ties	losses	mean	max
gate cannot be opened ($B < c_r^+ \tau$)	80	14	53	13	+.004	+.061
gate can be opened	90	0	81	9	.000	+.001
all	170	14	134	22	+.002	+.061

Budget sweeps (full suite). Figure 4 traces the evaluated worst case of clean-trained models and of RLR and HGB with CARS-D, planned at each budget, from B=1 to 50 at P1. With the full suite, the clean-trained curves cross the graph-free lines at budgets of one to two edits on every graph, except trees on Amazon, which cross at B=5; the CARS-D curves stay within 0.002 of them except RLR + CARS-D on FDCompCN at $B \leq 2$ (0.014 below, within the validation radius).

Random cost vectors (full suite). Table 48 draws 24 cost vectors log-uniformly in $[0.25, 16]^{2R}$ per dataset (B=5, seed 0) and compares the observed trap of RLR (exact attack) and HGB-Graph with the plug-in prediction (the plug-in worst case of the full linear detector below that of the graph-free one).

Further threat variations under the full suite. The following checks were rerun with the exhaustive-greedy attack in evaluation and in the validation check of CARS-D (seed 0; YelpChi, Tolokers, FDCompCN; entity-level attacks on YelpChi with three seeds).

Attacker knowledge. Table 49 compares the white-box attacker with labels (default), the white-box attacker without the label oracle, and a black-box attacker who transfers edits from a surrogate R-GCN trained on its own labeled sample (seed 0).

Entity-level attacks. Table 50 attacks YelpChi detectors through entities: a review is reassigned to another (aged) account, a fresh account, another rating group, or another posting time, and every induced edge is rebuilt (three seeds). CARS-D is planned with the edge-level profile that matches the entity profile (Q1 with P1, Q2 with P2). Under Q1, entity actions trap R-GCN and HGB-Graph at both budgets and RLR at B=5, as edge actions do under P1; each CARS-D variant is within 0.002 of its graph-free counterpart under Q1 (LR, not in the table, scores .771 on YelpChi; MLP .837; HGB-X .935) and above it under Q2.

Persistent camouflage. Before training, every fraud node (training, validation, and test) applies the exact attack of a clean RLR surrogate at P1, B=2, and the edits stay in the graph; detectors and CARS are then built on the poisoned graph, and test fraud nodes attack again with the same budget (seed 0). When the camouflage edges concentrate on the surrogate's most benign-looking nodes (c), they leave a signature that graph models detect almost perfectly; when they are diversified over benign nodes (d), the trap persists. In both cases, attacked fraud nodes in the training set act as adversarial augmentation.

Sybil accounts. The attacker creates 1,000 fake accounts that copy the features of random labeled benign nodes with small noise and connects them to targets; each Sybil edge costs c_s on every relation, and deletions keep their P1 costs. Detectors are trained without the Sybils (seed 0).

Feature and structure attacks. Fraudsters may also perturb their own standardized features within an ℓ_∞ ball of radius ϵ; graph-free detectors face this feature attack, and graph detectors face it followed by the structural suite at P1, B=2 on

Table 19. Degree gating per relation: share of isolated test fraud nodes, wins/losses/cases at P1 and P2, and largest gain.

Graph	Relation	isolated	P1	P2	max
Amazon	U-P-U	15%	0/0/5	0/0/5	.000
Amazon	U-S-U	1%	0/0/5	0/0/5	+.001
Amazon	U-V-U	4%	0/0/5	0/0/5	+.001
FDCompCN	C-I-C	72%	0/3/5	0/5/5	−.001
FDCompCN	C-P-C	88%	0/0/5	0/0/5	.000
FDCompCN	C-S-C	88%	0/0/5	0/0/5	−.001
S-FFSD	T-S-T	26%	1/0/5	5/0/5	+.061
S-FFSD	T-M-T	0%	0/0/5	0/0/5	.000
S-FFSD	T-L-T	0%	0/0/5	0/0/5	.000
S-FFSD	T-Y-T	0%	0/0/5	0/0/5	.000
Synthetic	identity	0%	0/0/5	0/0/5	.000
Synthetic	behavior	0%	0/3/5	0/3/5	.000
Synthetic	dense	0%	0/0/5	0/0/5	.000
Tolokers	W-T-W	0%	0/5/5	0/2/5	.000
YelpChi	R-U-R	82%	3/1/5	5/0/5	+.021
YelpChi	R-T-R	1%	0/0/5	0/0/5	.000
YelpChi	R-S-R	0%	0/0/5	0/0/5	.000

Table 20. Ablation at equal attack budget: number of candidates, mean and worst test difference with the graph-free candidate, runs below it, and mean regret to the test-optimal candidate.

Rule	cands	mean Δ	worst Δ	below	regret
plug-in subset (CARS)	1	−.028	−.499	23/90	.036
plug-in gates	1	−.054	−.579	56/90	.061
val: free/subset	2	.007	−.014	2/90	.001
val: free/gates	2	.007	−.014	3/90	.000
val: subset/gates	2	−.018	−.372	21/90	.025
CARS-D	3	.007	−.014	3/90	.000

the perturbed graph (seed 0). Feature attacks hurt every detector; they hurt graph-free ones more when the structure is untouched, which shows that relations do buy robustness against feature manipulation as long as they are expensive to edit.

Label scarcity. Table 54 trains with 5%, 10%, and 20% of the evaluated nodes (validation 20%, test the rest; seed 0).

E. Earlier Robustness Checks

The checks in this section (gate-aware attacks, stronger evaluation of adversarial training, sensitivity, and the phase diagram) were run before the exhaustive-greedy attack was added, with the earlier suite (query, gradient, and transfer attacks, and the exact attack for linear models); Table 8 lists every experiment and its protocol. The exhaustive-greedy attack includes gate moves, which supersedes the gate-aware attack, and the stronger evaluation of adversarial training is superseded by the main table and Table 39. For clean-trained nonlinear models, the stronger attack can only lower the reported worst cases, so the traps shown here are, if anything, understated; rows of CARS-D and adversarially trained models were not revalidated under the exhaustive-greedy attack and should be read as evidence about the earlier threat model only.

Gate-aware attacks. For every gated configuration that CARS-D proposed or kept, Table 55 adds an attack that targets the gates: for every combination of keep, open (inject the missing edges), and close (delete accomplices until the degree falls below the threshold) over the gated relations, the remaining budget is spent on injections into the opened relation or the cheapest one, with most-benign-looking or typical endpoints (seed 0).

Stronger evaluation of adversarial training. Table 56 re-attacks CA-AT with three random restarts of a 200-step gradient attack, an AUC-targeted gradient loss, and 40 extra random allocations in the query attack; the last column is R-GCN + CARS-D in the same cell.

Sensitivity of CARS-D. Table 57 varies the number of degree groups ($K_f \in \{4, 8, 16\}$), the candidate pool (100, 300, 1000), and the threshold grid (default or coarse $\{0, 2, 8, 34, \infty\}$), and reports the range of the evaluated worst case of the kept detector, the number of distinct configurations kept, and the median planning time (seed 0).

Table 21. Selection accuracy of CARS-D: runs, share picking the test-optimal candidate, runs where a relation should be kept, kept, both, and mean / largest regret.

Dataset	runs	optimal	should	kept	both	mean reg.	max reg.
YelpChi	45	100%	15	15	15	.000	.000
Amazon	45	93%	1	2	0	.000	.007
FDCompCN	45	89%	4	6	1	.001	.014
S-FFSD	45	100%	15	15	15	.000	.000
Tolokers	45	87%	8	2	2	.001	.010

Table 22. CARS-D with a fraction of the validation labels: median validation fraud nodes, mean, 5% quantile and worst test difference with the graph-free candidate, share of draws below it, and mean regret.

Labels	n_+	mean Δ	5%	worst	below	regret
5%	26	.005	−.014	−.099	9.4%	.002
10%	51	.005	−.008	−.076	8.9%	.002
25%	128	.006	−.002	−.064	5.3%	.001
50%	256	.007	−.000	−.052	4.1%	.001
100%	513	.007	.000	−.014	3.3%	.000

Table 23. Selection rules at equal evaluation budget: number of validated candidates K (median), mean test worst-case AUC difference with the graph-free detector and number of cases more than 0.002 below it, per graph family (full validation labels), mean regret to the best evaluated configuration on relation-rich graphs, and the difference and share of draws below the graph-free detector with 10% of the validation labels on relation-rich graphs (20 draws).

		Real		Augmented real		Relation-rich CSBM			Rich, 10% labels	
Rule	K	Δ	below	Δ	below	Δ	below	regret	Δ	below
CARS-D	2	.000	2/16	−.001	1/12	+.112	2/52	+.011	+.110	6%
CARS-D, cost-blind planning	2	.000	0/16	−.003	3/12	+.076	1/52	+.047	+.074	1%
without validation (plug-in subset)	1	−.005	4/16	−.078	4/12	+.101	10/52	+.022	+.101	19%
without validation (plug-in gates)	1	−.044	8/16	−.203	11/12	+.076	14/52	+.047	+.076	27%
without gates	2	.000	2/16	−.001	2/12	+.112	0/52	+.011	+.110	3%
without graph-free option	2	−.002	4/16	−.078	3/12	+.103	11/52	+.020	+.102	21%
VAL-FA	2	.000	0/16	.000	0/12	+.065	0/52	+.058	+.063	1%
VAL-RAND-3	3	−.001	2/16	.000	0/12	+.037	1/52	+.086	+.033	6%
VAL-RAND-8	8	−.002	4/16	.000	0/12	+.067	1/52	+.056	+.061	9%
VAL-SUB	16	.000	2/16	–	–	+.060	0/20	+.001	+.054	4%
VAL-POOL	13	−.002	4/16	−.004	4/12	+.121	2/52	+.002	+.116	8%

Table 24. Per-graph paired differences in test worst-case AUC (full validation labels): CARS-D minus the graph-free detector, minus cost-blind planning, minus the plug-in's gated choice without validation, minus CARS-D without the gated candidate, minus VAL-RAND-8, and minus VAL-POOL; mean with 95% bootstrap interval; seeds and cases (seeds × cells × families).

Graph	seeds	cases	vs graph-free	vs cost-blind	vs no validation	vs no gated candidate	vs VAL-RAND-8	vs VAL-POOL
Synthetic	1	4	+.005 [.000, +.016]	+.005 [.000, +.016]	+.010 [.000, +.031]	.000 [.000, .000]	+.005 [.000, +.016]	.000 [.000, .000]
FDCompCN	1	4	−.004 [−.011, .000]	−.004 [−.011, .000]	.000 [.000, .000]	+.001 [.000, +.002]	+.006 [−.005, +.020]	+.007 [.000, +.019]
Tolokers	1	4	.000 [.000, .000]	.000 [.000, .000]	+.017 [−.004, +.037]	.000 [.000, .000]	.000 [.000, .000]	.000 [.000, .000]
Amazon	1	4	−.001 [−.002, .000]	.000 [−.001, .000]	+.151 [+.001, +.435]	−.001 [−.002, .000]	.000 [−.001, .000]	.000 [.000, .000]
FDCompCN+A	1	4	−.004 [−.010, .000]	+.004 [−.011, +.020]	+.050 [.000, +.119]	+.001 [.000, +.002]	−.004 [−.010, .000]	+.008 [+.001, +.020]
Tolokers+A	1	4	.000 [.000, .000]	.000 [.000, .000]	+.069 [+.046, +.092]	.000 [.000, .000]	.000 [.000, .000]	.000 [.000, .000]
Amazon+A	1	4	.000 [.000, .000]	.000 [.000, .000]	+.487 [+.046, +.927]	.000 [.000, .000]	.000 [.000, .000]	.000 [.000, .000]
CSBM, R=4	5	20	+.050 [+.006, +.126]	+.010 [.000, +.019]	+.008 [.000, +.019]	.000 [.000, .000]	+.003 [−.004, +.013]	−.010 [−.025, .000]
CSBM, R=8	5	20	+.139 [+.012, +.267]	+.056 [+.009, +.105]	+.029 [+.002, +.056]	.000 [−.006, +.007]	+.064 [+.012, +.117]	.000 [.000, .000]
CSBM, R=12	3	12	+.172 [+.122, +.212]	+.047 [−.038, +.141]	+.097 [+.023, +.170]	+.001 [.000, +.004]	+.084 [+.036, +.159]	−.022 [−.038, .000]

Table 25. Which relation to harden: share of cases where each rule picks a relation whose hardening raises the exact trap threshold of RLR most, mean share of the best gain obtained, and hit rate on graphs with at least 8 relations.

Rule	right	value	right ($R \geq 8$)
exposure per unit cost (no attack)	15/16	94%	100%
uniform-cost exact analysis	4/16	25%	11%
most informative relation	5/16	31%	22%
largest weight norm	4/16	25%	11%
cheapest relation	7/16	44%	33%

Table 26. Decisions under cost uncertainty: box coverage and, for each rule, the mean exact worst-case AUC difference with LR at the true costs and the number of draws more than 0.002 below LR, by log half-width w of the box.

		Naive		Maximin		Minimax regret	
w	covered	Δ	below	Δ	below	Δ	below
0	0%	−.067	11/40	−.067	11/40	−.028	9/40
0.5	2%	−.067	11/40	−.036	8/40	−.041	7/40
1	32%	−.067	11/40	−.017	4/40	−.023	4/40
2	90%	−.067	11/40	.000	1/40	.000	0/40
3	100%	−.067	11/40	.000	0/40	.000	0/40

Table 27. CARS-D (and CARS+AT for trees) minus its graph-free counterpart: mean evaluated worst-case AUC difference with 95% bootstrap interval. Bold: interval above zero; underlined: interval below zero.

	RLR + CARS-D		R-GCN + CARS-D		HGB + CARS-D		CARS+AT (HGB)	
Dataset	P1, B=2	P2, B=20	P1, B=2	P2, B=20	P1, B=2	P2, B=20	P1, B=2	P2, B=20
YelpChi	+.000 [+.000, +.000]	**+.021 [+.019, +.024]**	+.000 [+.000, +.000]	**+.061 [+.057, +.066]**	+.000 [+.000, +.000]	**+.017 [+.015, +.020]**	+.000 [+.000, +.000]	**+.017 [+.015, +.020]**
Amazon	+.000 [+.000, +.000]	+.000 [+.000, +.000]	+.000 [+.000, +.000]	+.000 [+.000, +.000]	−.002 [−.006, +.001]	+.000 [+.000, +.000]	−.001 [−.005, +.002]	+.000 [+.000, +.000]
FDCompCN	−.002 [−.012, +.009]	+.000 [−.000, +.000]	+.000 [+.000, +.000]	+.000 [+.000, +.000]	−.000 [−.003, +.003]	+.000 [+.000, +.000]	−.000 [−.003, +.003]	+.000 [+.000, +.000]
S-FFSD	+.000 [+.000, +.000]	**+.020 [+.018, +.023]**	+.000 [+.000, +.000]	**+.070 [+.062, +.079]**	+.000 [+.000, +.000]	**+.018 [+.015, +.021]**	+.000 [+.000, +.000]	**+.018 [+.015, +.021]**
Tolokers	+.000 [+.000, +.000]	+.000 [+.000, +.000]	+.001 [−.001, +.005]	+.001 [−.001, +.005]	+.000 [+.000, +.000]	+.000 [+.000, +.000]	+.000 [+.000, +.000]	+.000 [+.000, +.000]

Table 28. Conclusions under random cost vectors (RLR, exact attack, 20 vectors per distribution, $B \in \{2, 20\}$): share of trapped cases among those where RLR has a clean gain, number of cases where CARS-D is more than 0.002 below LR and its worst difference, number of gains and their mean, median of the smallest cost (injection or deletion) of the relations kept when it gains, and median of the smallest cost over all actions.

Graph	Distribution	cases	trapped	below	worst	gains	mean gain	kept cost	all costs
Amazon	independent	40	0%	4	−.004	2	+.003	2.12	0.25
Amazon	structured	40	0%	8	−.002	0	–	–	0.25
FDCompCN	independent	40	0%	9	−.023	0	–	–	0.25
FDCompCN	structured	40	0%	25	−.014	0	–	–	0.25
S-FFSD	independent	40	100%	0	−.001	11	+.014	2.25	0.25
S-FFSD	structured	40	100%	0	.000	18	+.010	0.50	0.25
Synthetic	independent	40	90%	0	.000	10	+.087	1.88	0.25
Synthetic	structured	40	98%	0	.000	7	+.051	0.50	0.25
Tolokers	independent	40	75%	0	.000	7	+.036	1.75	1.00
Tolokers	structured	40	100%	0	.000	0	–	–	0.75
YelpChi	independent	40	100%	0	.000	10	+.010	0.88	0.25
YelpChi	structured	40	100%	0	.000	25	+.005	0.50	0.25

Table 29. Confirmatory comparisons of CARS-D with its graph-free counterpart (five seeds): mean difference, seed standard deviation, 95% hierarchical-bootstrap interval, Holm-adjusted p, non-inferiority at margin 0.01, and smallest margin with non-inferiority.

Dataset	Family	Cell	difference	SD	95% interval	p_{Holm}	NI(0.01)	margin
YelpChi	RLR	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
YelpChi	RLR	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
YelpChi	RLR	P2, B=20	+.0213	0.0026	[+.0187, +.0240]	0.00	yes	0.0000
YelpChi	R-GCN	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
YelpChi	R-GCN	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
YelpChi	R-GCN	P2, B=20	+.0612	0.0039	[+.0568, +.0658]	0.00	yes	0.0000
YelpChi	HGB	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
YelpChi	HGB	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
YelpChi	HGB	P2, B=20	+.0174	0.0028	[+.0150, +.0203]	0.00	yes	0.0000
Amazon	RLR	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
Amazon	RLR	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Amazon	RLR	P2, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Amazon	R-GCN	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
Amazon	R-GCN	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Amazon	R-GCN	P2, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Amazon	HGB	P1, B=2	−.0019	0.0032	[−.0058, +.0005]	1.00	yes	0.0058
Amazon	HGB	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Amazon	HGB	P2, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
FDCompCN	RLR	P1, B=2	−.0017	0.0083	[−.0122, +.0091]	1.00	no	0.0122
FDCompCN	RLR	P1, B=20	+.0000	0.0000	[−.0000, +.0000]	1.00	yes	0.0000
FDCompCN	RLR	P2, B=20	+.0000	0.0000	[−.0000, +.0000]	1.00	yes	0.0000
FDCompCN	R-GCN	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
FDCompCN	R-GCN	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
FDCompCN	R-GCN	P2, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
FDCompCN	HGB	P1, B=2	−.0001	0.0003	[−.0034, +.0025]	1.00	yes	0.0034
FDCompCN	HGB	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
FDCompCN	HGB	P2, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
S-FFSD	RLR	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
S-FFSD	RLR	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
S-FFSD	RLR	P2, B=20	+.0203	0.0021	[+.0179, +.0228]	0.00	yes	0.0000
S-FFSD	R-GCN	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
S-FFSD	R-GCN	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
S-FFSD	R-GCN	P2, B=20	+.0698	0.0093	[+.0617, +.0789]	0.00	yes	0.0000
S-FFSD	HGB	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
S-FFSD	HGB	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
S-FFSD	HGB	P2, B=20	+.0179	0.0013	[+.0152, +.0205]	0.00	yes	0.0000
Tolokers	RLR	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
Tolokers	RLR	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Tolokers	RLR	P2, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Tolokers	R-GCN	P1, B=2	+.0009	0.0020	[−.0012, +.0046]	1.00	yes	0.0012
Tolokers	R-GCN	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Tolokers	R-GCN	P2, B=20	+.0008	0.0019	[−.0013, +.0045]	1.00	yes	0.0013
Tolokers	HGB	P1, B=2	0	0.0000	[0, 0]	1.00	yes	0.0000
Tolokers	HGB	P1, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000
Tolokers	HGB	P2, B=20	0	0.0000	[0, 0]	1.00	yes	0.0000

Table 30. Cost-agnostic planning (RLR, exact, seed 0, 16 cost scales): mean and worst difference with LR of the plan at the lower corner, mean difference of the plan at the true costs, mean and largest price of not knowing the costs, and cells below LR.

Dataset	B	agn. mean	agn. worst	true mean	price mean	price max	below
Amazon	2	.000	.000	.000	.000	.001	0/16
Amazon	20	.000	.000	.000	.000	.000	0/16
FDCompCN	2	.000	.000	−.003	−.003	.000	0/16
FDCompCN	20	.000	.000	−.004	−.004	.000	0/16
S-FFSD	2	.000	.000	.005	.005	.016	0/16
S-FFSD	20	.000	.000	.001	.001	.010	0/16
Tolokers	2	.000	.000	.015	.015	.047	0/16
Tolokers	20	.000	.000	.000	.000	.000	0/16
YelpChi	2	.000	.000	.007	.007	.029	0/16
YelpChi	20	.000	.000	.001	.001	.006	0/16

Table 31. Noisy cost estimates (RLR + CARS-D, P1, B=2, ten draws per σ): plan at the true costs, LR, and for box half-widths $z \in \{0, 1, 2\}$ the draws below LR and the mean regret to the true plan.

Dataset	σ	true	LR	z=0		z=1		z=2	
Amazon	0.5	.969	.969	0/10	.000	0/10	.000	0/10	.000
Amazon	1	.969	.969	4/10	.027	1/10	.002	0/10	.000
FDCompCN	0.5	.663	.677	10/10	.000	6/10	−.006	1/10	−.013
FDCompCN	1	.663	.677	5/10	.125	5/10	.059	1/10	−.013
S-FFSD	0.5	.643	.643	3/10	.091	1/10	.035	0/10	.000
S-FFSD	1	.643	.643	4/10	.103	1/10	.035	0/10	.000
Tolokers	0.5	.716	.716	1/10	.018	0/10	.000	0/10	.000
Tolokers	1	.716	.716	3/10	.053	0/10	.000	0/10	.000
YelpChi	0.5	.771	.771	4/10	.255	1/10	.064	0/10	.000
YelpChi	1	.771	.771	7/10	.374	1/10	.060	0/10	.000

Table 32. Break-even costs (P1, RLR, seed 0): clean gain over LR of the single-relation model, its exact trap threshold, and the cost multiplier needed to resist B=2 and B=20.

Dataset	Relation	gain	threshold	× for 2	× for 20
Amazon	U-P-U	−.001	–	–	no clean gain
Amazon	U-S-U	+.003	1	2	20
Amazon	U-V-U	−.001	–	–	no clean gain
FDCompCN	C-I-C	−.005	–	–	no clean gain
FDCompCN	C-P-C	−.028	–	–	no clean gain
FDCompCN	C-S-C	−.006	–	–	no clean gain
S-FFSD	T-S-T	+.018	2	1	10
S-FFSD	T-M-T	−.001	–	–	no clean gain
S-FFSD	T-L-T	−.001	–	–	no clean gain
S-FFSD	T-Y-T	−.000	–	–	no clean gain
Tolokers	W-T-W	+.047	1	2	20
YelpChi	R-U-R	+.022	2	1	10
YelpChi	R-T-R	+.000	–	–	no clean gain
YelpChi	R-S-R	+.010	1	2	20

Table 33. S-FFSD rolling windows: evaluated worst-case AUC by window and cell (seed 0).

Window	Cell	MLP	R-GCN	R-GCN+D	HGB-X	HGB-Graph	HGB+D
0	P1, B=2	.531	.325	.531	.619	.308	.619
0	P2, B=20	.531	.132	.758	.619	.213	.717
1	P1, B=2	.572	.284	.572	.619	.272	.619
1	P2, B=20	.572	.064	.680	.619	.122	.731
2	P1, B=2	.648	.165	.648	.610	.190	.610
2	P2, B=20	.648	.050	.828	.610	.119	.784
3	P1, B=2	.461	.079	.461	.686	.108	.686
3	P2, B=20	.461	.161	.862	.686	.100	.812

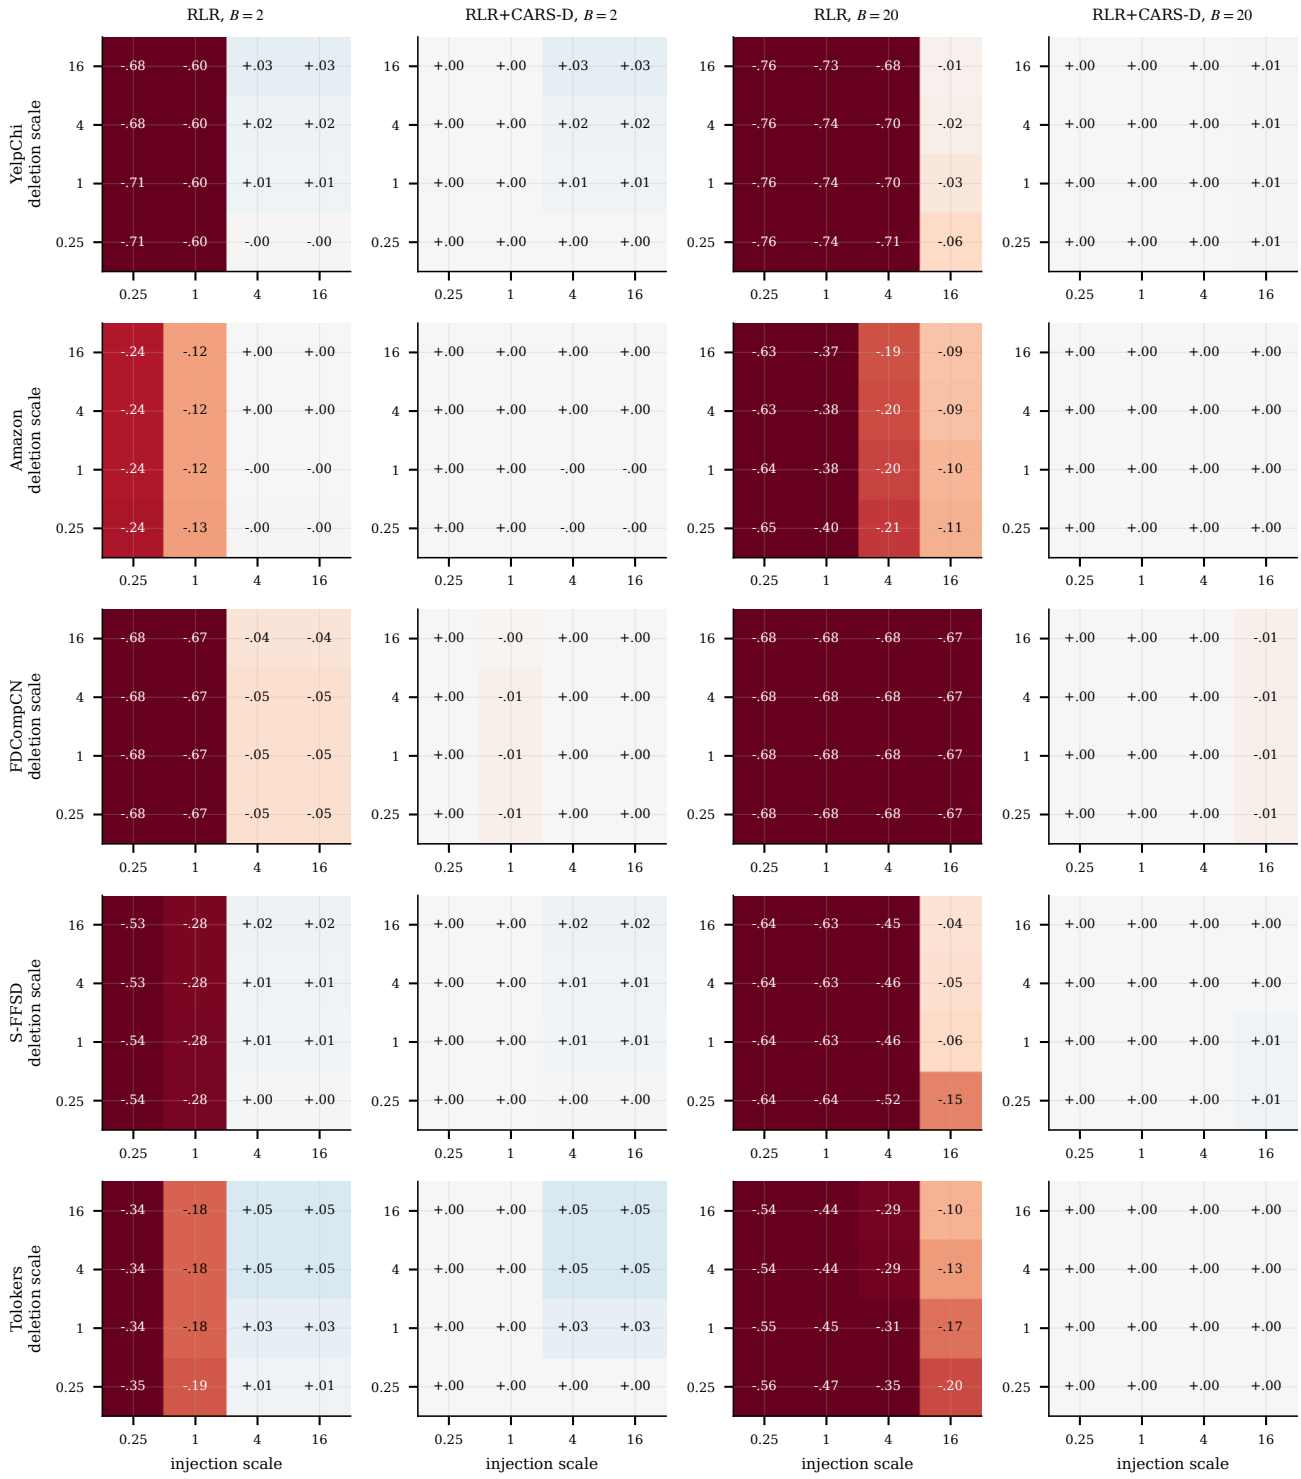

Figure 2. Exact cost maps: worst-case AUC minus LR of the clean-trained RLR and of RLR + CARS-D planned at the true costs, over injection and deletion cost scales relative to P1, at B=2 and B=20 (seed 0).

Table 34. Noisy planning costs (RLR + CARS-D, seed 0, five draws): plan at the true costs, mean and minimum of the noisy plans evaluated exactly at the true costs, LR, and share of draws with the same selection.

Dataset	B	σ	noise	true plan	noisy (mean)	noisy (min)	LR	same
Amazon	2	0.5	common	.969	.969	.969	.969	0%
Amazon	2	0.5	independent	.969	.969	.969	.969	20%
Amazon	2	1	common	.969	.945	.847	.969	20%
Amazon	2	1	independent	.969	.940	.851	.969	0%
Amazon	20	0.5	common	.969	.969	.969	.969	40%
Amazon	20	0.5	independent	.969	.969	.969	.969	60%
Amazon	20	1	common	.969	.969	.969	.969	20%
Amazon	20	1	independent	.969	.969	.969	.969	0%
FDCompCN	2	0.5	common	.663	.663	.663	.677	100%
FDCompCN	2	0.5	independent	.663	.663	.663	.677	100%
FDCompCN	2	1	common	.663	.671	.663	.677	40%
FDCompCN	2	1	independent	.663	.405	.003	.677	20%
FDCompCN	20	0.5	common	.677	.677	.677	.677	100%
FDCompCN	20	0.5	independent	.677	.677	.677	.677	100%
FDCompCN	20	1	common	.677	.677	.677	.677	100%
FDCompCN	20	1	independent	.677	.677	.677	.677	100%
S-FFSD	2	0.5	common	.643	.643	.643	.643	0%
S-FFSD	2	0.5	independent	.643	.460	.289	.643	0%
S-FFSD	2	1	common	.643	.587	.361	.643	0%
S-FFSD	2	1	independent	.643	.494	.289	.643	0%
S-FFSD	20	0.5	common	.643	.643	.643	.643	100%
S-FFSD	20	0.5	independent	.643	.643	.643	.643	100%
S-FFSD	20	1	common	.643	.643	.643	.643	100%
S-FFSD	20	1	independent	.643	.643	.643	.643	100%
Tolokers	2	0.5	common	.716	.716	.716	.716	60%
Tolokers	2	0.5	independent	.716	.680	.538	.716	0%
Tolokers	2	1	common	.716	.680	.538	.716	40%
Tolokers	2	1	independent	.716	.644	.538	.716	20%
Tolokers	20	0.5	common	.716	.716	.716	.716	100%
Tolokers	20	0.5	independent	.716	.716	.716	.716	100%
Tolokers	20	1	common	.716	.716	.716	.716	100%
Tolokers	20	1	independent	.716	.716	.716	.716	100%
YelpChi	2	0.5	common	.771	.517	.134	.771	60%
YelpChi	2	0.5	independent	.771	.517	.134	.771	40%
YelpChi	2	1	common	.771	.524	.134	.771	20%
YelpChi	2	1	independent	.771	.271	.134	.771	0%
YelpChi	20	0.5	common	.771	.771	.771	.771	80%
YelpChi	20	0.5	independent	.771	.771	.771	.771	60%
YelpChi	20	1	common	.771	.771	.771	.771	60%
YelpChi	20	1	independent	.771	.771	.771	.771	40%

Table 35. HGB + CARS-D planned at a profile with costs divided by ρ and evaluated under true profiles P1/P2 (evaluated worst-case AUC, seed 0); last column: HGB-X.

Dataset	B	plan	ρ=1	ρ=2	ρ=4	ρ=8	HGB-X
YelpChi	2	P1	.933/.933	.933/.933	.933/.933	.933/.933	.933
YelpChi	2	P2	.417/.948	.417/.948	.417/.948	.417/.948	.933
YelpChi	20	P1	.933/.933	.933/.933	.933/.933	.933/.933	.933
YelpChi	20	P2	.269/.948	.933/.933	.933/.933	.933/.933	.933

Table 36. Exact trap thresholds of RLR (seed 0): single relations (in the order of Table 3), the full model, the sum of single-relation thresholds, and the budget at which every subset is trapped.

Dataset	Profile	single relations	full	sum	all subsets
Amazon	P1	0, 1, 0	1	1	1
Amazon	P2	0, 1, 0	1	1	1
FDCompCN	P1	0, 0, 0	0	0	0
FDCompCN	P2	0, 0, 0	0	0	0
S-FFSD	P1	2, 0, 0, 0	1	2	2
S-FFSD	P2	25, 0, 0, 0	1	25	25
Tolokers	P1	1	1	1	1
Tolokers	P2	8	8	8	8
YelpChi	P1	2, 1, 1	1	4	2
YelpChi	P2	25, 1, 1	1	27	25

Table 37. Gangs (P1, seed 0): realized AUC at average budgets $\bar{B} \in \{0.5, 1, 2\}$ for oracle and surrogate allocation, by coordination cost κ; last column: graph-free trees.

Dataset	Model	κ	oracle			surrogate			free
Tolokers	HGB-CARS-D	0	.729	.729	.729	–	–	–	.729
Tolokers	HGB-CARS-D	0.5	.729	.729	.729	–	–	–	.729
Tolokers	HGB-CARS-D	1	.729	.729	.729	–	–	–	.729
Tolokers	HGB-Graph	0	.625	.561	.509	.652	.596	.540	.729
Tolokers	HGB-Graph	0.5	.659	.590	.524	.681	.622	.559	.729
Tolokers	HGB-Graph	1	.682	.614	.540	.702	.643	.576	.729
Tolokers	R-GCN	0	.700	.661	.615	.710	.674	.632	.729
Tolokers	R-GCN	0.5	.711	.671	.622	.719	.683	.638	.729
Tolokers	R-GCN	1	.720	.680	.630	.727	.690	.644	.729
YelpChi	HGB-CARS-D	0	.933	.933	.933	–	–	–	.933
YelpChi	HGB-CARS-D	0.5	.933	.933	.933	–	–	–	.933
YelpChi	HGB-CARS-D	1	.933	.933	.933	–	–	–	.933
YelpChi	HGB-Graph	0	.766	.610	.421	.784	.632	.438	.933
YelpChi	HGB-Graph	0.5	.801	.668	.473	.816	.691	.489	.933
YelpChi	HGB-Graph	1	.825	.710	.526	.839	.731	.546	.933
YelpChi	R-GCN	0	.673	.433	.048	.681	.455	.088	.933
YelpChi	R-GCN	0.5	.723	.529	.170	.728	.545	.198	.933
YelpChi	R-GCN	1	.756	.594	.283	.760	.604	.309	.933

Table 38. OLS coefficients of the trap depth on the CSBM factors (architecture dummies relative to HGB-Graph).

Term	coefficient
R-GCN (vs HGB-Graph)	-0.028
RLR (vs HGB-Graph)	-0.157
p_{id}	-0.190
isolated fraction	-0.334
$\log c_{\mathrm{id}}$	+0.254
$p_{\mathrm{id}} \times$ isolated	+0.173
intercept	+0.028
R^2	0.83

Table 39. R-GCN + CARS-D (wall time) and CA-AT by number of epochs (evaluated worst-case AUC, P1, B=2, seed 0).

Dataset	Cell	R-GCN + CARS-D	CA-AT (epochs: AUC)
FDCompCN	P1, B=2	.702 (16s)	37: .453, 50: .490, 100: .266, 200: .531, 400: .472
Synthetic	P1, B=2	.814 (21s)	36: .867, 50: .885, 100: .902, 200: .894, 400: .889
Tolokers	P1, B=2	.736 (50s)	28: .693, 50: .739, 100: .737, 200: .724, 400: .700

Table 40. Synthetic CSBM, c_{id}=0.5: evaluated worst-case AUC (three seeds). Bottom: the attacks on RLR, whose exact worst case is known.

Model	clean	B=1	2	4	8	16
MLP (graph-free)	.681	.681	.681	.681	.681	.681
LR (graph-free)	.702	.702	.702	.702	.702	.702
HGB-X (graph-free)	.616	.616	.616	.616	.616	.616
RLR	.893	.518	.170	.028	.004	.000
R-GCN	.993	.756	.280	.067	.012	.002
HGB-Graph	.888	.508	.389	.167	.042	.011
CA-AT (R-GCN)	.991	.952	.919	.880	.755	.486
AT-RLR (exact)	.855	.775	.725	.571	.394	.455
RLR + CARS	.843	.802	.702	.702	.702	.702
R-GCN + CARS	.951	.922	.828	.681	.681	.681
HGB + CARS	.820	.717	.530	.616	.616	.616
R-GCN + CARS-D	.951	.922	.828	.681	.681	.681
CARS+AT	.991	.952	.919	.880	.755	.681
Attack on RLR	**query**	**transfer**	**gradient**	**exact**		
RLR, B=1	.727	.753	.659	.518		
RLR, B=2	.610	.608	.415	.170		
RLR, B=4	.462	.486	.231	.028		
RLR, B=8	.422	.359	.108	.004		
RLR, B=16	.318	.211	.023	.000		

Table 41. Mean margin (AUC) over the graph-free counterpart across the three cells: plug-in subset (S), plug-in gated proposal (G), and configuration kept by the validation check (D).

	YelpChi			Amazon			FDCompCN			S-FFSD			Tolokers		
Family	S	G	D	S	G	D	S	G	D	S	G	D	S	G	D
RLR	.007	-0.020	.007	-0.006	-0.205	.000	-0.002	-0.048	-0.001	-0.040	-0.074	.007	.000	-0.011	.000
R-GCN	.020	-0.009	.020	-0.010	-0.058	.000	-0.015	-0.064	.000	-0.130	-0.069	.023	.000	-0.004	.001
HGB	.006	-0.028	.006	-0.000	-0.014	-0.001	-0.001	-0.006	-0.000	-0.126	-0.091	.006	.000	-0.036	.000

Table 42. S-FFSD, temporal split: clean and evaluated worst-case AUC (mean over three seeds). Models whose training depends on the cost cell (CA-HGB and the CARS variants, +C and +D) report the clean AUC of their P1, $B{=}2$ instance, as in Table 1.

Model	clean	P1, $B{=}2$	P1, $B{=}20$	P2, $B{=}20$
LR	.425	.425	.425	.425
MLP	.560	.560	.560	.560
HGB-X	.593	.593	.593	.593
RLR	.361	.202	.035	.091
R-GCN	.811	.145	.105	.178
HGB	.820	.186	.130	.178
CA-HGB	.791	.334	.230	.513
RLR+C	.361	.202	.425	.423
R-GCN+C	.811	.145	.560	.765
HGB+C	.820	.186	.593	.793
RLR+D	.425	.425	.425	.431
R-GCN+D	.560	.560	.560	.809
HGB+D	.593	.593	.593	.766

Table 43. Node-level exposure bounds: violations, median ratio of measured damage to bound, and Spearman correlation between bound and measured damage across nodes and across relations.

Dataset	action	n	viol.	ratio	ρ_{nodes}	$\rho_{\text{rel.}}$
Amazon	delete	1211	0	0.030	0.05	.500
Amazon	inject	1974	0	0.008	0.89	.500
FDCompCN	delete	35	0	0.000	-0.10	1.000
FDCompCN	inject	1344	0	0.032	0.68	.500
S-FFSD	delete	2225	0	0.018	0.42	1.000
S-FFSD	inject	3200	0	0.030	0.52	.800
Tolokers	delete	767	0	nan	0.59	–
Tolokers	inject	800	0	0.040	0.92	–
YelpChi	delete	1634	0	0.041	0.32	1.000
YelpChi	inject	2400	0	0.059	0.95	1.000

Table 44. CSBM thresholds as a function of the identity cost c_{id}.

c_{id}	B_{trap} (num.)	envelope	exposure LB	linear resp.	abandon (num.)	$\sum_r \kappa_r$
0.25	2.1	2.1	0.6	1.3	25.1	26.4
0.5	3.0	3.0	1.2	2.6	25.8	27.2
1	4.6	4.6	2.4	5.2	27.0	28.8
2	7.8	7.8	4.9	10.3	29.6	32.0
4	14.2	14.2	9.8	20.6	35.0	38.4
8	27.0	27.0	17.0	29.6	46.3	51.1
16	30.4	30.5	17.2	29.6	70.3	76.6

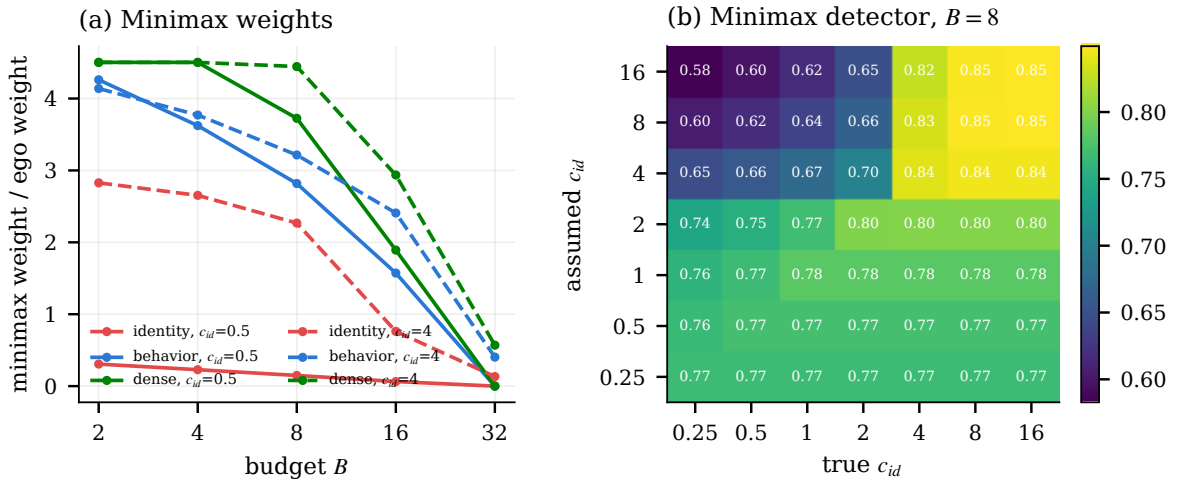

Figure 3. (a) Minimax weights relative to the ego weight as the budget grows. (b) Worst-case AUC at B=8 under the true identity cost (columns) of the minimax detector built for an assumed cost (rows).

Table 45. Fidelity of the plug-in model across configurations.

Dataset	Cell	configs	ρ(plug, test)	ρ(val, test)	regret plug	regret val	overest.
Amazon	P1, B=2	20	0.23	0.98	0.004	0.002	+0.035
Amazon	P1, B=20	20	0.11	1.00	0.039	0.000	+0.122
Amazon	P2, B=20	20	0.43	1.00	0.004	0.000	+0.116
FDCompCN	P1, B=2	20	0.52	0.71	0.017	0.017	+0.147
FDCompCN	P1, B=20	20	0.69	0.97	0.004	0.000	+0.217
FDCompCN	P2, B=20	20	0.72	0.88	0.000	0.026	+0.196
S-FFSD	P1, B=2	32	0.68	1.00	0.097	0.000	+0.101
S-FFSD	P1, B=20	32	0.68	1.00	0.000	0.000	+0.138
S-FFSD	P2, B=20	32	0.81	1.00	0.041	0.003	+0.071
Synthetic	P1, B=2	20	0.82	1.00	0.000	0.000	+0.103
Synthetic	P1, B=20	20	0.86	0.99	0.000	0.000	+0.016
Synthetic	P2, B=20	20	0.87	1.00	0.000	0.000	+0.037
Tolokers	P1, B=2	6	0.77	1.00	0.009	0.000	+0.086
Tolokers	P1, B=20	6	0.64	0.94	0.000	0.000	+0.058
Tolokers	P2, B=20	6	0.83	1.00	0.009	0.000	+0.082
YelpChi	P1, B=2	20	0.66	0.99	0.092	0.000	+0.001
YelpChi	P1, B=20	20	0.80	1.00	0.001	0.000	+0.079
YelpChi	P2, B=20	20	0.71	1.00	0.002	0.000	+0.093

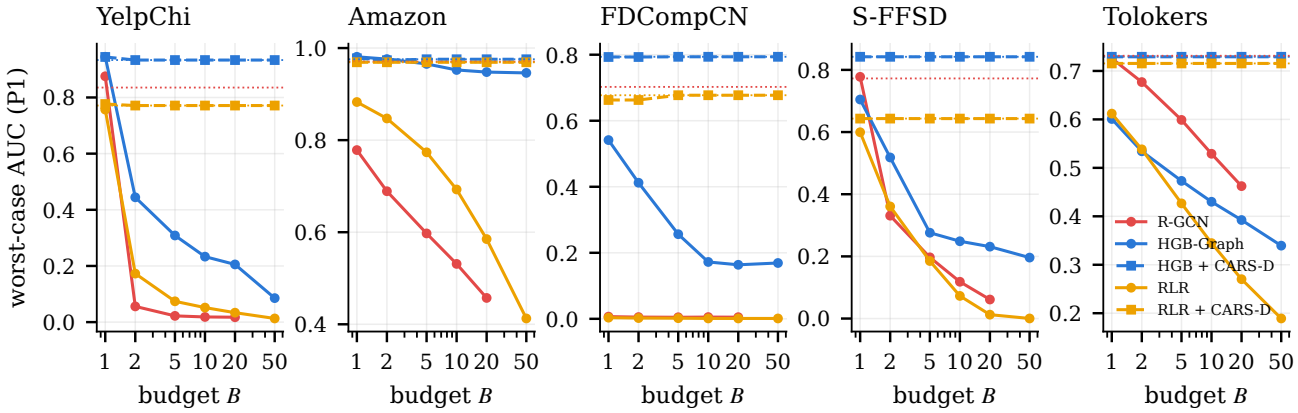

Figure 4. Evaluated worst-case AUC at P1 as a function of the per-node budget (seed 0). Solid: clean-trained; dashed: CARS-D planned at each budget; dotted: graph-free counterparts (HGB-X, MLP, LR). Gradient attacks are run up to B=20, so R-GCN is shown up to that budget.

Table 46. Paired comparison with the graph-free counterpart over seeds, and operating metrics under attack (ours / counterpart).

Dataset	Cell	Model	seeds	diff.	95% CI	at/above	AP	TPR@1%FPR
Amazon	P1, B=2	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.872 / .872	.769 / .769
Amazon	P1, B=20	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.872 / .872	.769 / .769
Amazon	P2, B=20	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.872 / .872	.769 / .769
FDCompCN	P1, B=2	RLR + CARS-D	5	-0.002	[-0.012, +0.009]	3/5	.207 / .224	.036 / .041
FDCompCN	P1, B=20	RLR + CARS-D	5	+0.000	[-0.000, +0.000]	5/5	.224 / .224	.041 / .041
FDCompCN	P2, B=20	RLR + CARS-D	5	+0.000	[-0.000, +0.000]	5/5	.224 / .224	.041 / .041
S-FFSD	P1, B=2	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.261 / .261	.002 / .002
S-FFSD	P1, B=20	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.261 / .261	.002 / .002
S-FFSD	P2, B=20	RLR + CARS-D	5	+0.020	[+0.018, +0.023]	5/5	.382 / .261	.084 / .002
Tolokers	P1, B=2	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.372 / .372	.042 / .042
Tolokers	P1, B=20	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.372 / .372	.042 / .042
Tolokers	P2, B=20	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.372 / .372	.042 / .042
YelpChi	P1, B=2	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.390 / .390	.099 / .099
YelpChi	P1, B=20	RLR + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.390 / .390	.099 / .099
YelpChi	P2, B=20	RLR + CARS-D	5	+0.021	[+0.018, +0.025]	5/5	.400 / .390	.090 / .099
Amazon	P1, B=2	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.889 / .889	.818 / .818
Amazon	P1, B=20	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.889 / .889	.818 / .818
Amazon	P2, B=20	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.889 / .889	.818 / .818
FDCompCN	P1, B=2	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.282 / .282	.074 / .074
FDCompCN	P1, B=20	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.282 / .282	.074 / .074
FDCompCN	P2, B=20	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.282 / .282	.074 / .074
S-FFSD	P1, B=2	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.582 / .582	.314 / .314
S-FFSD	P1, B=20	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.582 / .582	.314 / .314
S-FFSD	P2, B=20	R-GCN + CARS-D	5	+0.070	[+0.058, +0.081]	5/5	.686 / .582	.420 / .314
Tolokers	P1, B=2	R-GCN + CARS-D	5	+0.001	[-0.002, +0.003]	5/5	.391 / .385	.035 / .032
Tolokers	P1, B=20	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.385 / .385	.032 / .032
Tolokers	P2, B=20	R-GCN + CARS-D	5	+0.001	[-0.002, +0.003]	5/5	.391 / .385	.035 / .032
YelpChi	P1, B=2	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.525 / .525	.178 / .178
YelpChi	P1, B=20	R-GCN + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.525 / .525	.178 / .178
YelpChi	P2, B=20	R-GCN + CARS-D	5	+0.061	[+0.056, +0.066]	5/5	.637 / .525	.238 / .178
Amazon	P1, B=2	HGB + CARS-D	5	-0.002	[-0.006, +0.002]	3/5	.902 / .906	.810 / .822
Amazon	P1, B=20	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.906 / .906	.822 / .822
Amazon	P2, B=20	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.906 / .906	.822 / .822
FDCompCN	P1, B=2	HGB + CARS-D	5	-0.000	[-0.001, +0.000]	5/5	.412 / .413	.169 / .167
FDCompCN	P1, B=20	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.413 / .413	.167 / .167
FDCompCN	P2, B=20	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.413 / .413	.167 / .167
S-FFSD	P1, B=2	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.678 / .678	.430 / .430
S-FFSD	P1, B=20	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.678 / .678	.430 / .430
S-FFSD	P2, B=20	HGB + CARS-D	5	+0.018	[+0.016, +0.020]	5/5	.710 / .678	.448 / .430
Tolokers	P1, B=2	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.388 / .388	.027 / .027
Tolokers	P1, B=20	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.388 / .388	.027 / .027
Tolokers	P2, B=20	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.388 / .388	.027 / .027
YelpChi	P1, B=2	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.783 / .783	.458 / .458
YelpChi	P1, B=20	HGB + CARS-D	5	+0.000	[+0.000, +0.000]	5/5	.783 / .783	.458 / .458
YelpChi	P2, B=20	HGB + CARS-D	5	+0.017	[+0.014, +0.021]	5/5	.822 / .783	.489 / .458
Amazon	P1, B=2	CARS+AT	3	+0.000	[+0.000, +0.000]	3/3	.891 / .891	.818 / .818
Amazon	P1, B=20	CARS+AT	3	+0.000	[+0.000, +0.000]	3/3	.891 / .891	.818 / .818
Amazon	P2, B=20	CARS+AT	3	+0.000	[+0.000, +0.000]	3/3	.891 / .891	.818 / .818
FDCompCN	P1, B=2	CARS+AT	5	+0.000	[+0.000, +0.000]	5/5	.282 / .282	.074 / .074
FDCompCN	P1, B=20	CARS+AT	5	+0.000	[+0.000, +0.000]	5/5	.282 / .282	.074 / .074
FDCompCN	P2, B=20	CARS+AT	5	+0.000	[+0.000, +0.000]	5/5	.282 / .282	.074 / .074
S-FFSD	P1, B=2	CARS+AT	3	+0.000	[+0.000, +0.000]	3/3	.567 / .567	.291 / .291
S-FFSD	P1, B=20	CARS+AT	3	+0.000	[+0.000, +0.000]	3/3	.567 / .567	.291 / .291
S-FFSD	P2, B=20	CARS+AT	3	+0.074	[+0.047, +0.100]	3/3	.682 / .567	.418 / .291
Tolokers	P1, B=2	CARS+AT	5	+0.001	[-0.002, +0.003]	5/5	.391 / .385	.035 / .032
Tolokers	P1, B=20	CARS+AT	5	+0.000	[+0.000, +0.000]	5/5	.385 / .385	.032 / .032
Tolokers	P2, B=20	CARS+AT	5	+0.001	[-0.002, +0.003]	5/5	.391 / .385	.035 / .032
YelpChi	P1, B=2	CARS+AT	3	+0.000	[+0.000, +0.000]	3/3	.528 / .528	.180 / .180
YelpChi	P1, B=20	CARS+AT	3	+0.000	[+0.000, +0.000]	3/3	.528 / .528	.180 / .180
YelpChi	P2, B=20	CARS+AT	3	+0.059	[+0.052, +0.067]	3/3	.633 / .528	.240 / .180
Amazon	P1, B=2	CARS+AT-HGB	5	-0.001	[-0.006, +0.003]	4/5	.901 / .906	.812 / .822
Amazon	P1, B=20	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.906 / .906	.822 / .822
Amazon	P2, B=20	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.906 / .906	.822 / .822
FDCompCN	P1, B=2	CARS+AT-HGB	5	-0.000	[-0.001, +0.000]	5/5	.412 / .413	.169 / .167
FDCompCN	P1, B=20	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.413 / .413	.167 / .167
FDCompCN	P2, B=20	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.413 / .413	.167 / .167
S-FFSD	P1, B=2	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.678 / .678	.430 / .430
S-FFSD	P1, B=20	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.678 / .678	.430 / .430
S-FFSD	P2, B=20	CARS+AT-HGB	5	+0.018	[+0.016, +0.020]	5/5	.710 / .678	.448 / .430
Tolokers	P1, B=2	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.388 / .388	.027 / .027
Tolokers	P1, B=20	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.388 / .388	.027 / .027
Tolokers	P2, B=20	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.388 / .388	.027 / .027
YelpChi	P1, B=2	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.783 / .783	.458 / .458
YelpChi	P1, B=20	CARS+AT-HGB	5	+0.000	[+0.000, +0.000]	5/5	.783 / .783	.458 / .458
YelpChi	P2, B=20	CARS+AT-HGB	5	+0.017	[+0.014, +0.021]	5/5	.822 / .783	.489 / .458

Table 47. Median wall-clock time (seconds) of CARS-D (three families) and of one CA-AT model.

Dataset	CARS-D	CA-AT	ratio
YelpChi	187	1427	7.6
Amazon	221	1148	5.2
FDCompCN	45	58	1.3
S-FFSD	180	749	4.2
Tolokers	42	173	4.1

Table 48. Random cost vectors: number of traps observed, accuracy of the plug-in trap prediction, and rank correlation between the plug-in margin and the observed margin of RLR.

Dataset	vectors	traps RLR	acc.	traps HGB	acc.	Spearman
Amazon	24	0	0.58	0	0.58	0.71
FDCompCN	24	0	0.00	0	0.00	0.79
S-FFSD	24	24	1.00	24	1.00	0.72
Tolokers	24	20	0.79	12	0.88	1.00
YelpChi	24	24	1.00	22	0.92	0.76

Table 49. Evaluated worst-case AUC under three attacker knowledge levels (seed 0).

Dataset	Cell	Model	labels	no labels	black box
FDCompCN	P1, B=2	HGB + CARS-D	.794	.794	.794
FDCompCN	P1, B=2	HGB-Graph	.411	.420	.764
FDCompCN	P1, B=2	R-GCN	.005	.005	.264
FDCompCN	P1, B=2	R-GCN + CARS-D	.702	.702	.702
FDCompCN	P1, B=2	RLR	.002	.002	.370
FDCompCN	P1, B=2	RLR + CARS-D	.663	.663	.672
FDCompCN	P1, B=20	HGB + CARS-D	.794	.794	.794
FDCompCN	P1, B=20	HGB-Graph	.159	.162	.653
FDCompCN	P1, B=20	R-GCN	.005	.005	.186
FDCompCN	P1, B=20	R-GCN + CARS-D	.702	.702	.702
FDCompCN	P1, B=20	RLR	.001	.001	.321
FDCompCN	P1, B=20	RLR + CARS-D	.677	.677	.677
Tolokers	P1, B=2	HGB + CARS-D	.729	.729	.729
Tolokers	P1, B=2	HGB-Graph	.538	.539	.737
Tolokers	P1, B=2	R-GCN	.673	.674	.755
Tolokers	P1, B=2	R-GCN + CARS-D	.735	.735	.758
Tolokers	P1, B=2	RLR	.538	.538	.700
Tolokers	P1, B=2	RLR + CARS-D	.716	.716	.716
Tolokers	P1, B=20	HGB + CARS-D	.729	.729	.729
Tolokers	P1, B=20	HGB-Graph	.484	.483	.686
Tolokers	P1, B=20	R-GCN	.484	.490	.698
Tolokers	P1, B=20	R-GCN + CARS-D	.731	.731	.731
Tolokers	P1, B=20	RLR	.270	.269	.633
Tolokers	P1, B=20	RLR + CARS-D	.716	.716	.716
YelpChi	P1, B=2	HGB + CARS-D	.933	.933	.933
YelpChi	P1, B=2	HGB-Graph	.444	.434	.742
YelpChi	P1, B=2	R-GCN	.055	.055	.237
YelpChi	P1, B=2	R-GCN + CARS-D	.835	.835	.835
YelpChi	P1, B=2	RLR	.173	.172	.667
YelpChi	P1, B=2	RLR + CARS-D	.771	.771	.771
YelpChi	P1, B=20	HGB + CARS-D	.933	.933	.933
YelpChi	P1, B=20	HGB-Graph	.204	.200	.610
YelpChi	P1, B=20	R-GCN	.016	.016	.120
YelpChi	P1, B=20	R-GCN + CARS-D	.835	.835	.835
YelpChi	P1, B=20	RLR	.032	.032	.680
YelpChi	P1, B=20	RLR + CARS-D	.771	.771	.771

Table 50. YelpChi, entity-level attacks (evaluated worst-case AUC, 3 seeds).

Model	clean	Q1, B=2	Q1, B=5	Q2, B=5
MLP	.837	–	–	–
HGB-X	.935	–	–	–
RLR	.796	.781	.762	.746
R-GCN	.916	.520	.436	.850
HGB-Graph	.959	.777	.749	.935
RLR + CARS-D	.776	.769	.769	.790
R-GCN + CARS-D	.857	.837	.837	.896
HGB + CARS-D	.940	.935	.935	.951

Table 51. Persistent camouflage: evaluated worst-case AUC after poisoning (c: concentrated, d: diversified endpoints).

Model	FDCo.c	FDCo.d	Tolo.c	Tolo.d	Yelp.c	Yelp.d
LR	.677	.677	.716	.716	.771	.771
MLP	.702	.702	.731	.731	.835	.835
HGB-X	.794	.794	.729	.729	.933	.933
RLR	.012	.001	.870	.586	.872	.580
R-GCN	.777	.006	.995	.722	.976	.601
HGB	.991	.859	.984	.667	.968	.866
CA-HGB	.000	.615	.989	.731	.991	.950
RLR+C	.000	.663	.870	.716	.889	.771
R-GCN+C	.008	.662	.995	.731	.921	.835
HGB+C	.991	.790	.984	.729	.971	.933
RLR+D	.677	.663	.870	.716	.889	.771
R-GCN+D	.702	.702	.995	.755	.931	.835
HGB+D	.991	.790	.984	.729	.971	.933

Table 52. Sybil attacks: evaluated worst-case AUC for Sybil edge cost c_s and budget B.

Model	FDCo S1,2	FDCo S1,5	FDCo S4,20	Tolo S1,2	Tolo S1,5	Tolo S4,20	Yelp S1,2	Yelp S1,5	Yelp S4,20
LR	.677	.677	.677	.716	.716	.716	.771	.771	.771
MLP	.702	.702	.702	.731	.731	.731	.835	.835	.835
HGB-X	.794	.794	.794	.729	.729	.729	.933	.933	.933
RLR	.003	.002	.002	.571	.483	.456	.166	.131	.129
R-GCN	.005	.004	.004	.693	.621	.596	.040	.016	.009
HGB	.270	.178	.178	.541	.483	.473	.058	.034	.033
GNNGuard	.023	.011	.011	.679	.611	.607	.400	.391	.381
BWGNN	.204	.197	.200	.773	.768	.756	.659	.642	.640
RLR+D	.677	.677	.677	.716	.716	.716	.771	.771	.769
R-GCN+D	.702	.702	.702	.737	.731	.731	.835	.835	.833
HGB+D	.794	.794	.794	.729	.729	.729	.933	.933	.933

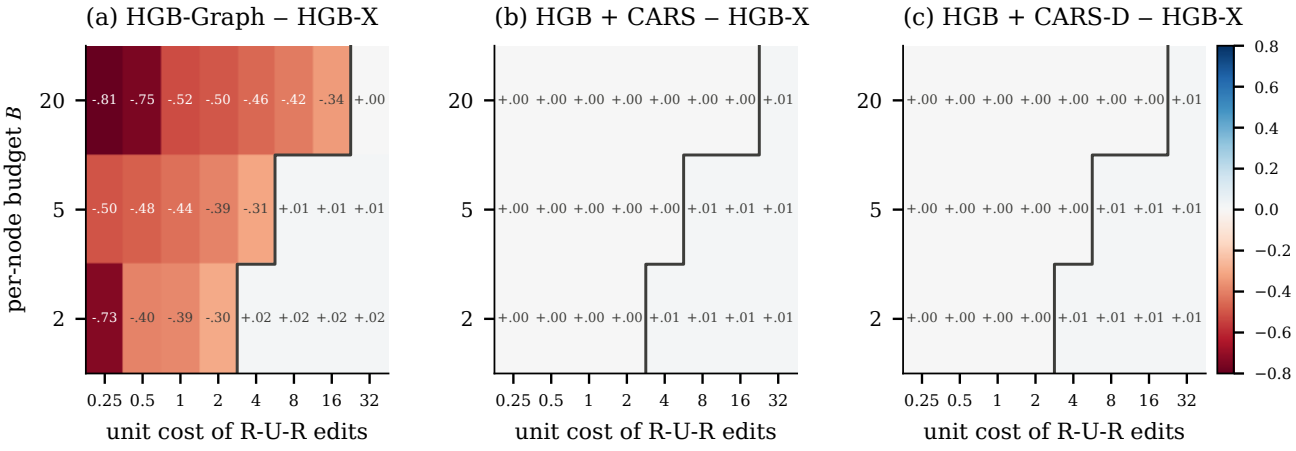

Figure 5. YelpChi: evaluated worst-case AUC minus graph-free trees for HGB-Graph, HGB + CARS, and HGB + CARS-D, each planned per cell; the line marks $B = c$.

Table 53. Evaluated worst-case AUC under feature attacks of radius $\epsilon \in \{0, 0.02, 0.05, 0.1\}$, combined with structural attacks for graph detectors.

Dataset	Model	0	0.02	0.05	0.1
YelpChi	LR	.771	.752	.721	.665
YelpChi	MLP	.835	.809	.767	.686
YelpChi	HGB-X	.933	.841	.749	.661
YelpChi	RLR	.173	.161	.146	.121
YelpChi	R-GCN	.056	.055	.050	.044
YelpChi	HGB	.444	.353	.315	.292
YelpChi	RLR+D	.771	.752	.721	.665
YelpChi	R-GCN+D	.835	.809	.767	.686
YelpChi	HGB+D	.933	.841	.749	.661
FDCompCN	LR	.677	.376	.137	.056
FDCompCN	MLP	.702	.401	.170	.104
FDCompCN	HGB-X	.794	.107	.051	.018
FDCompCN	RLR	.002	.002	.002	.002
FDCompCN	R-GCN	.005	.005	.005	.005
FDCompCN	HGB	.412	.058	.026	.007
FDCompCN	RLR+D	.663	.408	.176	.076
FDCompCN	R-GCN+D	.702	.401	.170	.104
FDCompCN	HGB+D	.793	.130	.062	.022
Tolokers	LR	.716	.697	.669	.625
Tolokers	MLP	.731	.688	.630	.560
Tolokers	HGB-X	.729	.605	.476	.442
Tolokers	RLR	.538	.528	.512	.486
Tolokers	R-GCN	.677	.661	.638	.599
Tolokers	HGB	.534	.485	.464	.462
Tolokers	RLR+D	.716	.697	.669	.625
Tolokers	R-GCN+D	.736	.705	.658	.593
Tolokers	HGB+D	.729	.605	.476	.442

Table 54. Label scarcity: clean AUC of the graph-free baselines and evaluated worst-case AUC at P1, B=2 / P2, B=20.

		graph-free (clean)		R-GCN		HGB-Graph		R-GCN + CARS-D		HGB + CARS-D	
Dataset	labels	MLP	HGB-X	P1,2	P2,20	P1,2	P2,20	P1,2	P2,20	P1,2	P2,20
YelpChi	0.05	.770	.869	.069	.565	.416	.782	.770	.856	.869	.908
YelpChi	0.1	.800	.893	.097	.662	.427	.857	.800	.873	.893	.926
YelpChi	0.2	.818	.920	.046	.651	.440	.902	.818	.885	.920	.944
FDCompCN	0.05	.602	.699	.001	.000	.695	.499	.602	.602	.699	.699
FDCompCN	0.1	.679	.760	.001	.000	.667	.433	.679	.679	.765	.760
FDCompCN	0.2	.694	.781	.001	.001	.628	.415	.694	.694	.782	.781
Tolokers	0.05	.718	.697	.579	.578	.475	.475	.718	.718	.697	.697
Tolokers	0.1	.719	.706	.586	.583	.437	.437	.719	.719	.706	.706
Tolokers	0.2	.731	.721	.621	.619	.506	.505	.731	.731	.721	.721

Table 55. Gate-aware attacks against gated detectors: evaluated worst case of the standard suite, of the gate attack alone, and of both.

Dataset	Cell	Detector	suite	gate	both
Amazon	P1, B=2	HGB+D (kept)	.979	.983	.979
Amazon	P1, B=2	HGB+D (prop.)	.979	.983	.979
Amazon	P1, B=2	R-GCN+D (prop.)	.975	.977	.975
Amazon	P1, B=2	RLR+D (prop.)	.967	.971	.967
Amazon	P1, B=20	HGB+D (prop.)	.959	.968	.959
Amazon	P1, B=20	R-GCN+D (prop.)	.808	.879	.803
Amazon	P1, B=20	RLR+D (prop.)	.405	.935	.406
Amazon	P2, B=20	HGB+D (prop.)	.963	.969	.962
Amazon	P2, B=20	R-GCN+D (prop.)	.830	.905	.828
Amazon	P2, B=20	RLR+D (prop.)	.390	.933	.390
FDCompCN	P1, B=2	HGB+D (kept)	.793	.794	.793
FDCompCN	P1, B=2	HGB+D (prop.)	.793	.794	.793
FDCompCN	P1, B=2	R-GCN+D (prop.)	.662	.664	.662
FDCompCN	P1, B=2	RLR+D (prop.)	.663	.668	.663
FDCompCN	P1, B=20	HGB+D (prop.)	.794	.794	.794
FDCompCN	P1, B=20	R-GCN+D (prop.)	.707	.707	.707
FDCompCN	P1, B=20	RLR+D (prop.)	.673	.677	.673
FDCompCN	P2, B=20	HGB+D (prop.)	.794	.794	.794
FDCompCN	P2, B=20	R-GCN+D (prop.)	.707	.707	.707
FDCompCN	P2, B=20	RLR+D (prop.)	.677	.677	.677
S-FFSD	P1, B=2	HGB+D (prop.)	.651	.714	.633
S-FFSD	P1, B=2	R-GCN+D (prop.)	.663	.785	.662
S-FFSD	P1, B=2	RLR+D (prop.)	.554	.621	.554
Tolokers	P1, B=2	HGB+D (prop.)	.723	.742	.721
Tolokers	P1, B=2	R-GCN+D (kept)	.744	.756	.744
Tolokers	P1, B=2	R-GCN+D (prop.)	.744	.756	.744
Tolokers	P1, B=2	RLR+D (prop.)	.719	.728	.719
Tolokers	P2, B=20	HGB+D (prop.)	.722	.742	.720
Tolokers	P2, B=20	R-GCN+D (kept)	.744	.756	.744
Tolokers	P2, B=20	R-GCN+D (prop.)	.744	.756	.744
Tolokers	P2, B=20	RLR+D (prop.)	.719	.728	.719
YelpChi	P1, B=2	HGB+D (prop.)	.867	.886	.863
YelpChi	P1, B=2	R-GCN+D (prop.)	.775	.799	.772
YelpChi	P1, B=2	RLR+D (prop.)	.683	.742	.683

Table 56. CA-AT under the standard suite and under stronger attacks (evaluated worst-case AUC; synthetic: 3 seeds).

Dataset	Cell	suite	AUC loss	restarts	random alloc.	all	CARS-D
Synthetic	P1, B=4	.887	.853	.874	.925	.842	.681
Tolokers	P1, B=2	.762	.762	.760	.770	.746	.741
Tolokers	P2, B=20	.764	.764	.763	.769	.752	.741
YelpChi	P1, B=2	.752	.759	.757	.808	.731	.836
YelpChi	P2, B=20	.901	.896	.895	.915	.892	.897

Table 57. Sensitivity of CARS-D to its hyperparameters: minimum and maximum worst-case AUC of the kept detector, number of distinct choices, median planning and checking time (seconds).

Dataset	Cell	Family	min	max	choices	s
Amazon	P1, B=2	HGB	.977	.980	2	123
Amazon	P1, B=2	RLR	.969	.969	1	123
Amazon	P2, B=20	HGB	.966	.976	3	98
Amazon	P2, B=20	RLR	.969	.969	1	98
S-FFSD	P1, B=2	HGB	.842	.842	1	180
S-FFSD	P1, B=2	RLR	.643	.643	1	180
S-FFSD	P2, B=20	HGB	.860	.860	1	108
S-FFSD	P2, B=20	RLR	.661	.661	1	108
Tolokers	P1, B=2	HGB	.729	.729	1	12
Tolokers	P1, B=2	RLR	.716	.716	1	12
Tolokers	P2, B=20	HGB	.729	.729	1	10
Tolokers	P2, B=20	RLR	.716	.716	1	10
YelpChi	P1, B=2	HGB	.933	.933	1	131
YelpChi	P1, B=2	RLR	.771	.771	1	131
YelpChi	P2, B=20	HGB	.948	.948	1	90
YelpChi	P2, B=20	RLR	.794	.794	1	90